

Unfriending can facilitate cooperation: Co-evolution of opinion and network dynamics

Mikkel Werling (201706722)

Abstract

Discovering what conditions can prevent political polarization is a vital and pressing issue, as political polarization can critically harm democratic institutions (Boxell, Gentzkow, & Shapiro, 2020; Levin, Milner, & Perrings, 2021). The co-evolution of opinion and network dynamics are known to be interacting processes (Bener, Çağlayan, Henry, & Prałat, 2016; McPherson, Smith-Lovin, & Cook, 2001; Kossinets & Watts, 2009). Nevertheless, previous models of opinion dynamics have largely neglected co-evolution in their models (Flache et al., 2017; Galesic, Olsson, Dalege, van der Does, & Stein, 2021). Moreover, the field of opinion dynamics suffers from not being able to incorporate empirical data in their theoretical models, which severely limits the field's ability to make reliable explanations and predictions (Flache et al., 2017; Galesic et al., 2021; Mäs, 2019). This thesis tries to fill both of these gaps in the literature by developing and investigating a co-evolutionary agent-based model of opinion and networks dynamics. Cutting edge techniques of hyperparameter optimization are used to investigate how well the patterns of generated networks of the model correspond to the patterns of 7 empirical social networks. The results show that triadic closure can explain the patterns of empirical networks better than the most used network generating algorithms. Moreover, when the empirical social networks are large and highly opinionated, co-evolutionary models are much better explanations than models that do not include opinion dynamics. Contrary to recent findings (Sasahara et al., 2021), the main result of this thesis is that avoiding polarization can be facilitated by the deletion of ties between dissimilar agents. This thesis offers novel perspectives on possible remedies for polarization, and improves upon the ability of previous models to explain how social networks are created.

Keywords: *agent-based modelling, opinion dynamics, social influence, co-evolution, social networks*

Contents

I	Introduction	6
1	Agent-based modelling of complex phenomena	10
2	Central Mechanisms	11
2.1	Homophily	11
2.2	Social Influence	13
3	Opinion Dynamics	14
3.1	Dynamical networks	14
4	Network Dynamics	15
4.1	Triadic Closure	16
5	Co-evolution via tie-deletion	18
5.1	Tie-deletions effect on cooperation	18
6	Model Description	20
6.1	Conceptual overview	21
6.2	Assumptions	21
6.3	The stages of the model	22
6.4	Initialization	23
6.5	Dynamics	24
6.5.1	Network dynamic stage	24

6.5.2	Opinion dynamic stage	25
6.5.3	Co-evolutionary stage	27
6.6	Data availability	27
II	Integrating empirical data	28
7	Data	30
8	Optimization methods	30
8.1	Specifying model parameters	32
9	Optimization results	33
10	Discussion of optimization	37
10.1	When co-evolution is a better explanation	37
10.2	The impact of neighborhoods	38
10.3	Including empirical data in agent-based models	39
III	Co-evolution's effect on opinion dynamics	41
11	Outcome metrics	41
11.1	Time-dependent Metrics	42
11.2	Final State Metrics	42
11.3	Correlations	42
12	Model Parameters	42

13 Results	43
13.1 The effect of the different model parameters	44
13.2 Correlations	46
14 Discussion	47
14.1 Reiterating previous literature	47
14.2 The importance of co-evolution	48
14.3 Tie-deletion facilitates cooperation	48
14.4 Distance in opinion and network space	49
15 The point of no return for polarization	54
IV Overall Discussion	54
16 Contrast to previous work	56
17 Lacking realism of the model	56
17.1 Non-cognitive agents	57
17.2 Shrinking Variance	58
17.3 Social influence exclusively	58
18 A broader perspective	58
19 Conclusion	59
20 Acknowledgements	60

References	60
A Appendix	71
A.1 Supplementary material for hyperparameter optimization	71
A.1.1 Diagnostic plots	71
A.1.2 Metrics of model performance	78
A.2 Supplementary information for co-evolution's effect on opinion dynamics	80
A.2.1 Negative tie-deletion	84
A.3 The effect of initial opinions	85

Part I

Introduction

Democracies around the world are experiencing increased amounts of political polarization (Boxell et al., 2020; McCoy, Rahman, & Somer, 2018; Somer & McCoy, 2018). As a result, cooperating with other people becomes harder, which can severely damage the ability of democratic systems to solve problems (Andris et al., 2015; Levin et al., 2021; McCoy et al., 2018). A striking recent example is the uniquely severe rise of polarization in the political system of the United States (Dimock & Wike, 2020). During the last two decades, the amount of overlap of opinions between the two political parties have decreased substantially, which has led to gridlock, government shutdowns and failure to enact new legislation (see Figure 1) (Andris et al., 2015; Center, 2014b). The cost of polarization was seen most severely by the fact that the global COVID-19 pandemic failed to generate a common effort from the two parties. Instead, preventative measures such as masks became party crests (Macy, Ma, Tabin, Gao, & Szymanski, 2021).

A related concept to political polarization is affective polarization. Affective polarization refers to the extent that citizens experience negative feelings for political parties other than their own (Boxell et al., 2020; Iyengar, Lelkes, Levendusky, Malhotra, & Westwood, 2019). Once again, the United States is reported to have uniquely high amounts of affective polarization (Boxell et al., 2020). Although the United States is unique in the extent of affective polarization, they are not the only country experiencing an increase in affective polarization. Countries such as France, Switzerland and Denmark have had increased levels of affective polarization in the last two decades (Boxell et al., 2020). The increase in polarization is in other words a general and global trend (McCoy et al., 2018; Somer & McCoy, 2018; Wilson, Parker, & Feinberg, 2020). With that said, polarization of opinions does not seem to be an inevitable state for democracies. During the last two decades, European democracies such as Sweden, Norway and Germany have experienced higher degrees of consensus and a decrease in affective polarization (Boxell et al., 2020).

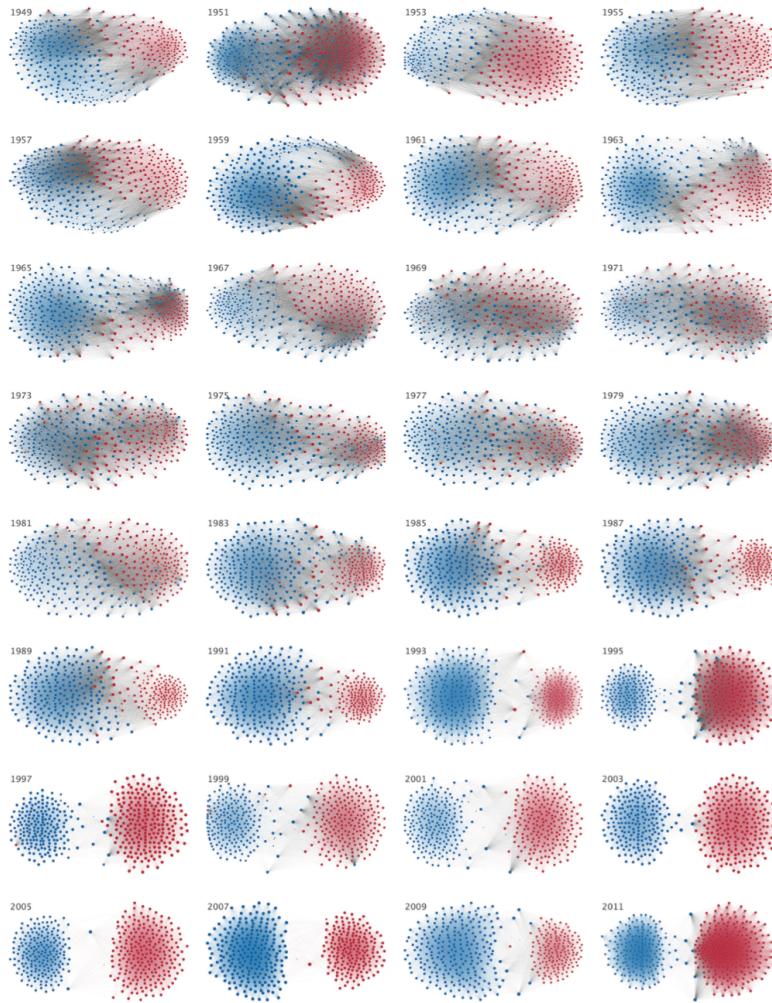


Figure 1: Polarization and collaboration in the House of Representatives in the United States. Adapted from Andris et al. (2015). Each node represents a member of the U.S. House of Representatives from 1949–2012. Each panel shows different two-year Congresses with the starting year displayed on top. Republican (R) representatives are in red and Democrat (D) representatives are in blue. When two members agree on a number of votes above the Congress' threshold, and edge is drawn between them. This threshold value reflects the number of agreements, where any pair exhibiting this number is equally likely to be comprised of two members of the same party (e.g. D-D or R-R), or a cross-party pair (e.g. D-R). The size of the nodes reflect their degree, with higher degree nodes being larger.

Despite the decrease in affective polarization in certain northern European countries, polarization globally is on the rise (Boxell et al., 2020; McCoy et al., 2018; Somer & McCoy, 2018) and there is reason to believe that the general problem of polarization might be getting worse due to technological advances. Specifically, a growing literature has reported that information is not distributed uniformly across the users of the internet and social media (Taylor, Mantzaris, & Garibay, 2018; Sasahara et al., 2021; Baumann, Lorenz-Spreen, Sokolov, & Starnini, 2020; Tsai, Tao, Chuan, & Hong, 2020). Instead, information is being catered to the individual based

on their previous search history and behavior on social media (Geschke, Lorenz, & Holtz, 2019). Very different people will therefore be exposed to very different pieces of information, as users are only exposed to a small subset of the available information. Heavily influenced by confirmation bias, this can lead to circumstances where almost all opinions one is exposed to are congruent with one's prior beliefs. Such circumstances are popularized under the term echo chambers on social media. Echo chambers are often cited as one of the main reasons for the increase in political polarization (Baumann et al., 2020; Sasahara et al., 2021; Tsai et al., 2020; Geschke et al., 2019).

Because of technology's impact on the level of polarization, the internet has been called a "mixed blessing" (Lev-On & Manin, 2009). On the one hand, individuals can access more diverse information via the internet, which can help decrease polarization. On the other hand, the internet also enables encounters of extremely opposing opinions by chance, which can increase polarization (Lev-On & Manin, 2009). Previous work paint a murky picture concerning the effect of being presented with very different opinions. The work of Levy (2021) suggests that when people are presented with diversifying views, they generally tend to decrease negative affects to other political parties. The opposite effect is found in Bail et al. (2018), where exposure to directly opposing views can cause opinions to become more polarized than they were before exposure. On a similar note, higher involvement in one's echo chamber correlates with more negative emotions in terms of valence, which suggests that echo chambers are driven largely by outrage. In other words, the most active users in an echo chamber are the ones who are most displeased with the opinions of people outside their echo chamber (Del Vicario et al., 2016).

Although polarization is usually described in the context of political polarization, polarization is not necessarily a political phenomenon. In fact, polarization can critically impair any social system's ability to cooperate successfully (Levin et al., 2021). It is therefore vital that we identify and understand what mechanisms can lead a system to polarization, to consensus, or to anywhere in between the two extremes. This is no easy task. Opinion formation and opinion dynamics are highly complex phenomena (Baumann, 2021). Despite the complexity of the problem, an entire field has been dedicated to study how opinions change. The study of the dynamics shaping opinions is aptly named opinion dynamics. The formal models of opinion dynamics often neglect the fact that agents can not only change their opinions but also change their peers (Flache et al., 2017). The models of opinion dynamics typically assume a static network structure in which agents are situated (Galesic et al., 2021). By doing so, these models fail to take into account that networks are dynamic and the structure of the networks is likely co-evolving with opinions (de Arruda et al., 2022; Galesic et al., 2021). The interdependence between the dynamics of social networks and opinion dynamics is largely understudied, despite the fact that the two processes are highly interacting (Asikainen, Iñiguez, Ureña-Carrión, Kaski, & Kivelä, 2020; Bruch &

Atwell, 2015; Galesic et al., 2021; Kossinets & Watts, 2009; Noorazar, Vixie, Talebanpour, & Hu, 2020). It is therefore not surprising that the dynamic relation between the network structure and social influence has often been reported as an important avenue for future research (Flache et al., 2017; Galesic et al., 2021).

Another important avenue to gain insights into is how to incorporate empirical data in theoretical models (Mäs, 2019). This is an especially pertinent problem in the field of opinion dynamics, which often fails to account for how the results of their models match empirical data (Galesic et al., 2021; Flache et al., 2017; Mäs, 2019). The lack of empirical integration is understandable, as opinions are notoriously hard to measure (Mäs, 2019). Moreover, most of the models of opinion dynamics are not meant to accurately predict how opinions evolve over time, but rather point out interesting interactions between key variables in formal models (Mäs, 2019). Such theoretical models are useful for theory building, but they often run the risk of being too detached from the real world to be useful (Smaldino, 2020; Mäs, 2019).

This thesis seeks to fill these two critical gaps in the current literature on opinion dynamics: the need to integrate both empirical data and the co-evolution of networks and opinion dynamics. This thesis is organized in four parts. The introduction provides an overview of previous methods and vital mechanisms to consider. This is done first by explaining how agent-based models are well-equipped to investigate complex systems. Next, the critical underlying mechanisms of opinion and network dynamics are identified and introduced. A co-evolutionary agent-based model of opinion and network dynamics are build on the basis of these mechanisms, drawing on insights from data science, social psychology, computational biology and network science. The model is specifically designed to help illuminate the co-evolution of opinion and network dynamics. The second part of the thesis integrates empirical data into the agent-based modelling process by investigating how well the agent-based model of the thesis can generate the patterns found in real world networks. To do so, cutting-edge techniques of hyperparameter optimization are employed. This thesis expands upon previous use of this technique in agent-based modelling by also implementing a functional analysis of variance (fANOVA) of the hyperparameter importance (Hutter, Hoos, & Leyton-Brown, 2014). The third part is focused on how co-evolution impacts opinion dynamics. It does so by exploring how polarization of the system is influenced by the model parameters. Beyond echoing the previous findings of the literature, this section focuses on understanding how tie-deletion to dissimilar individuals might facilitate cooperation. Finally, the fourth part of the paper discusses the overall findings and presents the key limitations of the model. We begin by discussing why agent-based models are particularly well-suited to model opinion dynamics.

1 Agent-based modelling of complex phenomena

The social sciences face the daunting task of trying to explain and predict the behavior of extremely complex systems. Historically, social science have explained such systems by relying on the power of language using verbal models (Smaldino, 2020). Verbal models typically explain a system by providing useful analogies, which can help illuminate the system of interest. These analogies are often ambiguous and several interpretations of the same model are possible. The benefit of this ambiguity is that verbal models often incorporate many realistic social and cognitive mechanisms (Fogarty, Ammar, Holding, Powell, & Kandler, 2022). But ambiguity is often something to be avoided in science, as clarity and precision is a necessity for developing useful theories of reality (Smaldino, 2020). Formal models offer an alternative to verbal models, as they offer the kind of precision that is lacking in verbal models at the cost of some realism. Typical examples of formal models are mathematical models, where different variables describe forces in the system of interest (Fogarty et al., 2022). Such models are especially hard to build for complex systems as the different constituents of the system are hard to reduce to single components (Smaldino, 2020; Poli, 2013). Formal models of complex systems often solve this problem by making simplifying assumptions, which reduces the realism of the models.

Agent-based models are well-equipped to strike a balance between the precision of formal models and the realism of verbal models (Flache, 2018; Galesic et al., 2021; Epstein, 1999; Mäs, Flache, & Kitts, 2014). This is done by investigating the macro-behaviors of a system, where the micro-behaviors are specified (Bruch & Atwell, 2015; Epstein, 1999; Flache, 2018). In particular, agent-based models are well-suited when analyzing systems where more is known about interactions between individuals instead of interactions between variables (Geschke et al., 2019). As is the case with any model, the results critically hinge on the assumptions of the model. It is the role of the modeler to provide as empirically plausible mechanisms as possible for the system of study (Crooks & Heppenstall, 2012; Epstein, 1999; Page, 2010). At the same time, the value of a model comes from the fact that it is a simplification of reality. The model should be as simple as it can be and as complicated as it needs to be in order to answer its questions (Smaldino, 2020). The hope is that by simplifying the system to a sufficient extent, we can observe and understand some important features of even extremely complex systems (Fogarty et al., 2022; Smaldino, 2020, 2016, 2022).

One of the complex systems that the social sciences have studied for decades is the dynamics governing how opinions are formed and changed. Describing the mechanisms which shape opinion dynamics constitute a complex problem (Mäs, 2019), in the sense that opinion dynamics is governed by multiple interacting systems. The change of opinions is both a cognitive matter of updating one's belief and a social matter of being exposed to different perspectives (Flache et al., 2017; Friedkin & Johnsen, 1990; Spears, 2021).

2 Central Mechanisms

Investigating the complex system of opinion dynamics is the central aim of this paper. In the following sections, the central mechanisms of opinion dynamics are introduced. The specific focus will be on the tendency for similar individuals to cluster together and on how social relations influence the individual's opinion.

2.1 Homophily

One of the most robust findings of the social world is summed up in the now famous phrase "birds of a feather flock together" (McPherson et al., 2001). This phrase refers to the fact that assortment in humans is non-random, and is often based on how similar individuals are (Asikainen et al., 2020; Crandall, Cosley, Huttenlocher, Kleinberg, & Suri, 2008; McPherson et al., 2001). The observation that similarity is predictive of connections is called observed homophily (McPherson et al., 2001; Kossinets & Watts, 2009). In all types of human social networks, observed homophily seems ubiquitous. For all demographic variables which have been investigated, including gender, race, religion, and socioeconomic class, there is always a tendency for more similar individuals to cluster together (Asikainen et al., 2020; McPherson et al., 2001; Taylor et al., 2018).

Observed homophily has mainly been explained using a psychological and a structural explanation. The psychological explanation seeks to explain observed homophily by the fact that people seem to exhibit a psychological preference to be with people like themselves. This psychological preference is referred to as choice homophily (Asikainen et al., 2020; McPherson et al., 2001; Winter & Kataria, 2020). Such a preference might have evolved as interacting with similar people ensures easier communication and enhances the individual's ability to predict the other person's behavior (Kossinets & Watts, 2009; Winter & Kataria, 2020). In this sense, assortment based on similarity could have evolved to facilitate cooperation, because coordination between similar individuals is simpler and less costly for the individual (Winter & Kataria, 2020; Carter, Lee, & Marshall, 2015; Smaldino, 2019).

Although choice homophily seems to be an important piece of the puzzle, observed homophily could also be explained by structural constraints. Such constraints could limit how dissimilar choices of interactions are available (Peixoto, 2022). As an example, if you work as a female nurse, chances are that most of your interactions at work will be with other females. In this case, even if you do not have a strong psychological preference for interacting with people like yourself, the social interactions available to you might primarily be people similar to you. The existing assortment in the network makes the probability of interacting with similar

individuals much higher than with non-similar individuals (Peixoto, 2022). Such structural factors contributing to high levels of observed homophily are referred to as structural homophily (Asikainen et al., 2020; McPherson et al., 2001; Winter & Kataria, 2020).

Observed homophily hints at the prevalent intuition that you have more in common with your friends than you have with strangers (McPherson et al., 2001). Despite its prevalence, few studies have been made to gauge the importance of the effect of homophily on social interactions. One such study was made by Kossinets and Watts (2009), which estimated how observed homophily manifested in real dynamical social networks. They recorded social interactions on a large university over the course of a year. These interactions allowed Kossinets and Watts (2009) to gauge not only how social ties changed over time, but also how similar individuals of the social networks were. Their findings are of significant importance to this thesis. First, the results confirm what one might intuitively expect - you have more in common with your friends than you have with strangers. But their findings suggests that this is in fact a special case of a more general phenomenon; that distance in similarity is a function of distance in the social network (see Figure 2). The further you are removed from someone in the social network, the less you will have in common (Kossinets & Watts, 2009). The connection between distance in similarity and distance in the social network seems to be a robust finding of social networks (Bener et al., 2016; Crandall et al., 2008).

Choice and structural homophily are normally used as separate explanations for the observed homophily in social networks, but they are not mutually exclusive. Rather, it seems that the two processes can facilitate each other (Asikainen et al., 2020). For instance, even low amounts of choice homophily could over time evolve into structural homophily (Asikainen et al., 2020; Kossinets & Watts, 2009; Taylor et al., 2018). In fact, both structural and choice homophily contribute to the link between distance in similarity and distance in social networks. Structural homophily plays a large role as new ties are primarily made to people close to you in the social network (Bianconi, Darst, Iacovacci, & Fortunato, 2014; Peixoto, 2022). When similarity and distance in the social networks are connected, structural homophily will increase the observed homophily of the system. Controlling for both distance in the network and shared social circles, generating new social ties is more likely between similar agents (Kossinets & Watts, 2009; Bener et al., 2016). Similarity is not only predictive of which new social ties will be created, but also which old ties will be deleted. Although homophily has almost exclusively been studied as a process which makes similar individuals more likely to generate new ties (Noel & Nyhan, 2011; Bener et al., 2016), homophily has been shown empirically to also play a vital role in tie-deletion. More similar individuals have much more stable connections, with a much lower propensity to delete ties over time (Noel & Nyhan, 2011; Bener et al., 2016; McPherson et al., 2001; Kossinets & Watts, 2009). Complementing the empirical studies on tie-deletion, the famous computational work of Schelling (1971) shows

that including even subtle tendencies of tie-deletion of dissimilar pairs can create global patterns of network segregation and community structures that are characteristic for many real-world social networks.

The connection between similarity and distance is also driven by this increased propensity to delete ties to dissimilar people (Kossinets & Watts, 2009; Bener et al., 2016). The effect becomes more extreme as interactions between peers tend to make like-minded individuals even more like-minded (Friedkin & Johnsen, 1990; Spears, 2021).

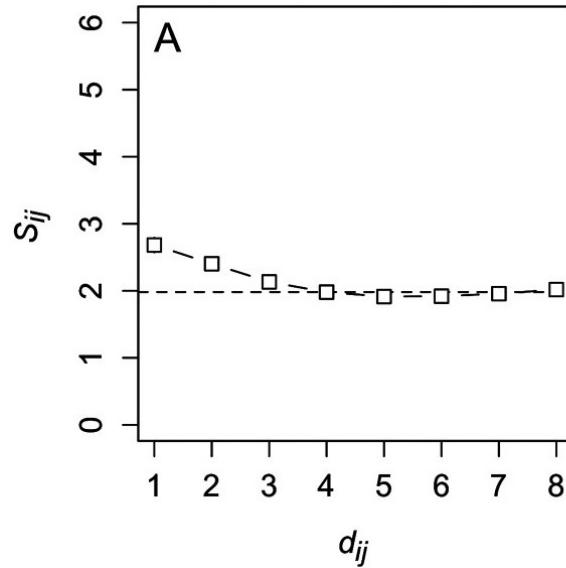


Figure 2: The link between distance and similarity. Adapted from Kossinets and Watts (2009). X-axis shows the distance in the network to other agents. Y-axis shows the average similarity between agents.

2.2 Social Influence

Humans are inherently a social species (Kurzban, Burton-Chellew, & West, 2015). We don't think and form our opinions in a vacuum. Either via social media, debates or discussions with friends, opinions are constantly being shared, debated and discussed with other individuals. As a result of these social interactions, the opinions of the people involved in the interaction might change. This influence that social encounters and social relations has on the opinions of the involved agents is referred to as social influence (Friedkin & Johnsen, 1990). Social influence has been proposed as the most important effect in shaping an individual's opinions (Chacoma & Zanette, 2015; Flache, 2018). In this thesis, we will distinguish between positive and negative social influence. Positive social influence refers to the force which drives similar individuals to become more similar through interactions (Flache et al., 2017; Levin et al., 2021). This can be thought of as instances of reaching a compromise, finding

common ground, or seeing the other person's perspective on a topic. Interactions which are driven by positive social influence will be referred to simply as positive interactions in this thesis. Negative social influence refers to the force which causes dissimilar agents to become more dissimilar as a result of interactions (Flache et al., 2017). Examples include moral outrage or out-group aversion. Interactions which are driven by negative social influence will be referred to as negative interactions in this thesis. It is important to note that although positive influence seems to be a robust empirical finding, negative influence is more elusive (Flache et al., 2017; Takács, Flache, & Maes, 2014; Turner & Smaldino, 2018). With that said, there is evidence that suggests that exposure to opposing views can lead to increased polarization (Bail et al., 2018).

3 Opinion Dynamics

The field of opinion dynamics attempts to identify key mechanisms which shape how opinions change over time using formal and agent-based models (Flache et al., 2017; Flache, 2018; Noorazar et al., 2020). Two of the most robustly identified mechanisms are social influence and homophily (Flache et al., 2017). The models from opinion dynamics are often analogous models, drawing inspiration from statistical physics in the structure and mechanisms of their scientific models (Galesic et al., 2021). The models of opinion dynamics have studied both continuous and discrete opinions, as well as studying systems where agents have multiple opinions of interest (Flache et al., 2017). The typical approach in opinion dynamics is to use empirical studies of psychological and sociological mechanisms to inform the assumptions of their theoretical models. What these models are typically used for is identifying how these empirically justified parameters interact with each other (Baumann, 2021; Chacoma & Zanette, 2015; Flache et al., 2017; Friedkin & Johnsen, 1990; Noorazar et al., 2020; Spears, 2021; Turner & Smaldino, 2018).

3.1 Dynamical networks

By modelling opinion dynamical systems in computational models, researchers in opinion dynamics create simple systems which represent the more complex and messy real social world. These simplified systems are normally studied to identify which conditions can give rise to polarization or consensus in terms of the opinions of simulated agents (Flache et al., 2017). As has already been alluded to in the introduction, previous models have largely only focused on social influence without considering how the social interactions change as a result of changes in opinions (Galesic et al., 2021; Holme & Newman, 2006; Jalili, 2015). Central to most classical models of opinion dynamics is that agents are situated in static, theoretical networks (Flache et al., 2017). The

assumption of static networks is consequential. Social networks are inherently dynamic, with new ties being formed and deleted over time. As has already been established, empirical work suggests that these two processes are largely interdependent. New social interactions are more likely between similar agents, and tie deletion is more likely between dissimilar agents (Kossinets & Watts, 2009; Bener et al., 2016). Additionally, it is worth noting that when computational models of static networks include negative social influence in the models, the system cannot overcome polarization of opinions (Flache et al., 2017; Kozma & Barrat, 2008). Previous work has already stressed the importance of the assumption of the static networks. Using the mechanism of deleting ties to dissimilar agents, Kozma and Barrat (2008) shows that co-evolution can create clusters of like-minded individuals which has a large impact on the convergence time of reaching consensus (Kozma & Barrat, 2008). Although all models do some violence on the world by making assumptions, models should try to include the vital mechanisms of the system they are representing, while simplifying the system enough to understand it (Epstein, 1999; Smaldino, 2016). It is one of this thesis's most central claims that the co-evolution of network and opinion dynamics should be considered a vital mechanism. In order to better explain how opinions are shaped over time, we need to include not only the basic underlying mechanisms of social influence, but also of network dynamics.

4 Network Dynamics

To understand how to include realistic network dynamics, a review of the most used algorithms is provided here. Classic examples of the static, theoretical networks used in most models of opinion dynamics include ring lattices, small-world networks and scale-free networks (Barabási & Bonabeau, 2003; Watts & Strogatz, 1998). Both small-world and scale-free networks are famous for being able to capture essential features of real social networks. Small-world networks are able to capture the general tendency of large average clustering coefficient and small average path length (Watts & Strogatz, 1998). This is achieved by starting with a ring lattice and rewiring a small percentage of the ties randomly. Including a few random rewiring of ties creates highways of information which decreases the average path length considerably (Watts & Strogatz, 1998). Scale-free networks can capture the general tendency of long tails in degree distributions of social networks. The degree distribution of social networks often follows a power law, where a few nodes have an extremely high amount of edges, while most nodes only have a few (Barabási & Bonabeau, 2003). Recent work seems to call into question how universal power laws are in empirical social networks, and finds instead that most networks follow a log-normal distribution (Broido & Clauset, 2019). Note that neither of the most used theoretical networks can generate all the essential features of social networks (Jackson & Rogers, 2004). Small-world networks have

unrealistic degree distributions, and scale-free networks have very low average clustering coefficients.

4.1 Triadic Closure

An alternative to the preferential attachment of scale-free networks and the rewiring of small-world networks is to consider what mechanisms govern how real social networks change over time. Both scale-free and small world networks attempt to answer the question of how networks characteristics could be generated, not why social networks have the characteristics that they do (Jackson & Rogers, 2004). An attempt at answering the why-question is the work by Jackson and Rogers (2004), which emphasizes the role of triadic closure in network generation. Triadic closure refers to making new connections by generating new ties to "friends of friends", i.e. the edges of one's edges (see Figure 3). Reliably, triadic closure is found to be the most important and robust mechanism for how new connections are made in social networks (Asikainen et al., 2020; Bianconi et al., 2014; Kossinets & Watts, 2009; Peixoto, 2022). Empirical studies on dynamical networks find that the probability of creating a new tie is a decreasing function of the distance in the network (Bener et al., 2016; Kossinets & Watts, 2009). The less separated you are from someone, the more likely you are to form a social tie to this person. Specifically, the empirical studies suggest that when you are removed one rather than two degrees of separation from someone, you are 30 times more likely to form a connection. This increase in likelihood only becomes more extreme when distance is increased. When you are 5 degrees of separation away, you are 2.500 times less like to form a connection than you would have been if you were only removed by one degree of separation (Kossinets & Watts, 2009).

Using triadic closure as the generating principle for network formation has shown great promise in explaining some key characteristics of social networks (Ilany & Akçay, 2016). In previous formal models of network formation, this is implemented by making most new connections via triadic closure while letting a few connections be made at random (Ilany & Akçay, 2016; Jackson & Rogers, 2004, 2007). These models can generate the important characteristic findings of high average clustering coefficient, low average path length and log-normal degree distributions (Jackson & Rogers, 2004, 2007).

It is worthwhile to pause to understand why these network metrics are important and why triadic closure might generate patterns akin to those of empirical social networks. Average clustering coefficient is a measure of how connected local communities are in the network (Watts & Strogatz, 1998). Social networks typically exhibit high average clustering coefficients, which is closely linked to the fact that social networks are often made of tight-knitted, well-connected local communities (Peixoto, 2022). Models of triadic closure are likely to exhibit

high average clustering coefficients, as triadic closure will increase the number of local links (Jackson & Rogers, 2007). Average path length refers to the average length of the shortest path between any node in the network (Watts & Strogatz, 1998). Social networks typically have low average path lengths, where no individual is many handshakes away from being connected to another individual (Backstrom, Boldi, Rosa, Ugander, & Vigna, 2012). The network formation models using triadic closure also has low average path length, benefiting from similar mechanisms of randomness as small-world networks to achieve low average path lengths (Jackson & Rogers, 2007; Watts & Strogatz, 1998). A small amount of randomness can create long-range connections, which decreases the average path length substantially (Watts, 1999). The average path length is important for how fast information or contagious diseases can spread in a network (Cowan & Jonard, 2004). Finally, the degree distribution is the distribution of the number of edges of the agents of a network. As mentioned, social networks typically have log-normal or scale-free degree distributions. In networks with such degree distributions, some nodes will have a very large degree, while most nodes will have a relatively low degree (Bianconi et al., 2014). These types of degree distributions are generated via triadic closure. The reason is that triadic closure is a mechanism where the probability of generating a new connection is proportional to how many connections a node already has. High degree nodes have more possibilities for being selected for tie generation simply because they have higher degrees (Jackson & Rogers, 2007). Similarly, low degree nodes will have fewer nodes that could generate new connections with them via triadic closure. This is similar to what is observed via preferential attachment in scale-free networks (Barabási & Bonabeau, 2003).

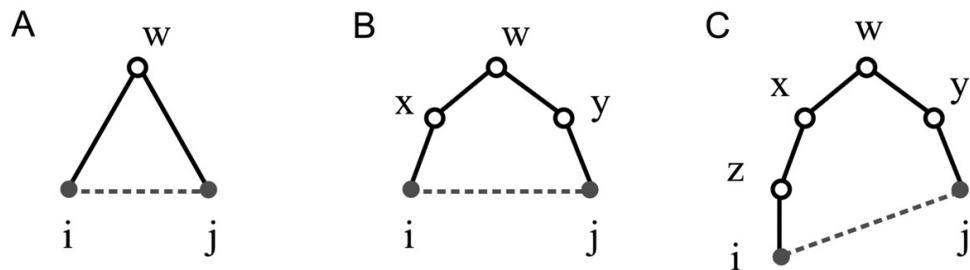


Figure 3: Different types of tie closure. Taken from Kossinets and Watts (2009). Panels show different possible tie-closures. Panel A depicts closing ties with a friend of a friend. This type of tie closure is called triadic closure. Panel B and C shows longer cycles of closure.

5 Co-evolution via tie-deletion

Up until now, opinion and network dynamics have been presented and discussed in isolation. As mentioned, these dynamics are known to be interdependent. It is easy to see how opinions could be influenced by network dynamics. As the social network changes, the available interactions for each agent will change. As one of the primary forces of opinion dynamics is social influence, changes in the social network will have large effects on the opinions of the social agents of the network. This is one side of the coin. As the social network already influences opinion dynamics, any mechanism that changes the ties in the network based on the opinions of the agents will cause opinion and network dynamics to co-evolve. There are good reasons to believe that such effects are numerous (Bener et al., 2016; Kossinets & Watts, 2009; Levin et al., 2021). In this thesis, the co-evolutionary mechanism considered will be the tie-deletion of dissimilar individuals. Tie deletion of dissimilar individuals have been shown empirically to be a robust phenomenon (Kossinets & Watts, 2009; Bener et al., 2016). Moreover, this mechanism has been studied in similar models previously (Santos, Pacheco, & Lenaerts, 2006; Sasahara et al., 2021). By using a similar mechanism, this thesis can better assess the robustness of the current claims in the literature.

5.1 Tie-deletions effect on cooperation

Opinion dynamics is not the only field interested in identifying which circumstances can lead a group of social agents to cooperate. One field especially keen on identifying which dynamics can facilitate the evolution of cooperation is the field of computational biology. Computational biology should be considered an adjacent field to opinion dynamics, as they share much in their interests and in their methods (Dakin & Ryder, 2018; Melamed & Simpson, 2016; Pepper & Smuts, 2002; Santos et al., 2006; Smaldino, 2019).

Intuitively, one might expect that when a network of social actors becomes more connected, levels of cooperation in the system should increase. However, the opposite effect is observed. In their paper, Santos et al. (2006) consider agents playing game-theoretic games, where agents either cooperate or defect. Santos et al. (2006) analyze what happens when agents adjust their ties based on the interactions made with other agents. When cooperators interact with defectors, the social tie between cooperators and defectors has a chance of being rewired to a different agent. This keeps the number of edges constant, but decreases how globally connected the network is (Santos et al., 2006). Their results indicate that cooperation flourishes when the propensity for deleting ties between cooperators and defectors increases (see Figure 4). The interpretation given by Santos et al. (2006) is that removing ties between cooperators and defectors increases positive assortment between coop-

erators and decreases positive assortment between defectors and cooperators. One of the striking attributes of this finding is that without any propensity for tie-removal, cooperation does not evolve for conditions equivalent to the Prisoner's Dilemma (Santos et al., 2006). More generally, positive assortment has been found to be a robust facilitator of cooperation in computational biology (Boyd, Gintis, & Bowles, 2010; Dakin & Ryder, 2018; Melamed & Simpson, 2016; Pepper & Smuts, 2002). As an aside, Santos et al. (2006) also shows that using even relatively simple principles of co-evolution where network dynamics reflect the individual agent's reaction to their social interactions can produce realistic social networks (Santos et al., 2006).

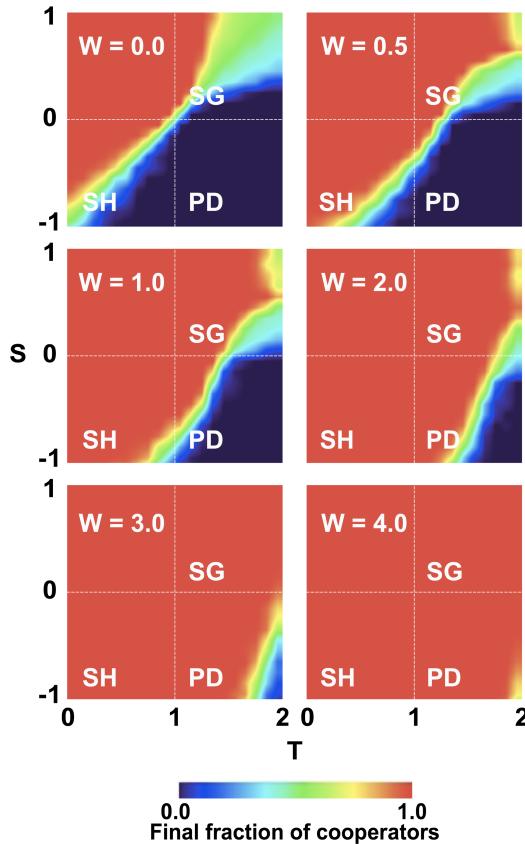


Figure 4: Tie-deletions facilitates cooperation. Taken from Santos et al. (2006). Panels show simulations of different values of W . In their model, W controls the propensity to delete ties between defectors and cooperators. Y-axis shows the disadvantage of being cheated, x-axis shows the pay-off of cheating. Letters on the panel represent well-known game theoretic dilemmas (Prisoners Dilemma, Stag Hunt and Snow Drift Game). Colors show the final fraction of cooperators after simulations are completed. When $W = 4.0$, cooperation effectively wipes out defectors completely.

Although the effect of positive assortment on cooperation seems general in the computational biology literature, models from opinion dynamics point to an opposite result. When combined with social influence, Sasahara et al. (2021) report that tie-deletion between dissimilar agents accelerates polarization by generating echo chambers that stifle cooperation (see Figure 5). This is in line with the more intuitive explanation, where a decrease in

the communication between agents leads to a decrease in cooperation. Similarly, this is also in line with the research on echo chambers, where isolation leads to further polarization (Tsai et al., 2020; Del Vicario et al., 2016).

It is noticeable that the same underlying mechanism gives rise to effects of opposite directions on a system's ability to cooperate. Furthermore, both studies find that tie-deletion is critical to facilitating cooperation, but with different directions of the effect (Santos et al., 2006; Sasahara et al., 2021). It is precisely because of this dispute that further research should be done on how the deletion of negative ties can impact cooperation.

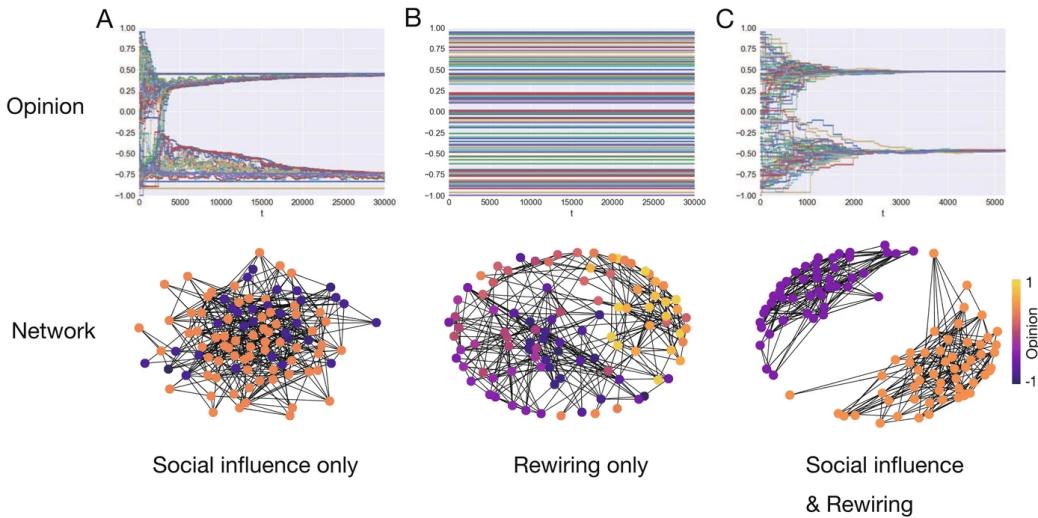


Figure 5: Tie-deletion accelerates polarization. Taken from Sasahara et al. (2021). First row shows the opinions for the simulation, second row shows the network after simulation. Columns show simulations of different conditions. Column A shows the results from only including social influence. Column B shows the results of only rewiring based on opinions. Column C shows the results of including both social influence and rewiring. Of special interest is the last column, which shows that tie-deletion can accelerate echo chamber formation, as the opinions stabilize faster than when only social influence is included.

6 Model Description

The following section describes the agent-based model of opinion and network dynamics based on the key mechanisms of homophily, social influence and triadic closure. First, a conceptual overview of the model is provided, followed by a description of the key assumptions of the model. Thereafter, the specific steps which govern the behavior of the agents of the model at each time-step is explained in detail.

6.1 Conceptual overview

The agent-based model of this thesis represents a system where social agents interact with their peers and discuss their opinions with each other. The system of interest is a social system where agents change not only their opinions, but also their social ties. To represent these social ties, the agents are situated in an undirected and unweighted network, where nodes of the network represent social agents and edges represent social ties between agents. The fact that the network is undirected means that all relations between agents are reciprocal. This way of representing connections between social agents results in a system where you cannot be friends with someone if they are not friends with you. Similarly, the fact that the network is unweighted simply means that no connections are stronger or weaker than others.

6.2 Assumptions

The model makes several assumptions to simplify the process of opinion and network dynamics. It is assumed that the opinion of an agent is shaped only by her initial opinion and the influence of her peers. The model does not assume that the agents of the system have any tasks to perform or strategies to follow. Agents are not trying to find any "true opinion" or act in a way to hide their own opinion from their peers. Instead, the agents are assumed to share their opinions truthfully to other agents, and to be able to perceive the exact opinion of their connections. In line with previous models of opinion dynamics, the model assumes positive social influence between agents of similar opinions. In other words, agents of similar opinions will reach a compromise which pulls their opinions closer together. Similarly, the model assumes negative or no social influence between dissimilar agents. Agents will either not be influenced at all, or will further distance themselves from agents of dissimilar opinions from their own. Regarding the social network of the model, it is assumed that agents will find new connections primarily through their already existing connections via the mechanism of triadic closure. Finally, it is assumed that agents will tend to delete ties to dissimilar agents. All these assumptions can vary in the strength of the proposed effect. Take for instance the effect of positive social influence. A strong force of positive social influence would result in agents reaching a perfect compromise, while a weaker force would result in agents that barely influence each other. As the results of the model depend on the strength of these assumptions, they are controlled for explicitly by the parameters of the model.

6.3 The stages of the model

The agent-based model is divided into three distinct and sequential stages, which are executed at every time-step. In essence, every time-step consists of an agent first creating new social ties, sharing their opinions with their existing connections and then deleting dissimilar ties. These stages are referred to as the network dynamic stage, the opinion dynamic stage and the co-evolutionary stage respectively:

- The network dynamic stage specifies how new edges are created in the social network of the model. The behavior of this stage of the model is controlled by one parameter, R . R specifies the probability of generating random ties rather than ties via triadic closure.
- The opinion dynamic stage describes how interactions between agents change their opinions via social influence. The opinion dynamics are controlled by three parameters, T , α and β : T specifies the threshold for what constitutes a similar agent, α specifies the power of positive social influence and β specifies the power of negative social influence.
- Finally, the co-evolutionary stage specifies the tendency for agents to delete connections to dissimilar agents. This is controlled by the parameter $P(D)$, which describes the probability of deleting ties to dissimilar agents.

The specifics of each stage of the process is introduced in detail in the section called *Dynamics*. While this thesis mainly focuses on one agent-based model, it develops two models. This is done to investigate how the co-evolution of opinion and network dynamics differs from only considering network dynamics. One of the models considered is an agent-based model which only contains the network dynamic stage as described above. This model serves as the baseline for the full co-evolutionary model, which includes all three stages described above. The comparison between these two models will be a central part of the discussion on the integration of empirical data. I will refer to the model that only includes the network dynamic stage as the Network Formation Model and the model with all three stages as the Co-evolutionary Model (see Figure 6).

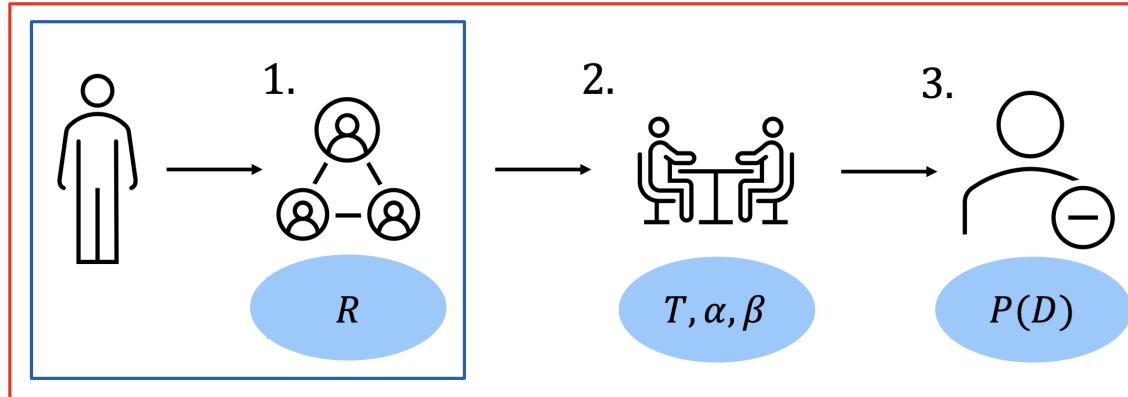


Figure 6: Flowchart of a single time-step. Numbers indicate different parts of the process of a single time-step. Squares of different colors indicate time-steps of different models. The blue square depicts a single time-step for the Network Formation Model and the red square indicates the Co-evolutionary Model. Cyan bubbles show what model parameters are relevant to the particular model stage. A random agent, A_t , is sampled from the network. Next, the network dynamic stage begins (1) where a new connection is made from A_t to another agent. Whether the connection is made randomly or via triadic closure is controlled by R . This is the only stage included in the Network Formation Model. After the network dynamics stage, the Co-evolutionary Model starts the opinion dynamic stage (2). The sampled agent, A_t , pushes and pulls on opinions of neighboring agents. This is controlled by a threshold for similarity, T , and forces of positive and negative social influence, α and β . After this process, the co-evolutionary stage begins (3). With a probability of $P(D)$, ties to agents with an opinion further than T apart from A_t 's opinion are deleted. This concludes one time-step of the Co-evolutionary Model.

6.4 Initialization

The model has a set of requirements that must be initialized before it can function as intended. The first of which is that for social agents to be social, they need a social scene to act on. Therefore, agents are initially situated in a random Watts-Strogatz small-world graph (Watts & Strogatz, 1998). For all models of this thesis, all initial small-world graphs are generated with the parameter $p = 0.5$, while the parameters N and k are varied throughout the thesis (Watts & Strogatz, 1998). The choice of initial network was made as they are a typical choice of networks in the literature (Turner & Smaldino, 2018; Flache et al., 2017). The second initial requirement of the model is that the agents must have some opinions to share with their peers. To model opinions, each agent is given a continuous number between -1 and 1, which is taken to represent her initial opinion, O_I . By having opinions represented continuously, opinions can be anything between extremely pro or extremely against some notion. In line with previous research, the initial opinions of agents are assumed to be diverse (Galesic et al., 2021; Flache et al., 2017; Flache, 2018). To implement diversity of initial opinions, all

values of O_I are initialized by drawing from a uniform distribution between -1 and 1:

$$O_I \sim U(-1, 1)$$

As a consequence, no agent will be initialized with an opinion lower than -1 or higher than 1. Additionally, the expected mean of the distribution of O_I is 0.

6.5 Dynamics

To investigate how network and opinion dynamics influence each other over time, agents iteratively interact with each other over the course of a number of discrete time-steps. In the real world, social agents are constantly interacting, with multiple interactions happening at the same time. This causes problems in simulation models, as synchronous updating of opinions have been reported to cause dramatic instability in the models (Flache et al., 2017; Sasahara et al., 2021; Galesic et al., 2021). For this reason, this model follows previous conventions and updates opinions asynchronously. To do so, a random agent is sampled from the network to be the agent on turn every time-step. Formally, let t specify the time-step of the agent-based model, A_t be the agent on turn at time-step t and let S_t be the network at time-step t .

Because the model acts asynchronously, the maximum number of time-steps, t_{max} , should be proportional to the number of agents in the system, N . This is standardized across different network sizes by specifying that all agents are expected to be the agent on turn 20 times during the course of one simulation. For all simulations considered in this thesis, the model terminates when $t = t_{max}$, where t_{max} is defined as

$$t_{max} = 20 \cdot N$$

6.5.1 Network dynamic stage

The first stage of the model represents how social agents find new connections in social networks. These new connections are primarily generated via triadic closure, while a subset of new connections will be generated randomly. The balance between triadic closure and randomness is controlled by the parameter R . With a probability of R , A_t connects to a random agent that A_t is not currently connected to. With a probability of $1 - R$, A_t connects to one of its edge's edges that it is not currently connected to. The lower the value of R , the more

new ties are generated via triadic closure.

Some intricacies concerning the network dynamic stage are worth noting, as later stages of the model can delete social ties. This is consequential as the number of ties in the network is likely to be an important factor and therefore needs to be controlled for. To keep the number of edges in the network approximately constant, new edges are either created or rewired to account for deleted edges. This is done by ensuring that when the number of edges in the network is lower than the initial number of edges of the network, new edges are created and not rewired. When the number of edges in the network is the same number as it had initially, new edges are rewired instead. As a consequence, the network can have fewer edges than it had initially but never more. Formally, this is achieved by letting E_1 be the number of edges of N_1 and letting E_t be the number of edges of S_t . If $E_t < E_1$, A_t will not rewire one of its existing edges, but instead create a new edge. If $E_t \geq E_1$, A_t will rewire one of its existing connections to the new agents.

Rewiring of edges or deletion of ties can cause the network to separate completely into distinct components where no links exist between members of different components. Although complete segregation of networks likely occurs in the real world (see Figure 1), having multiple components complicates network analysis extensively. For instance, the average path length of the social network is an important measure of this study, but it is only rigorously defined for networks of one component. When a network only has one component, there exists a path between any two nodes of the network. To simplify the network analysis and allow for the calculation of the average path length, the model is restricted to only have one component. Formally, this is done by letting C define the newly connected agent that A_t connected to via triadic closure or randomly at the start of her turn. Next, the degree of A_t and C is calculated. If the degree of either A_t or C is 0, these agents constitute a second component on their own. In this case, a new edge is created randomly from the agent that has a degree of 0 to a new node from S_t . If both A_t and C have degrees larger than 0, and S_t has more than one component, a new edge is created to merge the two components together. Specifically, we restore the edge that A_t rewired from C , while keeping the edge from A_t to C . I will refer to the process of ensuring only one connected component simply as component insurance.

6.5.2 Opinion dynamic stage

After finding a new connection, the agents interact with each other. This is meant to represent discussions or other social interactions between agents. Social interactions between similar individuals will lead to compromises. Up to a point, the more you disagree with a person, the more malleable your opinion is. If you reach a

compromise with a person with a very different opinion than yours, you will change your opinion much more than if you reach a compromise with someone that you almost agree with. To model this kind of behavior, the agent on turn, A_t , interacts with all her connected agents iteratively and changes her opinion and theirs. As the ordering of interactions between the edges of A_t could be important for the resulting change in opinions, the order of interactions is randomized at every time-step. Whether other social agents are similar enough to reach a compromise is decided by a threshold value, T . When the distance between the opinions is lower than T , the social influence will pull these opinions closer together. If the distance is larger than T , social influence will push the opinions further apart. Representing this formally, let B denote an edge of A_t and let $O(\cdot)$ define a function with agents as inputs and their opinions as outputs. When $T \geq |O(A_t) - O(B)|$, the interaction between the two agents will be positive and as a result the opinions of the two agents will be pulled closer together. The force with which they are pulled is the positive social influence, V_p , which is defined as a fraction of the distance between opinions, where the fraction is controlled by the parameter α :

$$V_p = (|O(A_t) - O(B)|) \cdot \frac{\alpha}{2}$$

Let $O_{max} = \max(O(A_t), O(B))$ and $O_{min} = \min(O(A_t), O(B))$ and update the values of opinions by:

$$O'_{max} = O_{max} - V_p$$

$$O'_{min} = O_{min} + V_p$$

where O'_{max} and O'_{min} are the updated values of O_{max} and O_{min} respectively after interacting with each other. When $T < |O(A_t) - O(B)|$, the interaction will be negative. The opinions of the two agents will be pushed further apart by the power of negative social influence, V_n , via a principle similar to positive social influence. With negative social influence, how much of the fraction they are pushed apart is controlled by the parameter β :

$$V_n = (|O(A_t) - O(B)|) \cdot \frac{\beta}{2}$$

Using the same definitions as above, the opinion of each agent is updated:

$$O'_{max} = O_{max} + V_n$$

$$O'_{min} = O_{min} - V_n$$

As the model represents opinions on the spectrum between -1 and 1, the opinions after interacting are also restricted to fall within this spectrum. If an agent's opinion becomes larger than 1 after interacting with other agents, their updated opinion is replaced with the number 1. Similarly, if their opinion becomes lower than -1, their opinion is replaced with the number -1. Before moving on to the co-evolutionary stage of the model, there are some aspects of this particular operationalization of social influence that are important to notice. The first of which is that the difference between how the power of positive and negative social influence is exerted is minimal. The main difference is a sign difference in how they influence updating. Positive social influence works by adding V_p to the opinion closest to -1 and subtracting V_p from the opinion closest to 1. With $\alpha = 1$, this would result in the agents reaching a perfect compromise, where their opinions would be identical after interacting, as each agent is moved closer by half of the distance between them. Negative social influence is operationalized as an opposite to positive social influence. Negative social influence works by subtracting V_n from the opinion closest to -1 and adding V_n to the opinion closest to 1. The second important aspect to notice is that social influence is operationalized as being symmetrical, in the sense that the opinion of agents are moved by the same amount.

6.5.3 Co-evolutionary stage

After interacting with her connected agents, A_t might delete the connection to some of her connections that are dissimilar to her. For the rest of the paper, this process is referred to as negative tie-deletion. To link this mechanism to the opinion dynamic stage, negative tie-deletion relies on the threshold, T , to determine which interactions positive or negative. Formally, if $T < |O(A_t) - O(B)|$, there is a chance that their tie is deleted after interacting. The probability of negative tie-deletion between dissimilar agents in the model is described by the parameter $P(D)$. When $P(D) = 1$, there is a 100% chance that dissimilar ties will be deleted, while with a $P(D) = 0.5$, there is 50% chance that dissimilar ties will be deleted. When ties are deleted, there is a possibility that the network will separate into multiple components. The process of component insurance is therefore performed exactly as described in the *Network dynamic stage* section. When A_t has removed her negative ties at a probability of $P(D)$, the time-step concludes. The simulation ends when $t = t_{max}$.

6.6 Data availability

To allow for full transparency of the model, its implementation and its conclusion, the entire code is available for anyone to reproduce. The code is fully documented and specifically designed to be easily readable. The

hope is that other researchers or interested parties will be able to build, expand, validate or refute the proposed framework. The code for generating the model, the analysis, the plots and the paper is all available on [Github](#).

Part II

Integrating empirical data

Arguably, the most pressing issue facing models of opinion dynamics is that they rarely consider how the results of their computational models compare with the real world (Flache, 2018). Most models of opinion dynamics are low-dimensional (Bener et al., 2016), in that they don't seek to explain exhaustively how opinions are formed, but rather to point out interesting interactions between key variables which govern opinion dynamics. This makes the inclusion of empirical data especially tricky for opinion dynamical models. Compounding this problem is that the time-scales considered by these models are massive. Longitudinal data which is needed to validate these models are very hard to come by (Mäs, 2019; Kossinets & Watts, 2009). Even if longitudinal data are acquired, it is often not clear what a single time-step of the model is meant to represent in the real world (Mäs, 2019). As a consequence, opinion dynamical models can rarely make specific predictions, which makes validation via empirical data even harder. In conclusion, incorporating empirical data in these models is a surprisingly daunting task - but not an impossible one. As the model considered in this thesis is a co-evolutionary model, it is possible to shift the focus of empirical data inclusion from opinions to networks. Instead of investigating how well the model generates real world opinions, it is possible to investigate how the generated networks of the model compare to real empirical social networks. This is a common practice in statistics, where a model is fitted to data, whereby it learns the values of the model parameters. It does so by minimizing some objective function of difference between the generated predictions and the observed values, typically a distance metric such as the mean squared error (Akiba, Sano, Yanase, Ohta, & Koyama, 2019).

A similar problem of finding the right parameter values to generate the best models is finding hyperparameters values for deep learning neural networks (Bergstra, Bardenet, Bengio, & Kégl, 2011). Getting the right combination of hyperparameters for these incredibly large models is often the difference between success and failure (Akiba et al., 2019; Bergstra et al., 2011). Finding the hyperparameters which minimize some objective function is referred to as hyperparameter optimization. Hyperparameter optimization is an active field of research, and it has already produced several effective algorithms to finding the best combination of parameters such as the tree-structured Parzen estimator (Akiba et al., 2019; Bergstra et al., 2011; Hutter et al., 2014).

Given some objective function, the tree-structured Parzen estimator can learn to navigate the parameter space essentially via intelligent trial-and-error. By sampling different parts of the parameter space, the estimator can learn what combination of parameters result in minimizing the objective function of interest. Only recently have agent-based modelers used hyperparameter optimization to integrate empirical data into their models (Kerr et al., 2021; Krivorotko, Sosnovskaia, Vashchenko, Kerr, & Lesnic, 2022). While these validation approaches are seminal for using hyperparameter optimization on agent-based models, none of the recent proposals analyze how important each model parameter is for the goodness of fit. This is despite the fact that effective methods exists for computing parameter importance (Hutter et al., 2014). By fitting a random forest classifier to the different samples of the parameter space that the optimization algorithm has tested, one can perform a fANOVA of the model parameters. The result of the fANOVA is a measure of the fraction of the variance that is explained by each model parameter (Hutter et al., 2014). This makes the importance of each parameter easy to interpret, enabling researchers to investigate which model parameters were the most critical for explaining the variance in the data.

This thesis uses hyperparameter optimization to integrate empirical data. Similar to traditional approaches in data science, the aim is to estimate how well generated networks match real empirical networks. The interesting results in this regard are whether the agent-based model can generate realistic networks, and if so which model parameters that lead to realistic networks. This is analogous to estimation of parameter values in statistics. Beyond knowing whether the model can generate the observed structure of real-world networks, this thesis is specifically interested in how critical co-evolution is to the explanatory power of these models. This is where having both the Network Formation model and the Co-evolutionary model comes into play. By comparing the results obtained from both these models, it is possible to identify the circumstances where the performance of the two models differ significantly. This will help answer the question of how vital co-evolution is as an organizing principle of social networks. In addition to the Network Formation model and Co-evolutionary model, the two most used network generating algorithms, the small-world (Watts & Strogatz, 1998) and the scale-free network (Barabási & Bonabeau, 2003) will be used as benchmarks. This is analogous to model comparison in statistics. Although these steps are analogous to model fitting and model comparison, there are important differences worth mentioning. Especially important is the fact that model comparison normally penalizes the inclusion of additional parameters in the model (Vrieze, 2012). In other words, simple explanations are preferred when the performance of these explanations are comparable (Emiliano, Vivanco, & De Menezes, 2014; Vrieze, 2012).

7 Data

To compare the generated results with real empirical networks, 7 empirical social networks are examined. These networks differ in many aspects, but of main interest to this thesis is that they differ in size and how opinionated they are expected to be. The 7 networks consist of a social networks of dolphins (Lusseau et al., 2003), the Karate Club Network (Zachary, 1977), a citation network of the field of network science (Newman, 2006), a co-purchase network of political books on Amazon (Shi, Shi, Dokshin, Evans, & Macy, 2017), a network of political blogs (Adamic & Glance, 2005), a network of politicians based on shared likes on Facebook and a similar network of TV shows based on shared likes on Facebook (Rozemberczki, Davies, Sarkar, & Sutton, 2019). For all networks, only the main component of the network was considered. This was done as average path length is not well-defined for graphs with multiple components. As the model assumes undirected networks, the network of political blogs was transformed from a directed to an undirected network. This was done by making any directed edge into a reciprocal edge between the two agents.

8 Optimization methods

Hyperparameter optimization was performed using the Optuna framework (Akiba et al., 2019) to gauge what model parameters most closely resemble empirical social networks. Optuna was chosen as it showed the best optimization results for previous studies using hyperparameter optimization for agent-based models (Kerr et al., 2021). Specifically, the use of a tree-structured Parzen estimator was used to optimize as previous evaluations of hyperparameter estimators suggest that it is the most effective estimator algorithm (Akiba et al., 2019; Bergstra et al., 2011; Hutter et al., 2014).

As discussed, the tree-structured Parzen estimator works by navigating the parameter space in search for combinations of parameters which minimizes an objective function (Akiba et al., 2019). It is therefore essential for the use of hyperparameter optimization to establish a proper objective function. For this thesis, such an objective function is a distance metric between networks. Although many features are important for characterizing a network, this thesis follows previous work which identifies the average path length, the average clustering coefficient and the degree distribution as the most defining aspects of a network (Jackson & Rogers, 2004, 2007). A meaningful distance metric between networks is therefore the mean difference of these three parameters between the generated and the empirical networks. Before we can simply calculate the mean difference, there are

some specifics regarding the degree distribution and the average path length needs to be clarified. In order to calculate the difference between two degree distributions, several metrics are possible to use. One such metric is the Jensen Shannon Divergence which is a distance metric between probability distributions (Fuglede & Topsoe, 2004). The Jensen Shannon Divergence between two degree distributions gives a measure of how distant these two distributions are from each other, and will be the metric used here.

For large networks, calculating the exact average path length is an expensive computational operation (Matsumura, Iwasaki, & Shudo, 2018). Instead of calculating the precise average path length, one can approximate it by taking 1000 samples of random pairs of nodes and calculating the average path length of those 1000 nodes. Despite not being the exact average path length of the network, this method has been shown to generally approximates the average path length well (Matsumura et al., 2018) and will therefore be used to approximate the average path length in this thesis. In addition to the approximation, the average path length must be normalized as well to ensure that all measures are weighted equally. Both the average clustering coefficient and the Jensen Shannon Divergence can only take on values between 0 and 1 by definition. The average path length on the other hand is not restricted in the same way. If it is not normalized, it will therefore be weighted much higher than the average clustering coefficient and the Jensen Shannon Divergence when the distance metric is just a simple average. We therefore normalize the average path length by

$$APL^* = \frac{|APL(G) - APL(A)|}{APL(A) + 2}$$

where APL^* is the normalized approximated average path length, $APL(G)$ is the approximated average path length of the generated network and $APL(A)$ is the approximated average path length of the actual network. It is important to note that this is not a true normalization, as APL^* can in principle still be above 1, although this is highly unlikely for the comparisons made in this thesis. This can happen if the average path length of the generated network, $APL(G)$ is much larger than that of the actual network, $APL(A)$. This is also the explanation for why there is an addition of 2 in the denominator. Although the specific value of 2 is somewhat arbitrary, the choice was made to be more certain that APL^* would still be below 1. With this in place, the objective function considered which quantifies the difference between the actual network, A , and the generated network, G , is

$$O(A, G) = \frac{1}{3} \cdot \left(JSD(D(A), D(G)) + |C(A) - C(G)| + APL^* \right)$$

where $O(A, G)$ is the difference between the two networks, JSD is the Jensen Shannon Divergence, $D(A)$ is the degree distribution of A , $D(G)$ is the degree distribution of G , $\bar{C}(A)$ is the average clustering coefficient

of A and $\bar{C}(G)$ is the average clustering coefficient of G . The sum of the differences is multiplied by $\frac{1}{3}$ to aid interpretation. $O(A, G)$ is the average of differences between two networks.

8.1 Specifying model parameters

The empirical networks considered vary greatly in size, which must be controlled for in fitting the agent-based model to the data. The primary concern is to initialize the network with an appropriate amount of edges. To this end, all models were fitted by first calculating the ratio between edges, E , and vertices, V . This ratio, K , is then rounded to the nearest integer:

$$K = \lfloor (2 \cdot \frac{E}{V}) \rfloor$$

The value K is then used as the K -parameter in the initially generated small-world network to ensure that the generated network has approximately the same number of edges per node as the target network.

To optimize the different models, the possible parameter values for each model must be specified (Akiba et al., 2019). For the small-world network, the rewiring probability, P , is specified as a free parameter between 0 and 1. For the scale-free network, the number of generated connections made from new agents, k , is specified as an integer value between 1 and 22. For both the Network Formation model and the Co-evolutionary model, the probability of generating new ties at random, R , is a free parameter between 0 and 1. For the Co-evolutionary model, the threshold parameter, T , is restricted to be between 0.1 and 1.3. The lower limit is chosen as values of T close to 0 results in agents that don't agree with anyone else but themselves. The upper limit was chosen to ensure that the behavior of the Co-evolutionary model differs significantly from the Network Formation model. The force of positive social influence, α , is restricted to be between 0.05 and 0.5. This is justified by values close to 0 being akin to no positive social influence, and values above 0.5 assumes social influence to be an unrealistically powerful force. Similarly, the force governing negative social influence, β , is restricted between values of 0 and 0.5. Letting β have values close to 0 reflects that negative social influence is still a less established fact (Takács et al., 2014; Turner & Smaldino, 2018). Finally, we force the model to exhibit co-evolution by restricting the probability of tie-deletion to be between 0.1 and 1.

With these restrictions on the parameters, hyperparameter optimization is performed with four models for all 7 empirical networks. This gives the best model parameters for generating every network. Optimization was done with 500 iterations per empirical network, which in all cases provided stable results (see Appendix A.1 to A.7).

As all the models considered are stochastic models, the best parameters found from optimization are used to generate 10 different networks using different random seeds. These 10 simulations act as a robustness check, and ensures that models reliably generate the patterns of the empirical networks. To calculate the importance of each parameter, a fANOVA was performed on the Co-evolutionary Model for each network.

9 Optimization results

The results show small-world networks can capture the average clustering coefficient and average path length of empirical networks very well, but fail to capture the degree distribution of empirical networks. Scale-free networks are better at approximating the degree distribution, but fail on the other two parameters(see Figure 8 & Table 2). Models based on triadic closure are generally better approximations of empirical networks than small-world and scale-free networks (see Figure 7 & Table 2). The Network Formation Model is on par with the Co-evolutionary Model for only the two smallest networks considered, the Dolphins network and the Karate Club Network (see Figure 7 & Table 1). The Co-evolutionary Model drastically outperforms the Network Formation model on larger networks, especially when large networks are expected to be highly opinionated. This is the case for the two networks Politicians and Political Blogs (see Figure 7 & Table 1). For these two networks, the Co-evolutionary model outperforms the Network Formation Model primarily due to lower Jensen-Shannon Divergence as well as lower difference in average clustering coefficient between simulated and empirical networks (see Appendix A.8, A.10).

The results of the fANOVA indicates that the R -parameter was the most important parameter for all networks except for the Political Blogs network. The Political Blogs network has the T -parameter as the most important instead. Three of the four networks with the highest importance of the parameter of T are the political networks (Political Books, Political Blogs, Politicians) (see Figure 9). The best parameters of the Co-evolutionary Model have low values of R for all networks except the Karate Club. Similarly, we see that the best parameter values for most networks have low values of β and relatively low values of T . This is especially the case for the Politicians and Political Blogs networks, which both have the best parameter values of $T \approx 0.1$ and $\beta \approx 0.01$ (see Table 3).

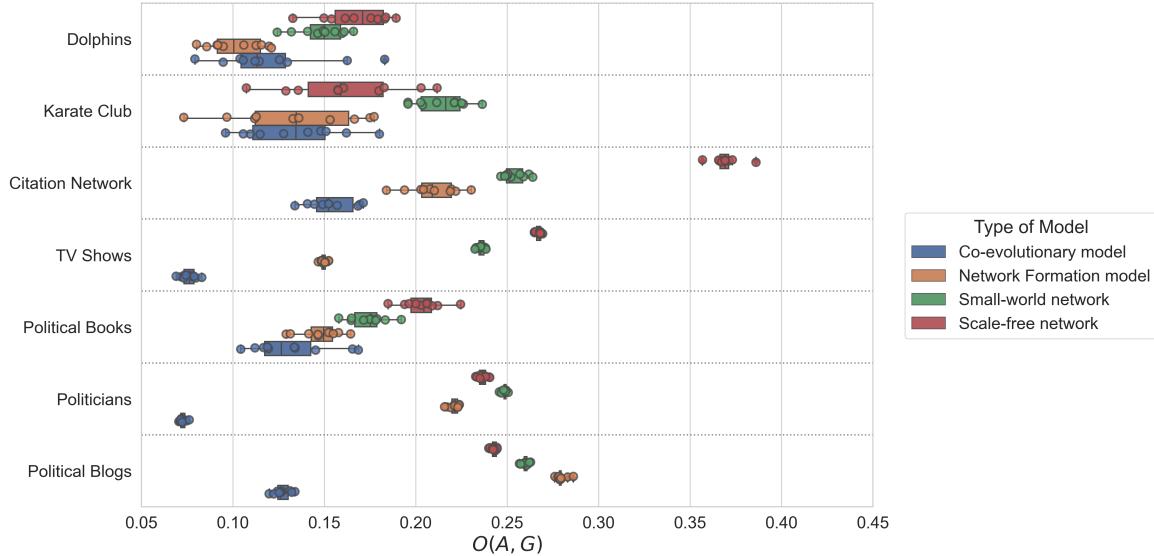


Figure 7: Difference in model performance. X-axis shows the mean difference between the generated network and actual networks, $O(A, G)$. Y-axis shows the different empirical networks considered. Colors show different algorithms for generating networks. Dots show the result from individual simulations, while boxplots show the median and quartiles of the distribution of values.

Table of network characteristics						
Network	N	E	D_μ	D_σ	\bar{C}	APL
Dolphins	62	159	5.129	2.932	0.259	3.257
Karate Club	34	78	4.588	3.820	0.571	2.326
Citation Network	379	914	4.823	3.927	0.741	6.052
TV Shows	3892	17262	8.871	12.557	0.374	6.324
Political Books	105	441	8.400	5.449	0.488	3.015
Politicians	5908	41729	14.126	20.096	0.385	4.628
Political Blogs	1222	16717	27.360	38.402	0.320	2.737

Table 1: Table of network characteristics. N is the number of nodes, E is the number of edges, D_μ is the mean degree, D_σ is the standard deviation of the degree distribution, \bar{C} is the average clustering coefficient, and APL is the average path length.

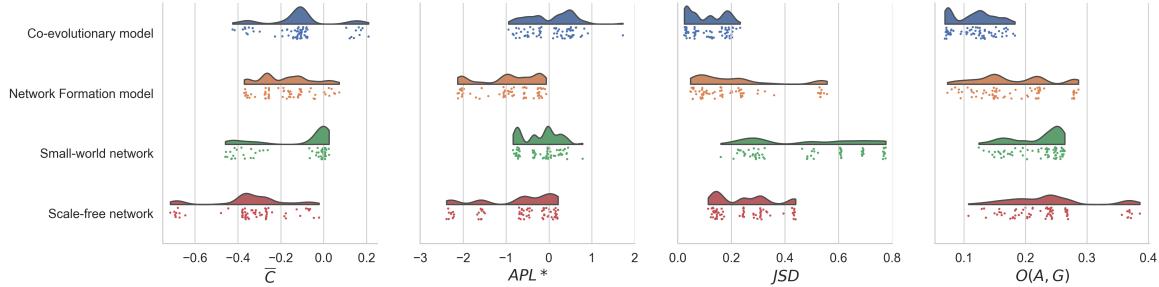


Figure 8: Rain cloud plot of the overall difference in model performance. The plot is divided into four panels. Each panel shows the performance on different metrics. Each metric is shown on the x-axis. The first two panels show the difference in regard to the average clustering coefficient, \bar{C} , and the average path length, APL^* , respectively. Both were calculated by subtracting the value of the empirical network from the generated networks. Third panel shows the Jensen Shannon Divergence, JSD , of the degree distribution from the generated and the empirical networks. Finally, the fourth panel shows the value of the objective distance function between the empirical and the generated network, $O(A, G)$. Performance is indicated by the distance to zero for all metrics. The y-axis and the colors show different network generating algorithms. Dots show the result from individual simulations with the probability density function drawn above the dots.

Summarized performance for different network formation algorithms								
Model	\bar{C}_{MED}	\bar{C}_{IQR}	APL_{MED}	APL_{IQR}	JSD_{MED}	JSD_{IQR}	O_{MED}	O_{IQR}
Scale-free network	-0.351	0.110	-0.480	1.531	0.245	0.163	0.236	0.079
Small-world network	0.005	0.298	-0.074	0.785	0.499	0.397	0.233	0.069
Network Formation model	-0.179	0.163	-0.837	1.210	0.161	0.148	0.158	0.086
Co-evolutionary model	-0.113	0.064	0.195	0.808	0.108	0.143	0.119	0.065

Table 2: Table for the performance of different network formation algorithms. Rows show different network generating models. Columns show the median (*MED*) and inter-quartile range (*IQR*) of the difference between the empirical and the generated networks. For the average clustering coefficient (\bar{C}) and the average path length (APL), difference was calculated by subtracting the value of the empirical network from the value of the generated network. The Jensen Shannon Divergence (JSD) was kept as is, as it already describes the difference between networks. Bold text indicates the value closest to zero in each column.

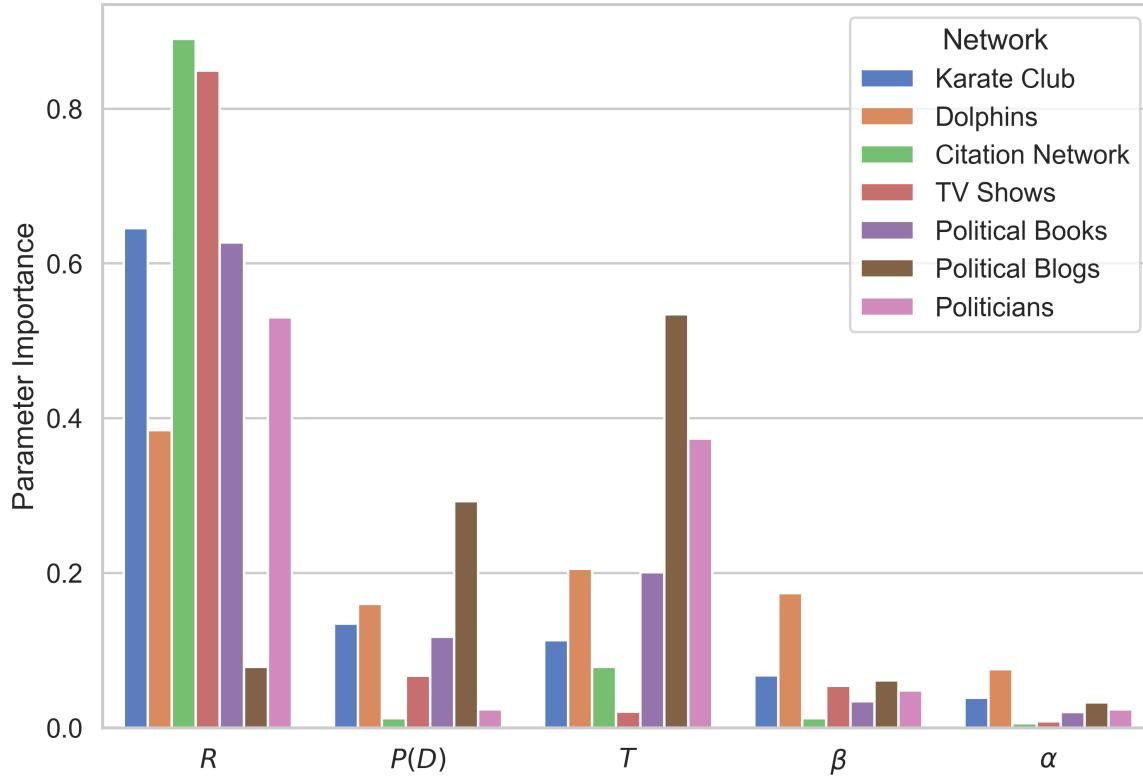


Figure 9: Hyperparameter importance. X-axis shows the different parameters of the Co-evolutionary model. Y-axis shows the fraction of the total variance explained by each parameter i.e. parameter importance. Colors indicate different networks. Parameter importance is calculated using an fANOVA (Hutter et al., 2014).

Table of best parameter values					
Network	R	T	α	β	P(D)
Dolphins	0.003	0.531	0.202	0.095	0.111
Karate Club	0.870	0.137	0.282	0.090	0.911
Citation Network	0.060	0.655	0.389	0.493	0.150
TV Show	0.038	0.157	0.349	0.093	0.751
Political Books	0.004	0.222	0.075	0.144	0.508
Politicians	0.014	0.101	0.050	0.014	0.229
Political Blogs	0.109	0.100	0.246	0.007	0.249

Table 3: Table of best parameter values for the Co-evolutionary model. Values correspond to the parameter combination, which minimized the objective function the most.

10 Discussion of optimization

The results from the hyperparameter optimization reiterates that triadic closure works well as a basic generating principle of social networks (Jackson & Rogers, 2007; Kossinets & Watts, 2009; Bianconi et al., 2014). The results reaffirm that small-world networks can generate realistic average clustering coefficients and average path lengths but fail dramatically concerning generating realistic degree distributions (see Figure 8)(Jackson & Rogers, 2007). Triadic closure is shown to be able to generate the patterns in empirical social networks better than the most used theoretical models for social networks (see Table 2). One of the main differences in performance stem from triadic closure reliably producing more realistic degree distributions, while still having large clustering coefficients. Scale-free networks cannot generate the high average clustering coefficients found in social networks. Interestingly, triadic closure creates more realistic degree distributions than scale-free networks. The fact that triadic closure outperforms both small-world and scale-free networks is notable in its own right, as these theoretical networks serve as the basis for many computational models (Flache et al., 2017; Turner & Smaldino, 2018).

For the remainder of the discussion, the focus will be on what separates the performance of the two triadic closure models, the Network Formation model and the Co-evolutionary model. The Network Formation model is on par with the Co-evolutionary model when the size of the network is small. This is the case in the Dolphins network and the Karate Club network (see Figure 7). The Karate Club network is especially interesting, since this network is notoriously a polarized network (Zachary, 1977). The fact that the two models show similar performance on this network seems to suggest that the Karate Club network is not large enough for the two models to perform significantly different. This is argued by the fact that the size of the network seems to be an important attribute difference between the networks where the models achieve similar performance and where the Co-evolutionary model performs better than the Network Formation model (see Figure 7).

10.1 When co-evolution is a better explanation

As the Co-evolutionary model is a more complicated model, it should be expected that it can fit more complicated patterns. The important question to answer is what ability the model has gained at the cost of higher complexity. The most plausible answer is that it enables the Co-evolutionary model to systematically generate networks with multiple neighborhoods. The notion of neighborhoods here refers to a set of nodes that are well-connected locally but not globally. The Co-evolutionary model can generate multiple neighborhoods in the network because of the co-evolution between opinion and network dynamics. When there is a propensity

to delete ties to dissimilar neighbors ($P(D) \gg 0$), neighborhoods of like-minded individuals will be created. This is especially the case when new ties are primarily generated via triadic closure ($R \ll 0.5$), as this will make new connections primarily between agents of the same neighborhood. If these two conditions are met, the Co-evolutionary model can in effect manipulate how many neighborhoods the network contains by adjusting the parameter T . The lower T is, the less open-minded agents are, as they will have stricter criteria for who to maintain ties with. With lower values of T , the resulting network will therefore contain more neighborhoods. Notice that for all large networks ($N > 1000$), the best parameters of the Co-evolutionary model have $P(D) \gg 0$, $R \ll 0.5$ and $T < 0.2$.

10.2 The impact of neighborhoods

The ability to systematically create networks containing multiple neighborhoods is important because it is highly unlikely that larger networks only contain one neighborhood. All members of large networks are unlikely to be equally well-connected to each other. It is far more likely that large networks contains several neighborhoods or subgroups, where social agents mostly interact with a small subset of all the agents of the network. This explanation can account for the large difference between the Co-evolutionary model and the Network Formation model, as this ability to generate neighborhoods is something that the Network Formation model doesn't have. Such an ability will obviously have a large effect on the performance of the model, as having multiple neighborhoods in the network will have a strong influence on all the considered network metrics.

To see the impact of neighborhoods on the considered network metrics, let's start by considering the impact of neighborhoods on the average path length. Having more neighborhoods will increase the average path length of the network. Neighborhoods are very well-connected locally, but not very well-connected globally. Few information highways will exist between neighborhoods, resulting in larger average path length than if there had not been multiple neighborhoods in the network. This is also evident from the results. On all networks considered, the average path length of the Co-evolutionary model is consistently higher than the average path length of the Network Formation model (see Appendix A.9).

Having multiple neighborhoods will have a large effect on the degree distribution of a network. When networks consist of multiple neighborhoods, the network will tend to have fewer very low degree nodes and fewer very high degree nodes. In other words, the network will have less extreme disparity in its degree distribution. Extremely high degree nodes become less plausible, simply due to the fact, that an agent is very unlikely to be part of many neighborhoods at the same time. It will therefore sever connections to more agents than had

there not been any neighborhoods. Having multiple neighborhoods will tend to reduce the number of extremely low degree nodes, because these neighborhoods are locally very well-connected. Low degree nodes will have more connections within these neighborhoods than they would have if there had not been any neighborhoods at all. These two effects combined will tend to make the degree distribution of the network less scale-free and more log-normal, which is more in-line with the degree distribution of most social networks (Broido & Clauset, 2019). It should be expected that the separation between groups becomes more stark in highly opinionated networks. This is also congruent with the much lower Jensen Shannon Divergence of the Co-evolutionary model, where the large political networks (Political Blogs and Politicians) show large disparities between the Network Formation model and the Co-evolutionary model in Jensen Shannon Divergence of the degree distributions (see Appendix A.10).

Finally, the clustering coefficient is also affected by the development of neighborhoods. For large networks, one should expect the average clustering coefficient to increase for networks with neighborhoods (see Appendix A.14). This is because the neighborhoods are well-connected locally, and exhibit very similar features to that of connected caveman networks (Watts, 1999). Again, the difference between the Network Formation model and the Co-evolutionary model is most vivid for the large, opinionated networks of Political Blogs and Politicians. Surprisingly, the average clustering coefficient for the TV Shows network is almost identical (see Appendix A.8).

Explaining the difference in results between the two models by the Co-evolutionary model's ability to generate multiple neighborhoods is also supported by the fANOVA analysis. Of the four networks that have T as their most important parameter, three of them are the political networks. Uniquely, the Political Blog network is the only network where opinion dynamics explains more of the variance than network dynamics. The parameter importance of the fANOVA analysis suggests that the ability to generate neighborhoods is an especially important part of the explanation for how political networks are generated. This suggests that political networks have more clearly delineated neighborhoods than the other considered networks.

10.3 Including empirical data in agent-based models

Although this explanation should be considered highly plausible, I have not exhausted all possible explanations for what the underlying difference between the different models are. A more complete explanation would be possible by improving on how empirical data is integrated in the model. For instance, the integration should be improved in regard to the quantity of data, as only 7 empirical networks are considered. This limits the certainty

with which conclusions can be drawn from the analysis. For instance, controlling for the variable of size was done primarily by including only one network, namely the TV Show network. The TV Show network is of a similar size as the Politicians network and was created using the same method. However, this doesn't however control for the variable of size explicitly. This could be done by varying the size of a network by for instance sampling a subset of the nodes of the network. However, it is not clear how such sampling should be done while maintaining the patterns of the original social network.

The analysis would also be more exhaustive by improving the quality of the data. The empirical validation of the agent-based models provided in this thesis compares static snapshots of empirical networks to the final result of a generative model. How well the patterns of the generative process fits with the static empirical networks is only interesting if it lends credence to the idea that the empirical networks could have been generated by the principles of the generative model. A better way to test how well the principles of the generative model describes reality would be to study empirical dynamical networks over time. A data-set which included the social connections over time as well as the social actor's opinions on key issues would be invaluable to the field of opinion dynamics. Such a data-set would allow us to gauge more directly how the mechanisms of co-evolution work in real social networks. However, such data is of course extremely hard to come by. One of the main reasons that opinions are primarily studied with agent-based models in the first place is that such data-sets are a rarity. Moreover, measuring the opinions of each social agent is a cumbersome process. Previous studies have relied on objective demographic measures to calculate similarity (Kossinets & Watts, 2009; Bener et al., 2016). If one tried instead to gauge opinions, it is unclear how one would get good measures without the inherent problems that come from self-reporting.

Although most researchers lack such high resolution in their data, there is still ample possibility to include empirical data in theoretical models. The analysis of this thesis shows that we can use recent insights from data science to not only answer how well our models match up with reality, but also why they perform well or not. This is possible specifically through the use of frontier methods such as the hyperparameter optimization used in this thesis. To the author's knowledge, the use of a fANOVA to investigate the importance of each parameter has not previously been done with agent-based models. It is the hope of this thesis that future papers will make use of this method, as it increases how interpretable the results of hyperparameter estimation are.

In conclusion, triadic closure alone can generate the characteristic patterns from social networks better than small-world and scale-free networks. This strengthens the notion that triadic closure is the basic generating mechanism for social networks (Ilany & Akçay, 2016; Jackson & Rogers, 2007, 2004). Moreover, including co-evolutionary mechanisms makes for a much better model of social networks, especially when these networks are

large and opinionated. The fact that including opinion dynamics substantially improves on network dynamical models suggests that there is a clear interdependence between the effects of network and opinion dynamics. We have seen that opinion dynamics can drastically alter the dynamics of networks. We now turn to the other side of this interplay by considering how network dynamics can change opinions.

Part III

Co-evolution's effect on opinion dynamics

The follow part of the paper explores how co-evolution plays a vital role in opinion dynamics. To explore the effect of co-evolution, we focus our attention on the Co-evolutionary model with a fixed number of agents, and explore how different parameter combinations interact. Of special interest is the effect of the parameter $P(D)$, which controls the probability of deleting negative ties.

For the following section, all simulations concluded after 10.000 time-steps ($t_{max} = 10.000$), and were made with an initial small-world network with $N = 500$ and $k = 7$.

To simplify analysis, the system focuses specifically on the case where opinions are primarily clustered in two groups. As seen in the previous part, this can be achieved by considering appropriate levels of T . This simplification was made as it allows for simpler interpretation of the system. Specifically, it allows for measuring the polarization of each agent simply as the absolute value of their opinion. When only two groups are considered, consensus is reached when agents cluster around the mean value, 0. Taking the mean of all absolute values of opinions therefore gives a measure of how polarized the system as a whole is.

11 Outcome metrics

Several outcome metrics are recorded to gauge the effect on the system that each parameter has. Some of these outcome metrics are time-dependent, while others are recorded at when the simulation concludes and the system is at its final state.

11.1 Time-dependent Metrics

Every 20th time-step, the state of the model is recorded. To track the polarization of opinions over time, the mean and standard deviation of the absolute value of the opinions of all agents is recorded. To track the effect of tie-deletion, the cumulative frequency of deleted ties is recorded. To characterize the network dynamics over time, the average clustering coefficient and the average path length of the network is recorded. At every 500th time-step, the opinion of every agent is recorded. This is done primarily to be able to visualize the trajectory of opinions directly instead of focusing on the absolute value of opinions.

11.2 Final State Metrics

After 10.000 time-steps, the network reaches its final state. At its final state, the initial opinion of every agent (O_I) and the final opinion of every agent (O_F) is recorded. Then the difference between the absolute value of the final and the initial opinion of every agent ($|O_F| - |O_I|$) is recorded. This gives a measure of whether the final opinion of agents are more extreme than their initial values were.

11.3 Correlations

To measure how much initial opinions dictate the final opinion of agents, the Pearson Correlation Coefficient between the final and the initial opinions of the agents (ρ_{O_I, O_F}) is calculated. Similarly, the Pearson Correlation Coefficient of the estimated average path length and mean of the absolute value of opinions of all agents over time ($\rho_{|O|, APL_*}$) is calculated. This correlation gives a measure for how the polarization of the agents correlates with the average distance between agents in the network.

12 Model Parameters

Only a subset of the possible values of each of the five parameter of the Co-evolutionary is considered to make the Co-evolutionary model computationally tractable. All different possible parameter combinations of this subset are simulated, resulting in 3.780 different parameter combinations. The considered parameter values are:

$$R \in \{0.1, 0.3, 0.5\}$$

$$\alpha \in \{0.05, 0.10, 0.15, 0.20, 0.25\}$$

$$\beta \in \{0.00, 0.05, 0.10, 0.15, 0.20, 0.25\}$$

$$T \in \{0.6, 0.7, 0.8, 0.9, 1.0, 1.1, 1.2\}$$

$$P(D) \in \{0.0, 0.2, 0.4, 0.6, 0.8, 1.0\}$$

Some justification of the different parameters are in order. On the power of social influence, it is fair to assume that its power comes from the frequency of social interactions. In terms of the parameters of this model, this corresponds to assuming that the values of α and β should be relatively low. One rarely reaches a perfect compromise when debating other people, which is what $\alpha = 1$ would reflect. It therefore makes sense to assume that it is realistic to have $\alpha \ll 0.5$. The same argument applies to β . Negative social influence is as mentioned less empirically validated than positive social influence (Takács et al., 2014). For this reason, simulations with $\beta = 0$ are also included. As mentioned, T was specifically chosen to limit the number of clusters in the resulting network. Within this limit, the specific parameters of T were chosen as they gave the most interesting results in preliminary analyses. For $T > 1.2$, the model more or less always converges to consensus, no matter what the values of all other parameters are. In the same vein, only very few simulations avoid polarization for $T < 0.6$. The parameters of interest for this thesis are therefore in between these two values, where $0.6 \leq T \leq 1.2$. For values of R , there is good reason to believe based on previous empirical work that most interactions are made via triadic closure and are not random (Kossinets & Watts, 2009). Therefore, the interesting values of R are values which reflect that at least half of new connections are made via triadic closure ($R \leq 0.5$). Finally, as co-evolution is one of this thesis's primary interests, I study the whole range of the co-evolutionary force of $P(D)$, the probability of tie-deletion. I therefore consider a range of values resulting in different levels of co-evolution in the system. Values of interest range from no co-evolution to perfect co-evolution i.e. $0 \leq P(D) \leq 1$.

13 Results

The effect of all the different model parameters as well as the correlations of interests outlined previously are all analyzed and reported. However, not all of these results are centered around the main argument of this thesis. For this reason, auxiliary visualizations and analysis is instead reported in the *Appendix*.

13.1 The effect of the different model parameters

All parameters directly associated with opinion dynamics in the model show clear effects on polarization. Higher values of the threshold, T , decrease polarization (see Appendix A.11). A stronger force of positive social influence, α , decreases polarization (see Appendix A.12). A stronger force of negative social influence, β , increases polarization (see Appendix A.13). Notably, when $\beta = 0$, all simulations reach consensus, regardless of all other parameters (see Figure 12). Similarly, increasing the probability of negative tie-deletion has no effect on polarization when $\beta = 0$ (see Figure 11). The drivers of polarization are low values of T and high values of β . When T is sufficiently small and β is sufficiently large, all agents will have more extreme opinions at the end of the simulations than they had in the beginning. Such conditions are both found with $T = 0.6$ and $\beta \geq 0.1$ as well as $T \leq 0.9$ and $\beta = 0.25$. With intermediate values of β and T , whether the opinions of a simulation becomes polarized or not is largely dictated by $P(D)$ (see Figure 12). The results indicate that higher values of $P(D)$ results in lower values of $|O|$ on average. The strength of the effect is modulated by how random new connections are made in the network. The effect of cooperation facilitation by deleting negative ties is diminished with increases in R (see Figure 10). When $P(D) = 0$, $T \leq 1.0$ and $\beta > 0$, the network always polarizes (see Figure 12).

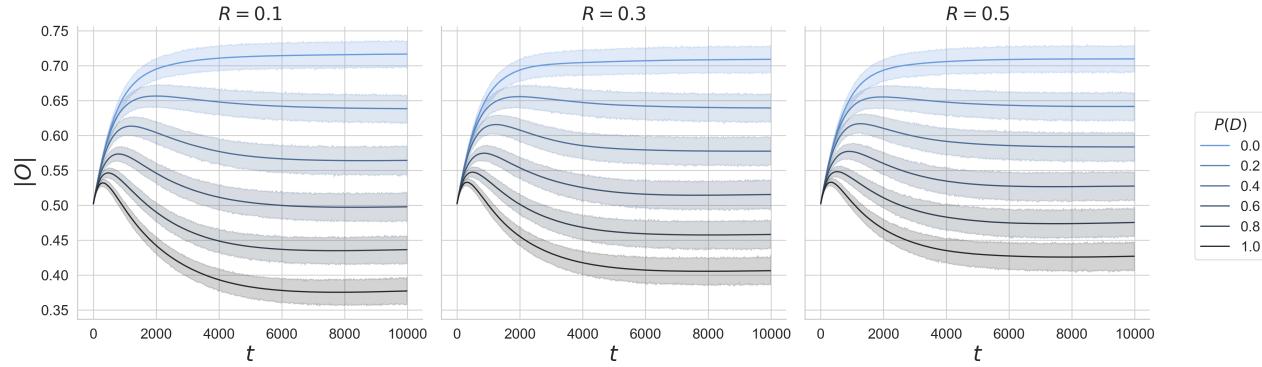


Figure 10: $P(D)$'s effect on polarization over time. The x-axis shows the time-steps, t and the y-axis shows the absolute value of the opinions of agents, $|O|$. Colors indicate different probabilities of tie-deletion of dissimilar agents. Lines indicate the mean value of each time-step, aggregated over all parameters excluding the probability of tie-deletion and randomness. Shaded regions show the 95% confidence intervals. The three different panels show simulations with different propensities for generating new ties via triadic closure rather than randomly, R .

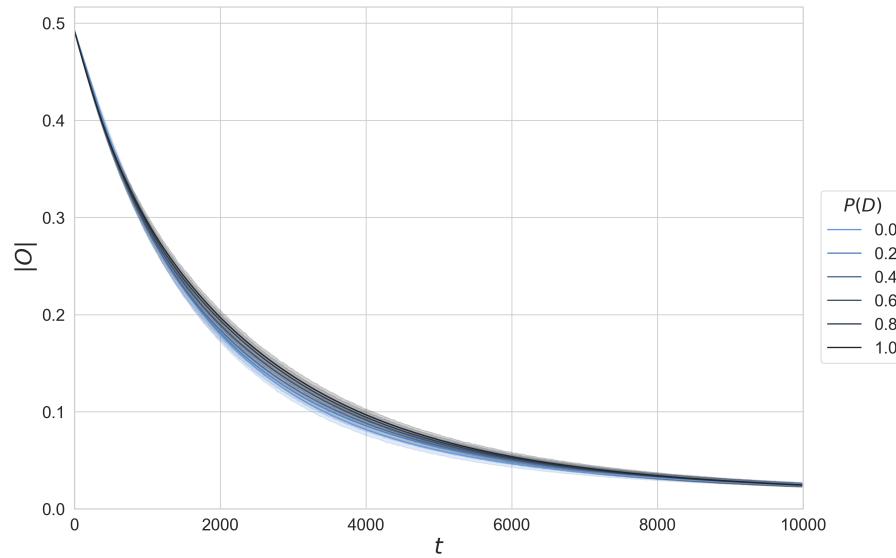


Figure 11: The effect of $P(D)$ without negative learning. The x-axis shows the time-steps, t and the y-axis shows the absolute value of the opinions of agents, $|O|$. Colors indicate different probabilities of tie-deletion of dissimilar agents. Lines indicate the mean value of each time-step, aggregated over all parameters with $\beta = 0$, excluding the probability of tie-deletion. Shaded regions show the 95% confidence intervals.

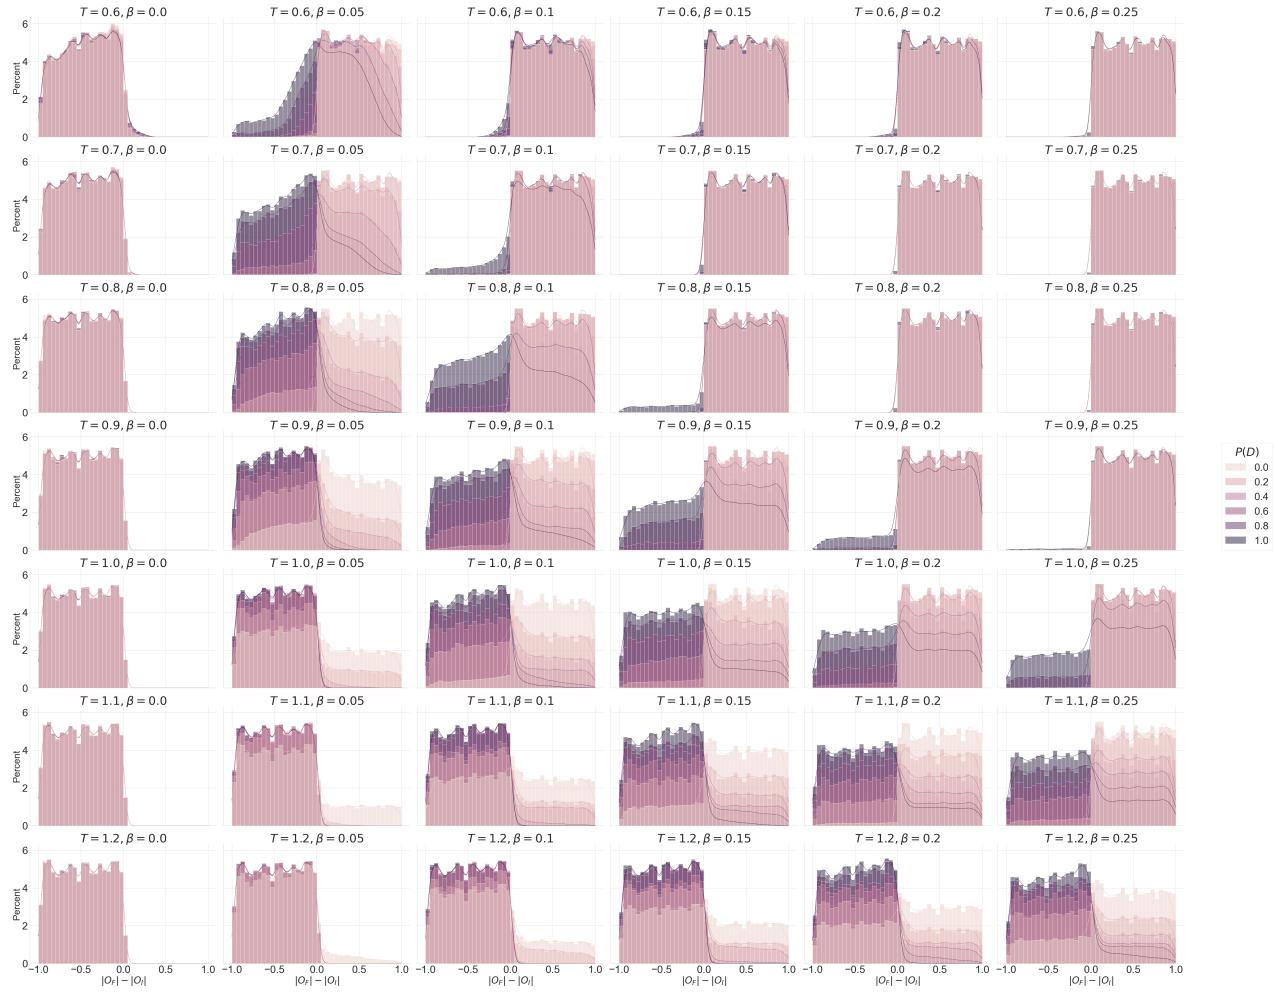


Figure 12: Distribution of opinion changes across different conditions. Rows display different values of thresholds, T , and columns indicate different values of negative social influence, β . Different colors show different values of the probability of tie-deletion, $P(D)$. The x-axis shows the difference between the absolute value of the final and initial opinions of agents, $|O_F| - |O_I|$. When this value is positive, the final opinion of agents is more extreme than the initial opinion of agents. When this value is negative, the final opinion of agents is less extreme than their initial opinion. Y-axis shows the percent of values within each bin.

13.2 Correlations

The correlation between the absolute value of opinions and the average path length, $\rho_{|O|,APL}$ is largely dictated by the probability of tie-deletion, $P(D)$. As $P(D)$ increases, so does $\rho_{|O|,APL}$. When $P(D) = 0$, there is no consistent correlation between the two variables (see Figure 13).

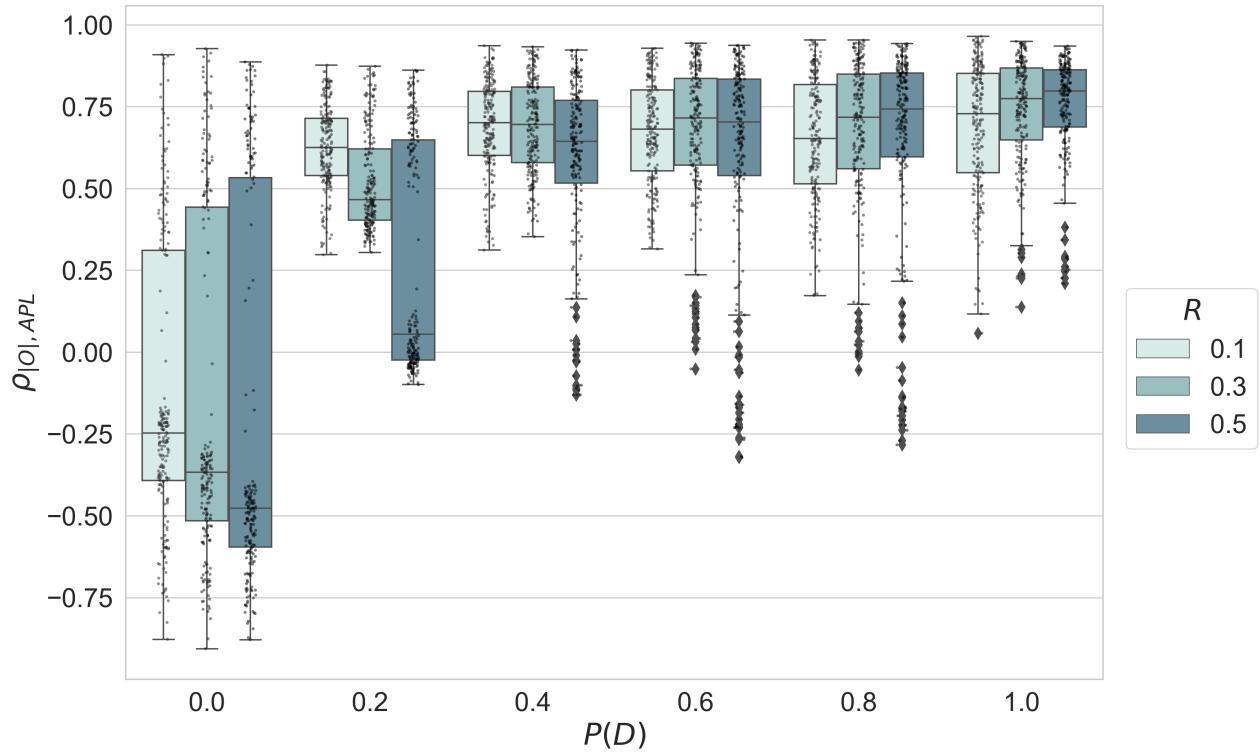


Figure 13: Correlations between absolute opinion and average path length. X-axis shows different values of the probability of deleting negative ties, $P(D)$. The Pearson Correlation Coefficient between the mean of absolute value of opinions and the estimated average path length, $\rho_{|O|,APL}$ is shown on the y-axis. Boxes indicate the quartiles and whiskers indicate the 1.5 of the inter-quartile range outside of the quartiles. Points outside this range is depicted as diamonds. Smaller dark dots show the correlation coefficient for individual simulations.

14 Discussion

14.1 Reiterating previous literature

The observed effect of positive social influence, α , negative social influence, β , and the threshold, T , are intuitive and in line with previous research on similar models (Flache et al., 2017). When α is high, there is a stronger force pulling the opinion of agent's closer together. Higher values of α therefore results in lower values of polarization (see Appendix A.12). Although α has an effect on the polarization of opinions, the effect is small in comparison to the effect of β and T (see Appendix A.11, A.13). This is primarily caused by the considered values of T and its operationalization. When $\beta = 0$, all simulations reach a consensus (see Figure 11, 12). If lower values of T had been considered, this would most likely not be true. This was seen in the second part of this thesis, where lower values of T generated neighborhoods of different opinions and

is also found in similar computational models (Flache et al., 2017; Sasahara et al., 2021). When β is large, almost all simulations polarize (see Figure 12). This is partly because it is a stronger force than α in how it is operationalized. The effect of pushing and pulling opinions is defined as being proportional to the distance between agents. By definition, agents that are pushed away from each other are further apart than agents that are pulling closer together. In this sense, β is a stronger force than α . Only in aggregate can the effect of α become stronger than the effect of β . This is only possible when the number of positive interactions is much higher than the number of negative interactions. The effect of T should therefore not be surprising. With larger values of T , agents are more open-minded in the sense that they can cooperate with more diverse agents, and will therefore have more positive interactions than negative. To see this, notice that if an agent has an opinion of 0 and $T \geq 1$, this agent will not be able to have any negative interactions in this model. The opposite effect is witnessed when T is low. This limits the subset of agents the agent can cooperate with, which makes negative interactions more plausible.

14.2 The importance of co-evolution

As mentioned in the *Introduction*, previous models have primarily focused on studying systems of static networks. The results of this thesis shows the consequences of this assumption. When the network does not co-evolve and there is some negative social influence in the system ($P(D) = 0$ and $\beta > 0$), the network always polarizes (see Figure 12). This is not the case when the network is dynamic and co-evolves, which can often prevent negative social influence from causing polarization in the network (see Figure 10). This in and of itself shows that the co-evolution of network and opinion dynamics is a vital piece of the puzzle of polarization.

14.3 Tie-deletion facilitates cooperation

One of the main findings of this thesis is that higher probabilities of tie-deletion can prevent polarization and facilitate cooperation. This is especially the case when new ties are created via triadic closure and not randomly (see Figure 10). This happens because of two effects in the system. The first is that deleting negative ties will cause fewer negative social interactions, which limits the effect of β . When negative ties are deleted, agents will assort based on their opinions. Deleting negative ties is important in its own right because it deletes ties that would otherwise cause further polarization. But because of assortment, the effect of tie-deletion interacts with how random new ties are created. When agents assort based on opinions, the heuristic that "a friend of a friend is also a friend" is very likely true. Triadic closure will result in even more positive interactions, because

similar to what is found in empirical networks, distance in the social network is highly indicative of distance in similarity (Kossinets & Watts, 2009). As such, new ties created randomly rather than via triadic closure will tend to connect more dissimilar agents. In other words, increasing randomness decreases assortment, which leads to more negative encounters (see Appendix A.19). This also echoes the finding from the opinion dynamics literature stating that when new connections are made randomly, opinions will tend to polarize more quickly when they include negative influence (Flache & Macy, 2006, 2011; Turner & Smaldino, 2018). Here it is showed that a plausible explanation for this finding is that social networks are assorted based on opinions. By increasing how random new connections are, one increases the probability of having negative interactions.

Positive assortment of opinions is also closely linked to why there is a strong positive correlation between the absolute value of opinions and the average path length when negative ties are deleted. As the network is co-evolutionary, changes in the network will reflect changes of opinions. The strong positive correlation between the absolute value of opinions and the average path length suggest that when opinions polarize, so does the network (see Figure 13). Importantly, this only happens when negative ties are deleted. When $P(D) \gg 0$, negative ties will tend to be deleted when the opinions of agents tend to polarize. Because of positive assortment, agents that have dissimilar opinions will tend to be further away from each other in the network. As such, long-range ties in particular will suffer when opinions polarize. This will make highways of information which connects distant parts of the network much less likely, which increases the average path length substantially. When the network depolarizes, the opposite happens. Some long range ties will be made by random, and because the network is depolarized, these will stay intact. This will make further long-range connections more probable, as the two clusters can connect via triadic closure.

14.4 Distance in opinion and network space

The results that have been presented and discussed have all been made on the basis of averages of several simulations. Although such averages are well-suited to diagnose the general tendencies of such a system, some intricacies can be lost in averages. Often a more detailed look at representative simulations are fruitful for understanding why the system behaves as it does (Turner & Smaldino, 2018). For this thesis in particular, focus on a subset of simulations makes it possible to visualize the raw trajectories of opinions and networks simultaneously (see Figure 14, 17). By describing representative simulations, it is much easier to understand how tie-deletion might facilitate cooperation in systems such as the one investigated in this thesis. To this end, a sharper focus is made on one of the conditions where tie-deletion and triadic closure is the difference between polarization and consensus. The representative conditions investigated in the following section are simulations

where the threshold is relatively low ($T = 0.8$), there is some negative social influence ($\beta = 0.1$) and positive social influence ($\alpha = 0.15$). This combination of parameters only avoids polarization when $P(D) \geq 0.8$ for the reasons highlighted in the previous section (see Figure 12).

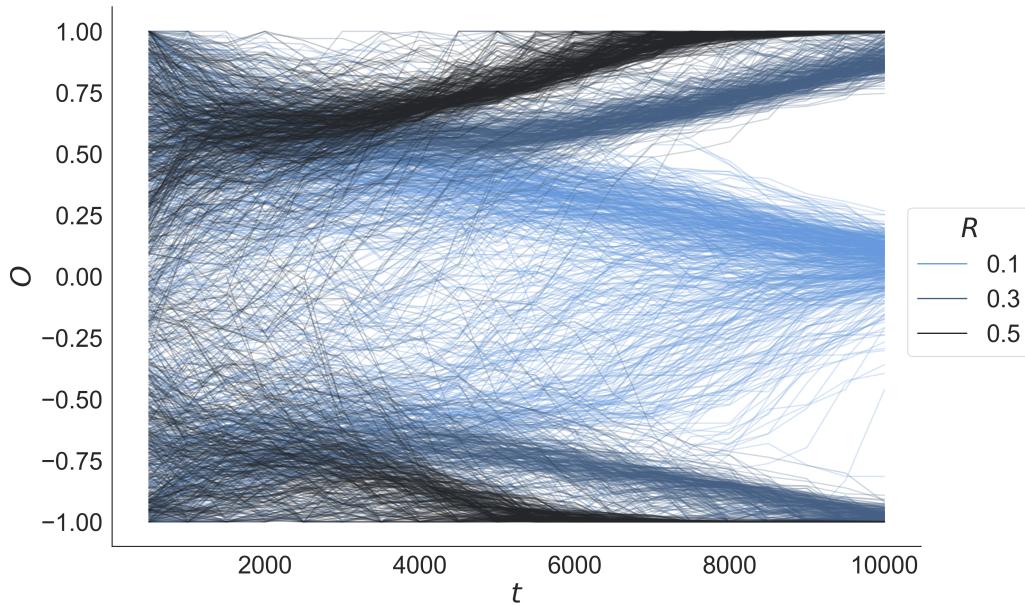


Figure 14: Lineplot of opinions over time. X-axis shows time-steps of the simulation, y-axis shows the opinion of agents in the simulation. Colors show different values of R . All simulations were generated with the parameter values of $T = 0.8$, $\alpha = 0.15$, $\beta = 0.1$, $P(D) = 1$ and with the same random seed.

In these representative conditions, how random new connections are plays a vital role for whether opinions polarize or not. There is a stark difference between the trajectories of opinions for different levels of R (see Figure 14). Only the simulation with $R = 0.1$ avoids polarization. Additionally, when $R = 0.5$, the polarization is much faster than when $R = 0.3$. For all the representative conditions, there is a strong connection between the average path length of the network and the polarization of opinions (see Figure 15, 16). As mentioned, the mechanism that links these two variables is positive assortment. This can clearly be seen by inspecting the opinions of the agents in the network directly (see Figure 17). After only 2000 time-steps, there is a clear gradient in the colors of the agents of the networks, which is the result of positive assortment. Positive assortment is achieved directly by negative tie-deletion in this model. Additionally, positive assortment is increased indirectly by triadic closure. This is exactly the kind of mechanism described in the literature as structural homophily (Asikainen et al., 2020; Peixoto, 2022). This also goes to show that tie-deletion is not necessarily an important mechanism in itself. Rather, it is important because it increases positive assortment. Presumably, any mechanism that will increase positive assortment in these kinds of models will decrease polarization. This point is also strengthened by the difference between the simulations considered here. The reason why the simulations

with $R > 0.1$ polarize is because increasing randomness decreases positive assortment. Counter intuitively, tie deletion in combination with triadic closure facilitates cooperation because it isolates groups, not in spite of it.

The representative simulations also clearly show how polarization in opinions leads to a less globally connected network, which increases the average path length. To see why this happens, one can use a spring as an analogy for the network. As long-range ties are deleted, the tension in the spring rises, and the network starts to be better characterized as a network of two factions, connected mostly by a group of agents with values around 0. As distant parts of the network reconnect, the tension of the spring is lowered, and all the parts of the network become closely linked again. When polarization increases, there will be less of these long-range ties, but there might also be fewer "middlemen" to facilitate cooperation across the aisle. By middlemen, I refer to the agents of moderate opinions, which tend to be situated in-between the two main clusters of agents. Visually, these middlemen are the nodes of cyan colors, that are situated in-between the two major clusters of dark and light colors (see Figure 17). I draw attention to these agents, as they are critical for the cooperation of the network. Without them, the spring snaps. It does so because the different clusters cannot have positive interactions with each other, as the average opinion of each cluster will tend to be too dissimilar. By interacting only through these middlemen, there is a chance of bridging the gap by isolating the different groups from each other. If these middlemen become part of either cluster before polarization has decreased sufficiently for the different groups to have positive interactions, the only possible outcome is polarization. The existence of middlemen is therefore a necessary condition for avoiding polarization in this model.

Another way to understand why cooperation is facilitated by tie-deletion, is to realize that the probability of tie-deletion can be seen as the rate at which the structure of the network adjusts to opinions. When the $P(D) = 1$, the structure of the network adjusts at the same speed of opinions, while lower values will indicate that the network adjusts at a slower pace than opinions. In this interpretation, consensus can be achieved if the social network adjusts fast enough. Having faster adjustment of the structure of social interactions is critical, because the system is highly path dependent (Turner & Smaldino, 2018; Macy et al., 2021). When the network doesn't adjust fast enough, agents will already have had several negative interactions, which will have pushed them further towards an extreme value of opinions. This will make them more likely to polarize other agents in the future. This will be the case for all agents including middlemen. Adjusting too slowly can therefore burn the only bridge that is available for facilitating cooperation.

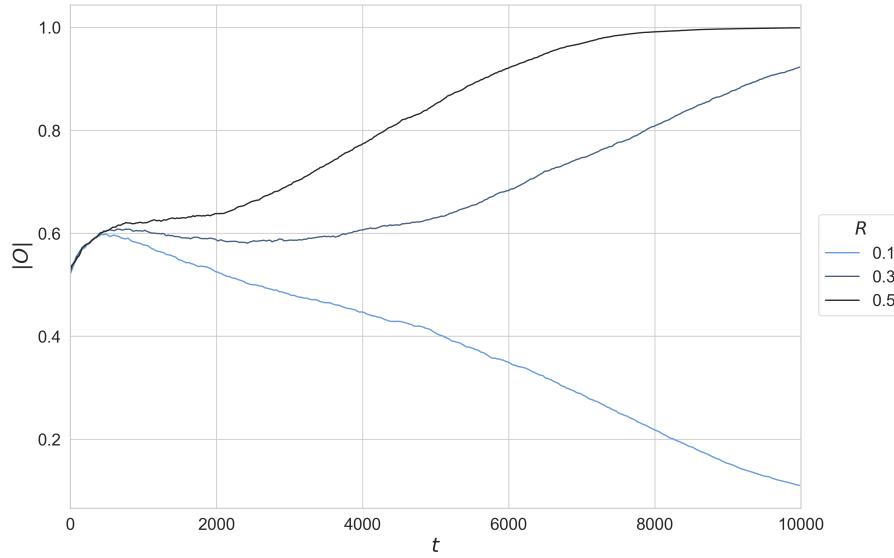


Figure 15: Absolute opinions of representative simulations over time. X-axis shows time-steps, and the y-axis shows the absolute opinion of agents. Colors indicate different levels of randomness. Lines show single simulations over time which were simulated using the same random seed. These simulations had parameter values of $T = 0.8$, $\alpha = 0.15$, $\beta = 0.1$ and $P(D) = 1$.

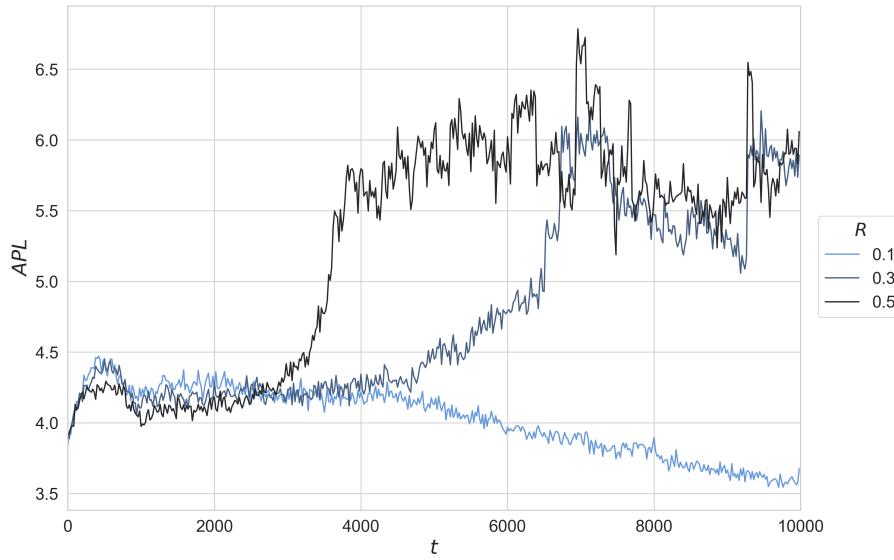


Figure 16: Average path length of representative simulations over time. X-axis shows time-steps, and the y-axis shows the average path length of the network. Colors indicate different levels of randomness. Lines show single simulations over time which were simulated using the same random seed. These simulations had parameter values of $T = 0.8$, $\alpha = 0.15$, $\beta = 0.1$ and $P(D) = 1$.

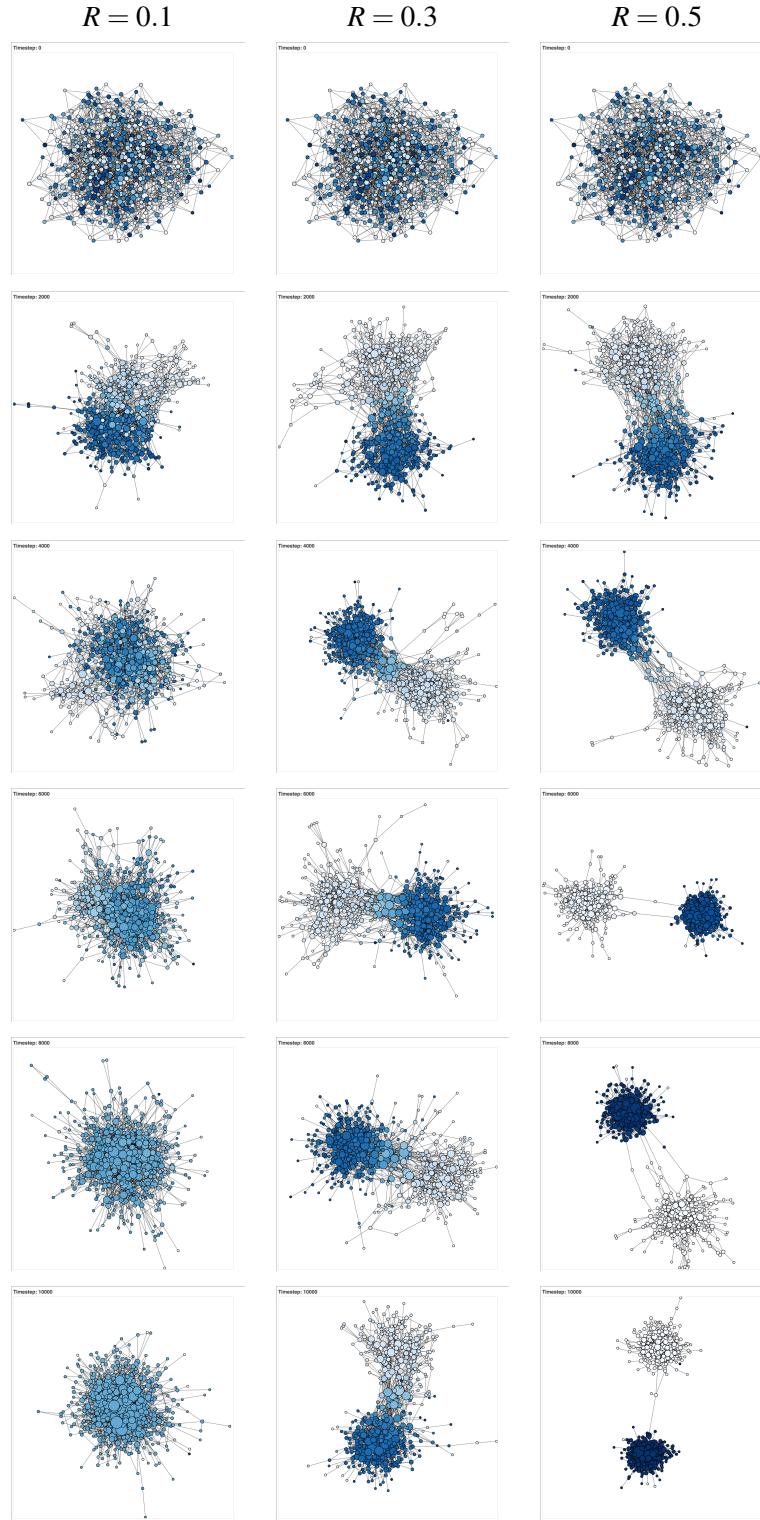


Figure 17: Network evolution for different levels of randomness. Columns show different values of randomness, R . Rows are different time-steps of the simulations, increasing in increments of 2000. Nodes represent agents, lines represent edges between agents. Agents are colored by their opinion. Dark blue colors represent values close to -1, and lighter colors represent opinions close to 1. The size of the nodes of agents reflects their degree, with higher degrees shown as larger nodes.

15 The point of no return for polarization

The simulations investigated in detail here also highlight the fact that there seems to be points of no return in regard to polarization. This is largely due to the fact that polarization breeds more polarization. If the system as a whole reaches a critical tipping point of polarization, the system will not be able to reach consensus as polarization becomes self-reinforcing (see Appendix A.17). The same conclusion is found in recent models of polarization where exogenous shocks to the system, such as a global pandemic, create a common threat for all agents (Macy et al., 2021; Levin et al., 2021). They find that if a system has already reached a certain level of polarization, it is almost impossible to introduce measures that can depolarize the system. Even in the face of deadly exogenous shocks, the system is unable to take action as polarization paralyzes their ability to act. Although identifying such tipping points is not the focus of this thesis, the results still show that early countermeasures are essential to avoid polarization.

Part IV

Overall Discussion

The central finding of both parts of this thesis is that co-evolution of opinion and network dynamics is a vital factor to consider for explaining both systems. For the first part of the paper, this is argued by the fact that models that include opinion dynamics drastically outperform network formation models when the target networks are large and opinionated. For the second part of the paper, this is argued by the fact that co-evolution is vital for explaining how social systems can avoid polarization in the face of mechanisms such as negative social influence. Without co-evolution, even small amounts of negative social influence can cause the system to polarize (see Figure 12). In both cases, the reason why co-evolution becomes so vital is that co-evolution allows for systematically generating neighborhoods in the network. Not only is this the phenomenon that makes the Co-evolutionary model a better explanation for large, opinionated networks, but it is also the phenomenon that explains how tie-deletion can facilitate cooperation. By generating two distinct neighborhoods, social exposure to the opposing group is minimized, which limits the effect of negative social influence (see Figure 17).

This thesis also highlights that many of the emergent properties of social networks can be generated by surprisingly simple mechanisms. The model considered here does not include mechanisms of choice homophily nor does it include agents with a goal or a strategy. By including only a propensity to delete ties to dissimilar agents as a co-evolutionary mechanism, the model can generate realistic social networks. This includes emergent properties such as the connection between distance in similarity and distance in the social network as seen in empirical work (Kossinets & Watts, 2009) (see Figure 2). However, one should be hesitant to conclude that tie-deletion of dissimilar agents is the general generating mechanism which gives rise to these properties in social networks. As was alluded to in the case-study of representative simulations, tie-deletion is important because it increases positive assortment. Previous results show that similar results could be achieved by implementing a mechanism of choice homophily instead, where agents connect to other agents based on their similarity (Asikainen et al., 2020). Regarding the effect tie-deletion has on avoiding polarization, such a model would only produce similar results if these new connections were rewired primarily from dissimilar agents. If this is not the case, the model would not be able to avoid the initial negative social influence which comes about from having opinions be randomly distributed.

If social agents choose new social ties based on how similar they are, i.e. choice homophily, one must assume that these agents can process information about the other agent before interacting with them. This is not an unrealistic assumption, as choice of clothes and other social markers provide effective ways of identity signaling (Smaldino, 2022). In this thesis, it is shown that observed homophily can occur even when agents do not share information, but purely through structural homophily. Structural homophily in the model emerges because of the important interaction between tie-deletion and triadic closure. This also highlights how important triadic closure is as a generating mechanisms for social systems. As we saw in the second part of this thesis, triadic closure alone can generate many of the characteristics associated with social networks. This serves as a robustness check from previous implementations of similar mechanisms in other models (Jackson & Rogers, 2004, 2007). When triadic closure is combined with homophily and social influence, the interaction between these mechanisms can explain more complicated social phenomena such as neighborhood formation. Previous work suggest that the interaction of triadic closure with homophily can also generate core-periphery structures in social networks, which is another stable feature of social networks (Asikainen et al., 2020). The same interaction also has a substantial impact on how opinions are shared in computational simulations (see Figure 10, 17). This is an important note as previous work focusing on dynamical networks in opinion dynamics do not include empirically motivated network dynamics, but simply generate new ties at random (Kozma & Barrat, 2008). This thesis shows that the inclusion of even simple, realistic network dynamics can have large effects on how opinions change over time.

16 Contrast to previous work

The agent-based model in this thesis is similar in many ways to the recent agent-based models which investigate the emergence of echo chambers in social networks (Sasahara et al., 2021). Despite their similarity, they find completely opposite results regarding negative tie-deletion. Contrary to the findings of this thesis, Sasahara et al. (2021) reports that increasing negative tie-deletion accelerates polarization. It is therefore important to understand why these results are in so stark contrast. First, the network considered in Sasahara et al. (2021) is a directed network, meaning that relations need not be reciprocal. This could have important implications because it allows for social influence to be asymmetric. Secondly, the number of agents considered in Sasahara et al. (2021) is only 100 whereas it is 500 in the model presented in this thesis. This could also be consequential. As we've seen in the representative simulations, tie-deletions can only facilitate cooperation if there are enough middlemen that can reunite the different clusters. By simply having more agents in the simulation, it increases the number of expected middlemen agents, which can help in facilitating cooperation between agents of very different opinions.

Perhaps the most critical difference between the models is that Sasahara et al. (2021) does not include negative social influence in their model. In the model of Sasahara et al. (2021), social relations outside the threshold of agents have no social influence. In other words, negative interactions are assumed to be without a cost for the agents involved. For the parameter values considered in the model of this thesis, tie-deletion is only effectively facilitating cooperation when there is negative social influence in the system (see Figure 11). As the main difference between these models is in whether forces akin to negative social influence push distant opinions apart, more research needs to be done in regard to how pervasive forces like negative social influence are in shaping opinions. In fact, whenever the costs outweighs the benefits associated with encountering dissimilar individuals, tie-deletion of dissimilar individuals will be beneficial to the individuals. Identifying these costs and benefits of social interactions is needed to develop better models and to decipher the role tie-deletion might play in facilitating cooperation or accelerating polarization.

17 Lacking realism of the model

As was highlighted in the very beginning of this thesis, a model's value comes from the fact that it simplifies reality. But its usefulness comes from the fact that it hasn't left out crucial factors contributing to the system of interest. Although this thesis shows that including co-evolution is vital to understanding opinion dynamics

and network formation, it leaves out many important mechanisms. This was done in the name of simplicity, but some of these simplifications could potentially have large impacts on results of this thesis.

17.1 Non-cognitive agents

The first important simplifying assumption is that agents are not cognitive agents. Including more realistic cognitive features for the agents might provide important changes to the dynamic of the system. The agents considered in this thesis have no agenda and therefore do not try to convince each other of their opinions. In this way, there is never a case where a persuasive argument pulls one of the agent without changing the opinion of the other agent. The fact that some arguments are better than others and that some agents are more stubborn than others likely has an important part to play in opinion dynamics (Flache et al., 2017; Ghaderi & Srikant, 2014; Yildiz, Ozdaglar, Acemoglu, Saberi, & Scaglione, 2013).

The agents from the model of this thesis have perfect information regarding their peers. All agents perceive each other's opinions perfectly, as there is no noise in the interaction between agents. Social interactions are filled with misspoken words and misinterpreted arguments (Jussim & Osgood, 1989). These misspoken words and misinterpretations are likely not random, but rather skewed by the same cognitive heuristics and biases that shape other parts of how we include information into our decisions (Arceneaux, 2012). This includes effects such as confirmation bias, which play an important role in how we update our opinions (Allahverdyan & Galstyan, 2014). Other models of opinion dynamics have previously included noisy transmissions between agents to represent imperfect communication (Sîrbu, Loreto, Servedio, & Tria, 2017; Su, Chen, & Hong, 2017). Noisy interactions are highlighted here, as they could be important for facilitating cooperation in the types of models discussed in this thesis. Noisy interactions allow agents that would not normally interact positively to perceive each other as more similar by chance (Allahverdyan & Galstyan, 2014; Su et al., 2017). Including more realistic cognitive mechanisms might change the dynamics of the system in important ways which are not considered in this model. Such a mechanism would be modelling agents as having not only a value describing what opinion they have, but also how certain they are of their opinions. This could be implemented by re-imagining updating opinions as Bayesian updating, considering other's opinions as new evidence (Allahverdyan & Galstyan, 2014). Instead of describing the opinions of agents using a single number, their opinions could be represented by a distribution of values. Such a distribution could be a simple Gaussian distribution, where the mean indicates the expected opinion of an agent, and the standard deviation describes the certainty of the agent's opinion.

17.2 Shrinking Variance

Because no new agents are added to the network and because interactions are non-noisy, the variance in opinions shrinks over time (see Appendix A.15). At first glance, this might not seem very problematic. People tend to conform to the majority of their social groups (McPherson et al., 2001), but one shouldn't expect near perfect consensus or perfect polarization on most issues. Without the introduction of new variance into the system, the diversity of opinions will vanish. Introducing mechanisms that keep a sufficient amount of variance would make for a more realistic model. Such a mechanism would also decrease the path dependency of the system. Previous models have sustained diversity of opinions with individualistic agents, which strive to be different from their peers (Mäs, Flache, & Helbing, 2010). This can produce cycles of opinions, which would radically change the behavior of the model, as neighborhoods would likely balkanize when they hit a certain number of members.

17.3 Social influence exclusively

The model investigated here assumes that one's opinion is only changed by local interactions. Although this is in line with most models of opinion dynamics, there are good reasons to believe that global trends and tendencies are powerful forces that shapes opinion (Bener et al., 2016). These global tendencies might shape our opinions more and more, as media consumption rises. Global perturbations to the system could be powerful forces that potentially could have large impacts on whether the system can reach consensus. These forces would be especially counterproductive for reaching consensus, if they only influence certain parts of the system. The power of such effects is not hard to imagine. Partisan news media are critical for how opinions are formed (Pennycook & Rand, 2019). As media consumption rises and becomes a more and more important force for what information we receive, these force might already be a more important variable to consider than the local interactions of agents (Bener et al., 2016; Strömbäck, Djerf-Pierre, & Shehata, 2013).

18 A broader perspective

As a closing remark in the discussion, the model points to some ill-omens in regard to real-world polarization and the structure of our social networks. As we've seen, the agent-based model of this thesis points to the fact that co-evolution of networks and opinions leads to distance in opinions being closely linked to distance in social networks. It is therefore interesting to see whether how connected we are and how much we disagree corresponds in empirical networks. As was pointed out in the introductory paragraph, polarization has been on

the rise in many western democracies over the last two decades (Boxell et al., 2020; Center, 2014a). On the other hand, estimations of the average path length between users on Facebook has decreased from 2011 to 2016 (Bhagat, Burke, Diuk, Edunov, & Filiz, 2016). The co-evolutionary model investigated in this thesis suggests that the increase in connectivity and the simultaneous increase in polarization could be a bad omen. If anything like negative influence characterizes interactions on social media (i.e. outrage), the decrease in average path length will make negative encounters more likely, and most likely increase rather than decrease polarization. Although being more connected might sound good on paper, it might be detrimental to our ability to cooperate in the long run. This effect is congruent with previous literature, where the introduction of long-range ties sharply increases polarization of opinions in the population (Flache & Macy, 2011; Turner & Smaldino, 2018).

19 Conclusion

This thesis attempts to fill two critical gaps in the existing literature on opinion dynamics; lack of integration of empirical data and lack of investigations on the effect of co-evolution on opinion dynamics. To better integrate empirical data, this thesis uses and develops upon the method of hyperparameter optimization for agent-based models. The results of the optimization reaffirm previous studies which show that the mechanism of triadic closure can generate many of the important characteristics of social networks. This thesis also finds that including opinion dynamics into models of network formation significantly improves how well these models can generate patterns from real world networks. This is especially the case when the considered social networks are large and opinionated. Such networks are constituted by distinct neighborhoods, which the Co-evolutionary model can generate systematically but the Network Formation model cannot. This explanation is supported by the previously unused method of fANOVA of hyperparameter importance, highlighting this method's usefulness in interpreting the results from hyperparameter optimization algorithms.

To study the impact of co-evolution on opinion dynamics, the results of the Co-evolutionary model were investigated in greater detail. The results reiterate known dynamics of social influence and homophily found in past work investigating similar models. But the results also show that co-evolution has a large effect in the simulations. The results show that the speed of co-evolution is frequently the difference between whether a simulation polarizes or reaches consensus. When social interactions between dissimilar individuals are costly, positive assortment of similarity can help agents overcome these costs by primarily interacting with agents like themselves. This is also why the findings of this thesis are in stark contrast to the results from a recent model in opinion dynamics, which show that tie-deletion accelerates polarization. This model has no such cost associated

with negative interactions, and therefore agents do not stand to gain anything by avoiding such interactions.

This thesis shows that by combining simple mechanisms from network and opinion dynamics, it is possible to generate key aspects of the complex systems of social networks. Using only tie-deletion as a co-evolutionary mechanism, the network can generate neighborhoods of like-minded agents as well as establishing the link between similarity and distance in the network, which is a robust empirical finding from social networks. This thesis highlights the important link between opinion and network dynamics, and shows that explanations of both dynamics stand to gain from incorporating the other.

20 Acknowledgements

This thesis was made in collaboration with the Cognitive and Information Science department of the University of California, Merced. I would like to express my gratitude to Cody Moser, Chanuwas Aswamenakul and Alejandro Perez Velilla for their valuable feedback. I am also grateful for the guidance and advice offered by Paul Smaldino throughout the project.

References

- Adamic, L. A., & Glance, N. (2005, August). The political blogosphere and the 2004 U.S. election: divided they blog. In *Proceedings of the 3rd international workshop on Link discovery* (pp. 36–43). New York, NY, USA: Association for Computing Machinery. Retrieved 2022-04-06, from <https://doi.org/10.1145/1134271.1134277> doi: 10.1145/1134271.1134277
- Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019, July). Optuna: A Next-generation Hyperparameter Optimization Framework. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining* (pp. 2623–2631). New York, NY, USA: Association for Computing Machinery. Retrieved 2022-02-24, from <https://doi.org/10.1145/3292500.3330701> doi: 10.1145/3292500.3330701
- Allahverdyan, A. E., & Galstyan, A. (2014, July). Opinion Dynamics with Confirmation Bias. *PLOS ONE*, 9(7), e99557. Retrieved 2022-03-31, from <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0099557> (Publisher: Public Library of Science) doi: 10.1371/journal.pone.0099557

- Andris, C., Lee, D., Hamilton, M. J., Martino, M., Gunning, C. E., & Selden, J. A. (2015, April). The Rise of Partisanship and Super-Cooperators in the U.S. House of Representatives. *PLOS ONE*, 10(4), e0123507. Retrieved 2022-02-24, from <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0123507> (Publisher: Public Library of Science) doi: 10.1371/journal.pone.0123507
- Arceneaux, K. (2012). Cognitive Biases and the Strength of Political Arguments. *American Journal of Political Science*, 56(2), 271–285. Retrieved 2022-03-31, from <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1540-5907.2011.00573.x> (eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1111/j.1540-5907.2011.00573.x>) doi: 10.1111/j.1540-5907.2011.00573.x
- Asikainen, A., Iñiguez, G., Ureña-Carrión, J., Kaski, K., & Kivelä, M. (2020, May). Cumulative effects of triadic closure and homophily in social networks. *Science Advances*. Retrieved 2022-01-26, from <https://www.science.org/doi/abs/10.1126/sciadv.aax7310> (Publisher: American Association for the Advancement of Science) doi: 10.1126/sciadv.aax7310
- Backstrom, L., Boldi, P., Rosa, M., Ugander, J., & Vigna, S. (2012). Four degrees of separation. In *Proceedings of the 4th annual acm web science conference* (pp. 33–42).
- Bail, C. A., Argyle, L. P., Brown, T. W., Bumpus, J. P., Chen, H., Hunzaker, M. B. F., ... Volfovsky, A. (2018, September). Exposure to opposing views on social media can increase political polarization. *Proceedings of the National Academy of Sciences*, 115(37), 9216–9221. Retrieved 2022-03-19, from <https://www.pnas.org/doi/10.1073/pnas.1804840115> (Publisher: Proceedings of the National Academy of Sciences) doi: 10.1073/pnas.1804840115
- Barabási, A.-L., & Bonabeau, E. (2003, May). Scale-free networks. *Scientific American*, 288(5), 60–69. Retrieved 2022-03-17, from <https://doi.org/10.1038/scientificamerican0503-60> doi: 10.1038/scientificamerican0503-60
- Baumann, F. (2021). *Modeling opinion dynamics on networks: How social influence shapes the formation of consensus and polarization* (Unpublished doctoral dissertation). Humboldt Universitaet zu Berlin (Germany).
- Baumann, F., Lorenz-Spreen, P., Sokolov, I. M., & Starnini, M. (2020, January). Modeling Echo Chambers and Polarization Dynamics in Social Networks. *Physical Review Letters*, 124(4), 048301. Retrieved 2022-01-26, from <https://link.aps.org/doi/10.1103/PhysRevLett.124.048301> (Publisher: American Physical Society) doi: 10.1103/PhysRevLett.124.048301
- Bener, A. B., Çağlayan, B., Henry, A. D., & Prałat, P. (2016, October). Empirical Models of Social Learning in a Large, Evolving Network. *PLOS ONE*, 11(10), e0160307. Retrieved 2022-04-01, from <https://doi.org/10.1371/journal.pone.0160307>

- journals.plos.org/plosone/article?id=10.1371/journal.pone.0160307 (Publisher: Public Library of Science) doi: 10.1371/journal.pone.0160307
- Bergstra, J., Bardenet, R., Bengio, Y., & Kégl, B. (2011). Algorithms for Hyper-Parameter Optimization. In *Advances in Neural Information Processing Systems* (Vol. 24). Curran Associates, Inc. Retrieved 2022-04-04, from shorturl.at/eu047
- Bhagat, S., Burke, M., Diuk, C., Edunov, S., & Filiz, I. O. (2016, February). *Three and a half degrees of separation - Meta Research*. Retrieved 2022-04-01, from <https://research.facebook.com/blog/2016/2/three-and-a-half-degrees-of-separation/>
- Bianconi, G., Darst, R. K., Iacoviacci, J., & Fortunato, S. (2014). Triadic closure as a basic generating mechanism of the structure of complex networks. *Physical review. E, Statistical, nonlinear, and soft matter physics*. doi: 10.1103/PhysRevE.90.042806
- Boxell, L., Gentzkow, M., & Shapiro, J. M. (2020, January). *Cross-Country Trends in Affective Polarization* (Working Paper No. 26669). National Bureau of Economic Research. Retrieved 2022-02-24, from <https://www.nber.org/papers/w26669> (Series: Working Paper Series) doi: 10.3386/w26669
- Boyd, R., Gintis, H., & Bowles, S. (2010, April). Coordinated Punishment of Defectors Sustains Cooperation and Can Proliferate When Rare. *Science*, 328(5978), 617–620. Retrieved 2022-03-30, from <https://www.science.org/doi/full/10.1126/science.1183665> (Publisher: American Association for the Advancement of Science) doi: 10.1126/science.1183665
- Broido, A. D., & Clauset, A. (2019, March). Scale-free networks are rare. *Nature Communications*, 10(1), 1017. Retrieved 2022-03-17, from <https://www.nature.com/articles/s41467-019-08746-5> (Number: 1 Publisher: Nature Publishing Group) doi: 10.1038/s41467-019-08746-5
- Bruch, E., & Atwell, J. (2015, May). Agent-Based Models in Empirical Social Research. *Sociological Methods & Research*, 44(2), 186–221. Retrieved 2022-04-05, from <https://doi.org/10.1177/0049124113506405> (Publisher: SAGE Publications Inc) doi: 10.1177/0049124113506405
- Carter, A. J., Lee, A. E., & Marshall, H. H. (2015). Phenotypic assortment in wild primate networks.
- Center, P. R. (2014a, June). *Political Polarization and Growing Partisan Antipathy*. Retrieved 2022-03-11, from shorturl.at/kuyMZ
- Center, P. R. (2014b, June). *Political Polarization in the American Public*. Retrieved 2022-03-11, from shorturl.at/qxCD3
- Chacoma, A., & Zanette, D. H. (2015, October). Opinion Formation by Social Influence: From Experiments to Modeling. *PLOS ONE*, 10(10), e0140406. Retrieved 2022-02-24, from <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0140406> (Publisher: Public Library of Science) doi: 10.1371/journal.pone.0140406

- Cowan, R., & Jonard, N. (2004). Network structure and the diffusion of knowledge. *Journal of economic Dynamics and Control*, 28(8), 1557–1575.
- Crandall, D., Cosley, D., Huttenlocher, D., Kleinberg, J., & Suri, S. (2008, August). Feedback effects between similarity and social influence in online communities. In *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining* (pp. 160–168). New York, NY, USA: Association for Computing Machinery. Retrieved 2022-04-08, from <https://doi.org/10.1145/1401890.1401914> doi: 10.1145/1401890.1401914
- Crooks, A. T., & Heppenstall, A. J. (2012). Introduction to agent-based modelling. In *Agent-based models of geographical systems* (pp. 85–105). Springer.
- Dakin, R., & Ryder, T. B. (2018, December). Dynamic network partnerships and social contagion drive cooperation. *Proceedings of the Royal Society B: Biological Sciences*, 285(1893), 20181973. Retrieved 2022-03-30, from <https://royalsocietypublishing.org/doi/full/10.1098/rspb.2018.1973> (Publisher: Royal Society) doi: 10.1098/rspb.2018.1973
- de Arruda, H. F., Cardoso, F. M., de Arruda, G. F., Hernández, A. R., da Fontoura Costa, L., & Moreno, Y. (2022). Modelling how social network algorithms can influence opinion polarization. *Information Sciences*, 588, 265–278.
- Del Vicario, M., Vivaldo, G., Bessi, A., Zollo, F., Scala, A., Caldarelli, G., & Quattrociocchi, W. (2016, December). Echo Chambers: Emotional Contagion and Group Polarization on Facebook. *Scientific Reports*, 6(1), 37825. Retrieved 2022-04-09, from <https://www.nature.com/articles/srep37825> (Number: 1 Publisher: Nature Publishing Group) doi: 10.1038/srep37825
- Dimock, M., & Wike, R. (2020, November). *America is exceptional in the nature of its political divide*. Retrieved 2022-03-11, from shorturl.at/wKMR4
- Emiliano, P. C., Vivanco, M. J., & De Menezes, F. S. (2014). Information criteria: How do they behave in different models? *Computational Statistics & Data Analysis*, 69, 141–153.
- Epstein, J. M. (1999). Agent-based computational models and generative social science. *Complexity*, 4(5), 41–60.
- Flache, A. (2018, December). Between Monoculture and Cultural Polarization: Agent-based Models of the Interplay of Social Influence and Cultural Diversity. *Journal of Archaeological Method and Theory*, 25(4), 996–1023. Retrieved 2022-01-26, from <https://doi.org/10.1007/s10816-018-9391-1> doi: 10.1007/s10816-018-9391-1
- Flache, A., & Macy, M. W. (2006, April). Why more contact may increase cultural polarization. *arXiv:physics/0604196*. Retrieved 2022-03-20, from <http://arxiv.org/abs/physics/0604196> (arXiv: physics/0604196)

- Flache, A., & Macy, M. W. (2011, January). Small Worlds and Cultural Polarization. *Journal of Mathematical Sociology*, 35(1-2), 146–176. doi: 10.1080/0022250X.2010.532261
- Flache, A., Mäs, M., Feliciani, T., Chattoe-Brown, E., Deffuant, G., Huet, S., & Lorenz, J. (2017). Models of Social Influence: Towards the Next Frontiers. *Journal of Artificial Societies and Social Simulation*, 20(4), 2.
- Fogarty, L., Ammar, M., Holding, T., Powell, A., & Kandler, A. (2022, March). Ten simple rules for principled simulation modelling. *PLOS Computational Biology*, 18(3), e1009917. Retrieved 2022-04-05, from <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1009917> (Publisher: Public Library of Science) doi: 10.1371/journal.pcbi.1009917
- Friedkin, N. E., & Johnsen, E. C. (1990, January). Social influence and opinions. *The Journal of Mathematical Sociology*, 15(3-4), 193–206. Retrieved 2022-01-26, from <https://doi.org/10.1080/0022250X.1990.9990069> (Publisher: Routledge _eprint: <https://doi.org/10.1080/0022250X.1990.9990069>) doi: 10.1080/0022250X.1990.9990069
- Fuglede, B., & Topsøe, F. (2004, June). Jensen-Shannon divergence and Hilbert space embedding. In *International Symposium on Information Theory, 2004. ISIT 2004. Proceedings*. (pp. 31–). doi: 10.1109/ISIT.2004.1365067
- Galesic, M., Olsson, H., Dalege, J., van der Does, T., & Stein, D. L. (2021, March). Integrating social and cognitive aspects of belief dynamics: towards a unifying framework. *Journal of The Royal Society Interface*, 18(176), 20200857. Retrieved 2022-03-11, from <https://royalsocietypublishing.org/doi/10.1098/rsif.2020.0857> (Publisher: Royal Society) doi: 10.1098/rsif.2020.0857
- Geschke, D., Lorenz, J., & Holtz, P. (2019). The triple-filter bubble: Using agent-based modelling to test a meta-theoretical framework for the emergence of filter bubbles and echo chambers. *British Journal of Social Psychology*, 58(1), 129–149.
- Ghaderi, J., & Srikant, R. (2014, December). Opinion dynamics in social networks with stubborn agents: Equilibrium and convergence rate. *Automatica*, 50(12), 3209–3215. Retrieved 2022-03-31, from <https://www.sciencedirect.com/science/article/pii/S0005109814004154> doi: 10.1016/j.automatica.2014.10.034
- Holme, P., & Newman, M. E. J. (2006, November). Nonequilibrium phase transition in the coevolution of networks and opinions. *Physical Review E*, 74(5), 056108. Retrieved 2022-03-17, from <https://link.aps.org/doi/10.1103/PhysRevE.74.056108> (Publisher: American Physical Society) doi: 10.1103/PhysRevE.74.056108
- Hutter, F., Hoos, H., & Leyton-Brown, K. (2014). An efficient approach for assessing hyperparameter importance. In *International conference on machine learning* (pp. 754–762).

- Ilany, A., & Akçay, E. (2016, June). Social inheritance can explain the structure of animal social networks. *Nature Communications*, 7(1), 12084. Retrieved 2022-01-26, from <https://www.nature.com/articles/ncomms12084> (Bandiera_abtest: a Cc_license_type: cc_by Cg_type: Nature Research Journals Number: 1 Primary_atype: Research Publisher: Nature Publishing Group Subject_term: Social evolution;Theoretical ecology Subject_term_id: social-evolution;theoretical-ecology) doi: 10.1038/ncomms12084
- Iyengar, S., Lelkes, Y., Levendusky, M., Malhotra, N., & Westwood, S. (2019). The Origins and Consequences of Affective Polarization in the United States. *Annual Review of Political Science*. doi: 10.1146/ANNUREV-POLISCI-051117-073034
- Jackson, M., & Rogers, B. (2004, June). Search and the Strategic Formation of Large Networks: When and Why do We See Power Laws and Small Worlds?
- Jackson, M., & Rogers, B. W. (2007). Meeting Strangers and Friends of Friends: How Random are Social Networks? doi: 10.1257/AER.97.3.890
- Jalili, M. (2015, November). Coevolution of opinion formation and network dynamics in complex networked systems. In *2015 International Conference on Information Society (i-Society)* (pp. 79–83). doi: 10.1109/i-Society.2015.7366863
- Jussim, L., & Osgood, D. W. (1989). Influence and Similarity Among Friends: An Integrative Model Applied to Incarcerated Adolescents. *Social Psychology Quarterly*, 52(2), 98–112. Retrieved 2022-04-08, from <https://www.jstor.org/stable/2786910> (Publisher: [Sage Publications, Inc., American Sociological Association]) doi: 10.2307/2786910
- Kerr, C. C., Stuart, R. M., Mistry, D., Abeyseuriya, R. G., Rosenfeld, K., Hart, G. R., ... others (2021). Covasim: an agent-based model of covid-19 dynamics and interventions. *PLOS Computational Biology*, 17(7), e1009149.
- Kossinets, G., & Watts, D. (2009, September). Origins of Homophily in an Evolving Social Network. *American Journal of Sociology*, 115(2), 405–450. Retrieved 2022-01-29, from <https://www.journals.uchicago.edu/doi/abs/10.1086/599247> (Publisher: The University of Chicago Press) doi: 10.1086/599247
- Kozma, B., & Barrat, A. (2008). Consensus formation on coevolving networks: groups' formation and structure. *Journal of Physics A: Mathematical and Theoretical*, 41(22), 224020.
- Krivorotko, O., Sosnovskaia, M., Vashchenko, I., Kerr, C., & Lesnic, D. (2022). Agent-based modeling of covid-19 outbreaks for new york state and uk: Parameter identification algorithm. *Infectious Disease Modelling*, 7(1), 30–44.

- Kurzban, R., Burton-Chellew, M. N., & West, S. A. (2015). The evolution of altruism in humans. *Annual review of psychology*, 66, 575–599.
- Levin, S. A., Milner, H. V., & Perrings, C. (2021, December). The dynamics of political polarization. *Proceedings of the National Academy of Sciences*, 118(50), e2116950118. Retrieved 2022-04-08, from <https://www.pnas.org/doi/10.1073/pnas.2116950118> (Publisher: Proceedings of the National Academy of Sciences) doi: 10.1073/pnas.2116950118
- Lev-On, A., & Manin, B. (2009, October). *Happy Accidents: Deliberation and Online Exposure to Opposing Views* (SSRN Scholarly Paper No. 1481869). Rochester, NY: Social Science Research Network. Retrieved 2022-04-09, from <https://papers.ssrn.com/abstract=1481869>
- Levy, R. (2021, March). Social Media, News Consumption, and Polarization: Evidence from a Field Experiment. *American Economic Review*, 111(3), 831–870. Retrieved 2022-04-10, from <https://www.aeaweb.org/articles?id=10.1257/aer.20191777> doi: 10.1257/aer.20191777
- Lusseau, D., Schneider, K., Boisseau, O. J., Haase, P., Slooten, E., & Dawson, S. M. (2003). The bottlenose dolphin community of Doubtful Sound features a large proportion of long-lasting associations - Can geographic isolation explain this unique trait? *Behavioral Ecology and Sociobiology*, 54, 396–405. Retrieved 2022-04-07, from shorturl.at/hiqIL (Publisher: Springer) doi: 10.1007/s00265-003-0651-y
- Macy, M. W., Ma, M., Tabin, D. R., Gao, J., & Szymanski, B. K. (2021). Polarization and tipping points. *Proceedings of the National Academy of Sciences*, 118(50).
- Mäs, M. (2019). Challenges to simulation validation in the social sciences. a critical rationalist perspective. In *Computer simulation validation* (pp. 857–879). Springer.
- Mäs, M., Flache, A., & Kitts, J. A. (2014). Cultural integration and differentiation in groups and organizations. In *Perspectives on culture and agent-based simulations* (pp. 71–90). Springer.
- Matsumura, T., Iwasaki, K., & Shudo, K. (2018, January). Average Path Length Estimation of Social Networks by Random Walk. In *2018 IEEE International Conference on Big Data and Smart Computing (BigComp)* (pp. 611–614). (ISSN: 2375-9356) doi: 10.1109/BigComp.2018.00107
- McCoy, J., Rahman, T., & Somer, M. (2018, January). Polarization and the Global Crisis of Democracy: Common Patterns, Dynamics, and Pernicious Consequences for Democratic Polities. *American Behavioral Scientist*, 62(1), 16–42. Retrieved 2022-04-06, from <https://doi.org/10.1177/0002764218759576> (Publisher: SAGE Publications Inc) doi: 10.1177/0002764218759576
- McPherson, M., Smith-Lovin, L., & Cook, J. M. (2001). Birds of a Feather: Homophily in Social Networks. *Annual Review of Sociology*, 27(1), 415–444. Retrieved 2022-03-14, from <https://doi.org/10.1146/annurev.soc.27.1.415> (.eprint: <https://doi.org/10.1146/annurev.soc.27.1.415>)

doi: 10.1146/annurev.soc.27.1.415

- Melamed, D., & Simpson, B. (2016, March). Strong ties promote the evolution of cooperation in dynamic networks. *Social Networks*, 45, 32–44. Retrieved 2022-03-31, from <https://www.sciencedirect.com/science/article/pii/S0378873315300605> doi: 10.1016/j.socnet.2015.11.001
- Mäs, M., Flache, A., & Helbing, D. (2010, October). Individualization as Driving Force of Clustering Phenomena in Humans. *PLOS Computational Biology*, 6(10), e1000959. Retrieved 2022-01-26, from <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1000959> (Publisher: Public Library of Science) doi: 10.1371/journal.pcbi.1000959
- Newman, M. E. J. (2006, September). Finding community structure in networks using the eigenvectors of matrices. *Physical Review E*, 74(3), 036104. Retrieved 2022-04-06, from <http://arxiv.org/abs/physics/0605087> (arXiv: physics/0605087) doi: 10.1103/PhysRevE.74.036104
- Noel, H., & Nyhan, B. (2011). The “unfriending” problem: The consequences of homophily in friendship retention for causal estimates of social influence. *Social Networks*, 33(3), 211–218.
- Noorazar, H., Vixie, K. R., Talebanpour, A., & Hu, Y. (2020, July). From classical to modern opinion dynamics. *International Journal of Modern Physics C*, 31(07), 2050101. Retrieved 2022-02-24, from <https://www.worldscientific.com/doi/abs/10.1142/S0129183120501016> (Publisher: World Scientific Publishing Co.) doi: 10.1142/S0129183120501016
- Page, S. (2010). Diversity and complexity. In *Diversity and complexity*. Princeton University Press.
- Peixoto, T. P. (2022, January). Disentangling homophily, community structure and triadic closure in networks. *Physical Review X*, 12(1), 011004. Retrieved 2022-01-26, from <http://arxiv.org/abs/2101.02510> (arXiv: 2101.02510) doi: 10.1103/PhysRevX.12.011004
- Pennycook, G., & Rand, D. G. (2019, July). Lazy, not biased: Susceptibility to partisan fake news is better explained by lack of reasoning than by motivated reasoning. *Cognition*, 188, 39–50. Retrieved 2022-04-01, from <https://www.sciencedirect.com/science/article/pii/S001002771830163X> doi: 10.1016/j.cognition.2018.06.011
- Pepper, J. W., & Smuts, B. B. (2002, August). A Mechanism for the Evolution of Altruism among Nonkin: Positive Assortment through Environmental Feedback. *The American Naturalist*, 160(2), 205–213. Retrieved 2022-03-30, from <https://www.journals.uchicago.edu/doi/full/10.1086/341018> (Publisher: The University of Chicago Press) doi: 10.1086/341018
- Poli, R. (2013). A note on the difference between complicated and complex social systems. *Cadmus*.
- Rossi, R., & Ahmed, N. (2015, March). The Network Data Repository with Interactive Graph Analytics and Visualization. In *Twenty-Ninth AAAI Conference on Artificial Intelligence*. Retrieved 2022-03-31, from <https://www.aaai.org/ocs/index.php/AAAI/AAAI15/paper/view/9553>

- Rozemberczki, B., Davies, R., Sarkar, R., & Sutton, C. (2019). Gemsec: Graph embedding with self clustering. In *Proceedings of the 2019 ieee/acm international conference on advances in social networks analysis and mining* (pp. 65–72).
- Santos, F. C., Pacheco, J. M., & Lenaerts, T. (2006, October). Cooperation Prevails When Individuals Adjust Their Social Ties. *PLOS Computational Biology*, 2(10), e140. Retrieved 2022-02-24, from <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.0020140> (Publisher: Public Library of Science) doi: 10.1371/journal.pcbi.0020140
- Sasahara, K., Chen, W., Peng, H., Ciampaglia, G. L., Flammini, A., & Menczer, F. (2021, May). Social influence and unfollowing accelerate the emergence of echo chambers. *Journal of Computational Social Science*, 4(1), 381–402. Retrieved 2022-01-26, from <https://doi.org/10.1007/s42001-020-00084-7> doi: 10.1007/s42001-020-00084-7
- Schelling, T. C. (1971). Dynamic models of segregation. *Journal of mathematical sociology*, 1(2), 143–186.
- Shi, F., Shi, Y., Dokshin, F. A., Evans, J. A., & Macy, M. W. (2017, April). Millions of online book co-purchases reveal partisan differences in the consumption of science. *Nature Human Behaviour*, 1(4), 1–9. Retrieved 2022-04-07, from <https://www.nature.com/articles/s41562-017-0079> (Number: 4 Publisher: Nature Publishing Group) doi: 10.1038/s41562-017-0079
- Sîrbu, A., Loreto, V., Servedio, V. D., & Tria, F. (2017). Opinion dynamics: models, extensions and external effects. In *Participatory sensing, opinions and collective awareness* (pp. 363–401). Springer.
- Smaldino, P. E. (2016, January). Models Are Stupid, and We Need More of Them.. doi: 10.4324/9781315173726-14
- Smaldino, P. E. (2019). Social identity and cooperation in cultural evolution. *Behavioural Processes*, 161, 108–116.
- Smaldino, P. E. (2020, July). How to Translate a Verbal Theory Into a Formal Model. *Social Psychology*, 51(4), 207–218. Retrieved 2022-04-06, from <https://econtent.hogrefe.com/doi/10.1027/1864-9335/a000425> (Publisher: Hogrefe Publishing) doi: 10.1027/1864-9335/a000425
- Smaldino, P. E. (2022, January). *Models of Identity Signaling* (Tech. Rep.). PsyArXiv. Retrieved 2022-01-10, from <https://psyarxiv.com/xv9j8/> (type: article) doi: 10.31234/osf.io/xv9j8
- Somer, M., & McCoy, J. (2018, January). Déjà vu? Polarization and Endangered Democracies in the 21st Century. *American Behavioral Scientist*, 62(1), 3–15. Retrieved 2022-04-06, from <https://doi.org/10.1177/0002764218760371> (Publisher: SAGE Publications Inc) doi: 10.1177/0002764218760371
- Spears, R. (2021). Social Influence and Group Identity. *Annual Review of Psychology*, 72(1), 367–390. Retrieved 2022-01-26, from <https://doi.org/10.1146/annurev-psych-070620-111818> (eprint: <https://doi.org/10.1146/annurev-psych-070620-111818>) doi: 10.1146/annurev-psych-070620-111818

- Strömbäck, J., Djerf-Pierre, M., & Shehata, A. (2013, December). The Dynamics of Political Interest and News Media Consumption: A Longitudinal Perspective. *International Journal of Public Opinion Research*, 25(4), 414–435. Retrieved 2022-04-01, from <https://doi.org/10.1093/ijpor/eds018> doi: 10.1093/ijpor/eds018
- Su, W., Chen, G., & Hong, Y. (2017, November). Noise leads to quasi-consensus of Hegselmann–Krause opinion dynamics. *Automatica*, 85, 448–454. Retrieved 2022-03-31, from <https://www.sciencedirect.com/science/article/pii/S0005109817304296> doi: 10.1016/j.automatica.2017.08.008
- Takács, K., Flache, A., & Maes, M. (2014, June). *Is There Negative Social Influence? Disentangling Effects of Dissimilarity and Disliking on Opinion Shifts* (SSRN Scholarly Paper No. ID 2445649). Rochester, NY: Social Science Research Network. Retrieved 2022-02-24, from <https://papers.ssrn.com/abstract=2445649> doi: 10.2139/ssrn.2445649
- Taylor, C. E., Mantzaris, A. V., & Garibay, I. (2018, December). Exploring How Homophily and Accessibility Can Facilitate Polarization in Social Networks. *Information*, 9(12), 325. Retrieved 2022-01-26, from <https://www.mdpi.com/2078-2489/9/12/325> (Number: 12 Publisher: Multidisciplinary Digital Publishing Institute) doi: 10.3390/info9120325
- Tsai, W.-H. S., Tao, W., Chuan, C.-H., & Hong, C. (2020, March). Echo chambers and social mediators in public advocacy issue networks. *Public Relations Review*, 46(1), 101882. Retrieved 2022-03-29, from <https://www.sciencedirect.com/science/article/pii/S036381120300035> doi: 10.1016/j.pubrev.2020.101882
- Turner, M. A., & Smaldino, P. E. (2018, November). Paths to Polarization: How Extreme Views, Miscommunication, and Random Chance Drive Opinion Dynamics. *Complexity*, 2018, e2740959. Retrieved 2022-03-10, from <https://www.hindawi.com/journals/complexity/2018/2740959/> (Publisher: Hindawi) doi: 10.1155/2018/2740959
- Vrieze, S. I. (2012, June). Model selection and psychological theory: A discussion of the differences between the Akaike Information Criterion (AIC) and the Bayesian Information Criterion (BIC). *Psychological Methods*, 17(2), 228–243. Retrieved 2019-05-30, from <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3366160/> doi: 10.1037/a0027127
- Watts, D. J. (1999, September). Networks, Dynamics, and the Small-World Phenomenon. *American Journal of Sociology*, 105(2), 493–527. Retrieved 2022-03-17, from <https://www.journals.uchicago.edu/doi/abs/10.1086/210318> (Publisher: The University of Chicago Press) doi: 10.1086/210318
- Watts, D. J., & Strogatz, S. H. (1998, June). Collective dynamics of ‘small-world’ networks. *Nature*, 393(6684), 440–442. Retrieved 2022-02-02, from <https://www.nature.com/articles/30918>. (Number: 6684 Publisher: Nature Publishing Group) doi: 10.1038/30918

- Wilson, A. E., Parker, V. A., & Feinberg, M. (2020, August). Polarization in the contemporary political and media landscape. *Current Opinion in Behavioral Sciences*, 34, 223–228. Retrieved 2022-04-06, from <https://www.sciencedirect.com/science/article/pii/S2352154620301078> doi: 10.1016/j.cobeha.2020.07.005
- Winter, F., & Kataria, M. (2020, May). You are who your friends are?: An experiment on homophily in trustworthiness among friends. *Rationality and Society*, 32(2), 223–251. Retrieved 2022-01-26, from <https://doi.org/10.1177/1043463120919380> (Publisher: SAGE Publications Ltd) doi: 10.1177/1043463120919380
- Yildiz, E., Ozdaglar, A., Acemoglu, D., Saberi, A., & Scaglione, A. (2013, December). Binary Opinion Dynamics with Stubborn Agents. *ACM Transactions on Economics and Computation*, 1(4), 19:1–19:30. Retrieved 2022-03-31, from <https://doi.org/10.1145/2538508> doi: 10.1145/2538508
- Zachary, W. W. (1977, December). An Information Flow Model for Conflict and Fission in Small Groups. *Journal of Anthropological Research*, 33(4), 452–473. Retrieved 2022-04-06, from <https://www.journals.uchicago.edu/doi/abs/10.1086/jar.33.4.3629752> (Publisher: The University of Chicago Press) doi: 10.1086/jar.33.4.3629752

A Appendix

A.1 Supplementary material for hyperparameter optimization

A.1.1 Diagnostic plots

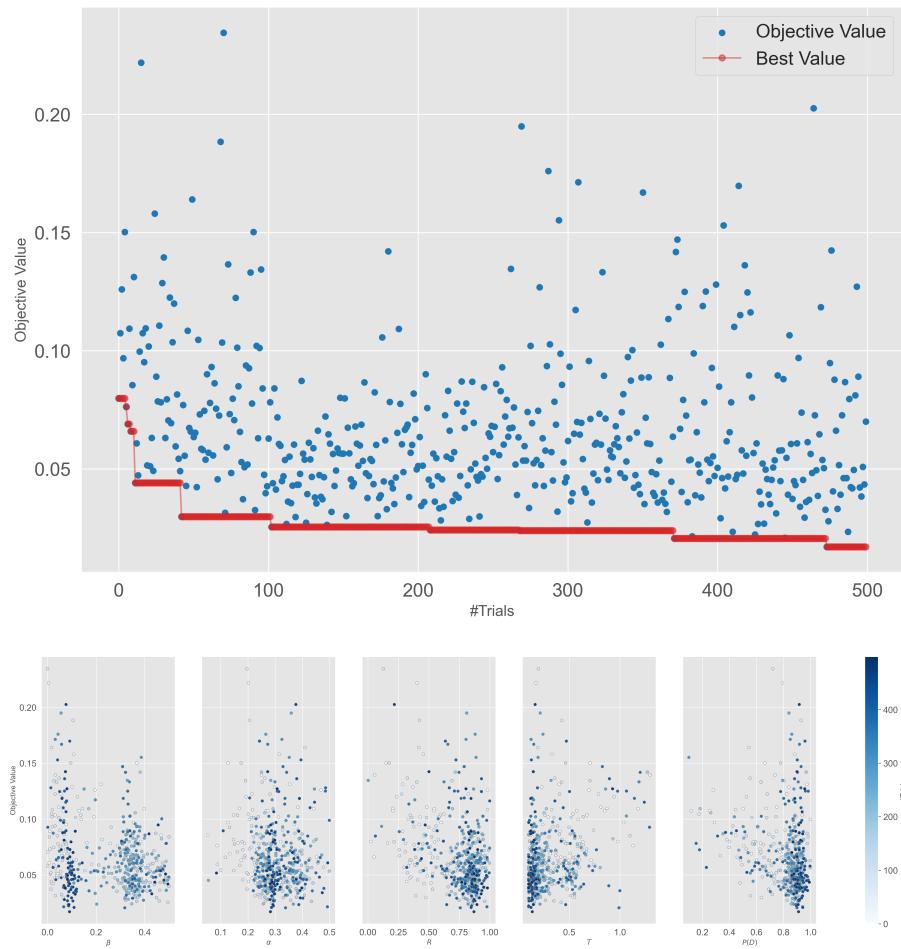


Figure A.1: Diagnostic plots of optimization for the Dolphin Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.

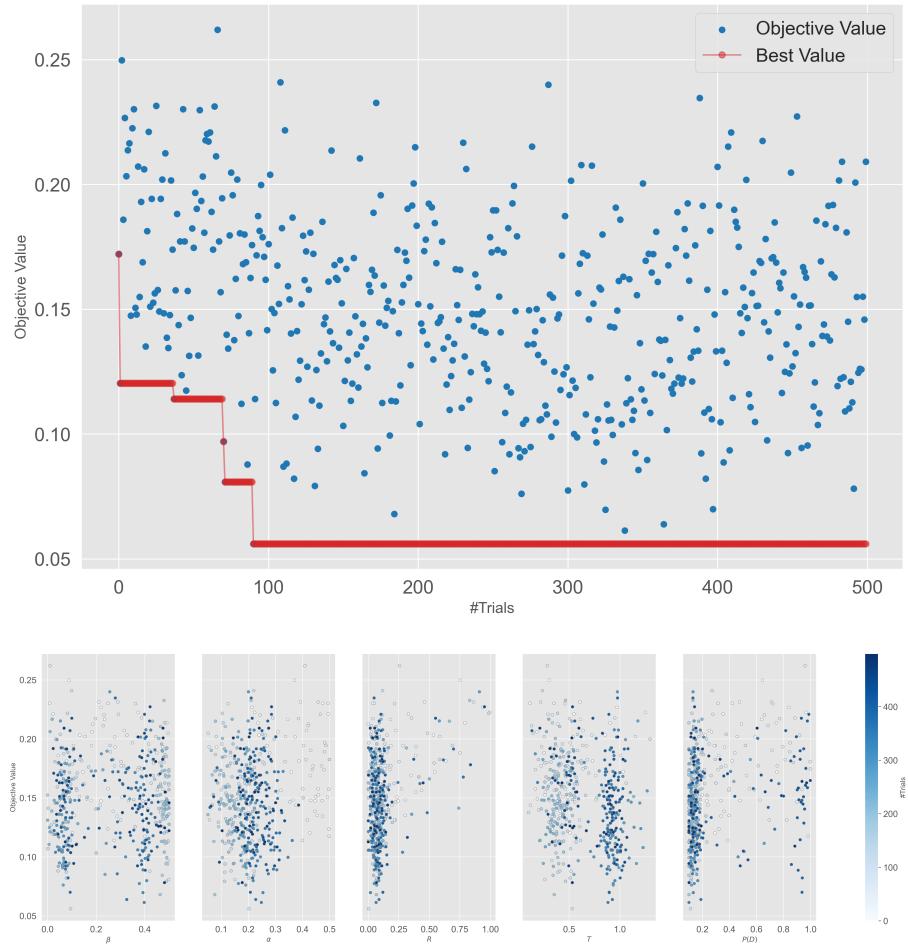


Figure A.2: Diagnostic plots of optimization for the Karate Club Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.

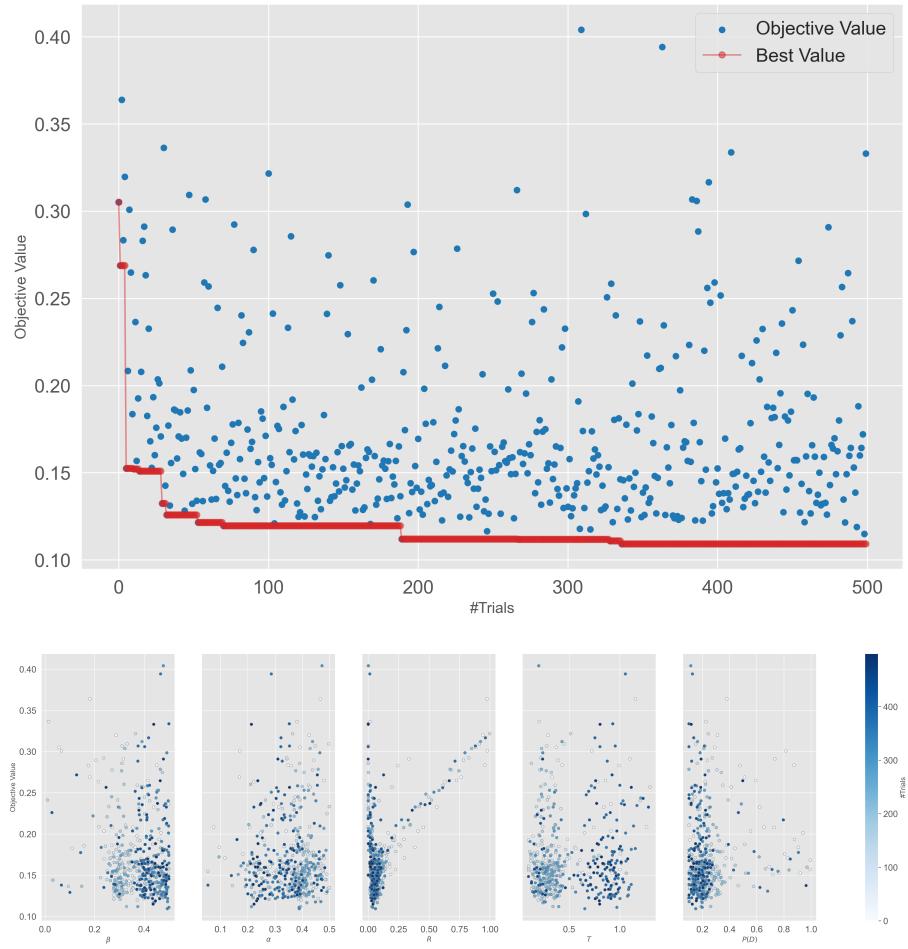


Figure A.3: Diagnostic plots of optimization for the Citation Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.

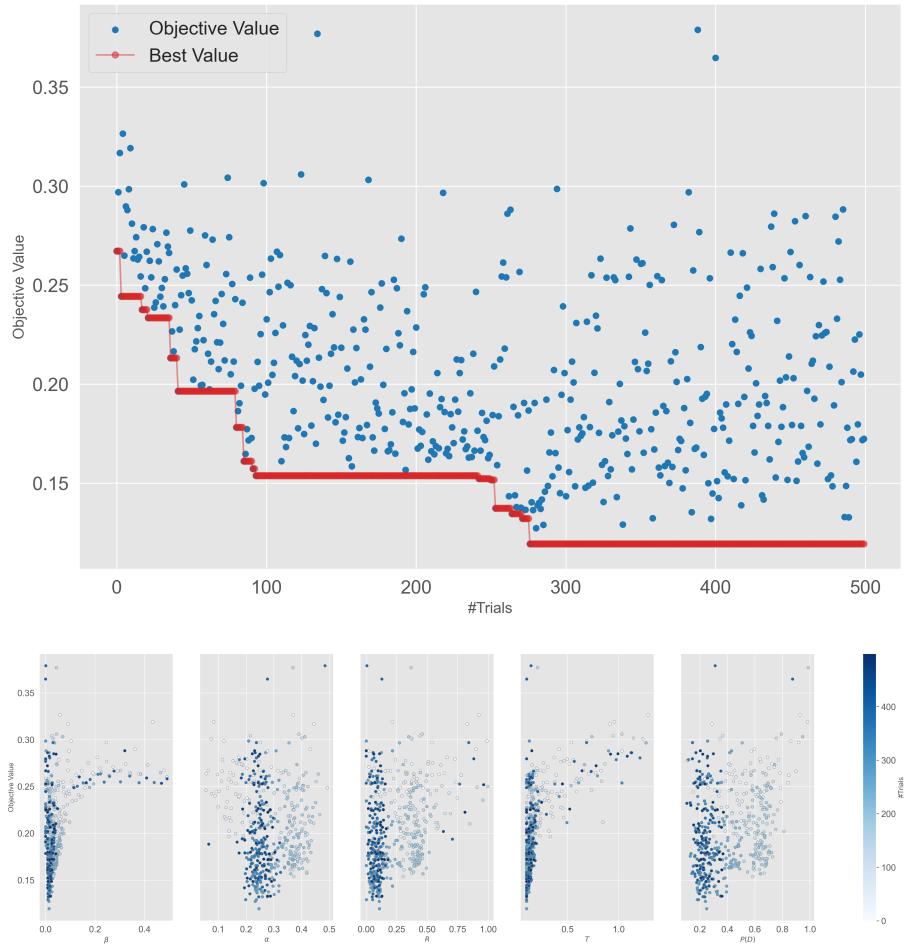


Figure A.4: Diagnostic plots of optimization for the Political Blogs Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.

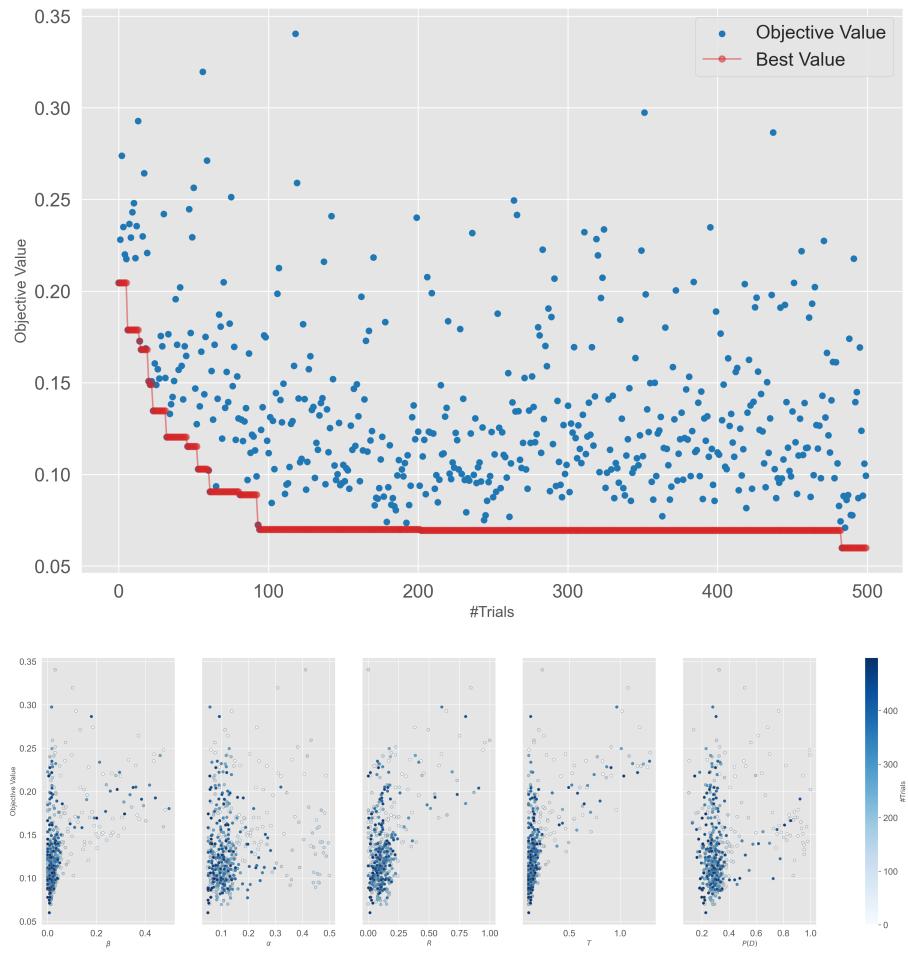


Figure A.5: Diagnostic plots of optimization for the Politicians Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.



Figure A.6: Diagnostic plots of optimization for the Political Books Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.



Figure A.7: Diagnostic plots of optimization for the TV Shows Network. The plot on the first row shows the optimization history, with x-axis showing the trial number and the y-axis showing the objective value of that trial. Red line shows the best value, blue values shows the objective value of each trial. Second row shows the search pattern for each parameter. The different columns show different model parameters. X-axis for all columns show the parameter's value. Y-axis shows the objective value. Dots are colored as a gradient by the trial number with darker colors indicating later trials.

A.1.2 Metrics of model performance

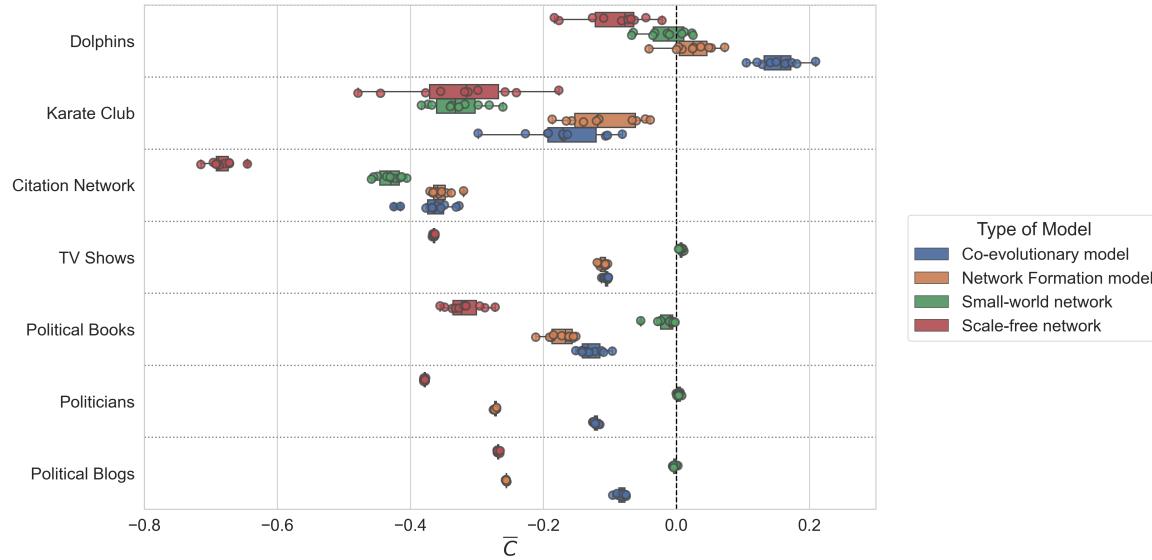


Figure A.8: Difference between models and target regarding average clustering coefficient. X-axis shows difference between the average clustering coefficient between the generated and the empirical networks, \bar{C} . Y-axis shows the different empirical networks considered. Colors show different algorithms for generating networks. Vertical dashed line shows the value 0.0, which is the point of no difference between the generated and the empirical network. Dots show the result from individual simulations, while boxplots show the median and quartiles of the distribution of values.

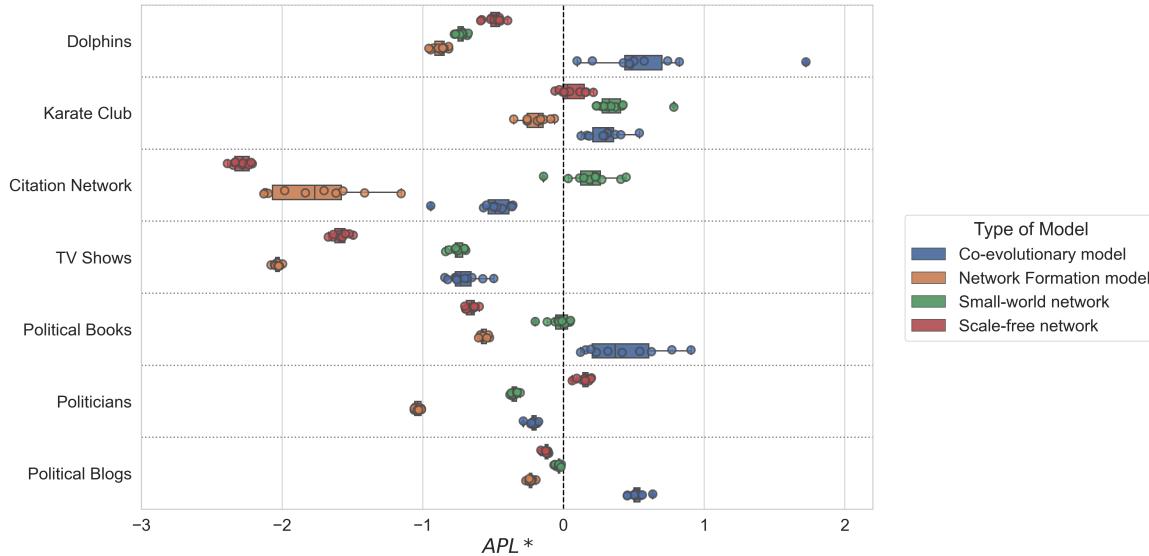


Figure A.9: Difference between models and target regarding average path length. X-axis shows difference between the estimated average path length between the generated and the empirical networks, APL^* . Y-axis shows the different empirical networks considered. Colors show different algorithms for generating networks. Vertical dashed line shows the value 0.0, which is the point of no difference between the generated and the empirical network. Dots show the result from individual simulations, while boxplots show the median and quartiles of the distribution of values.

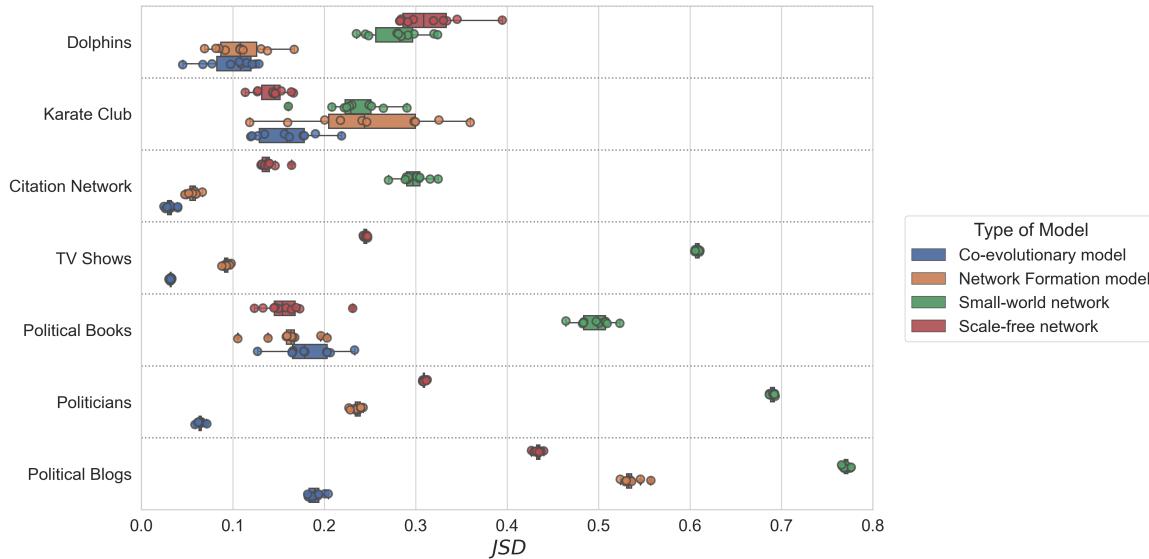


Figure A.10: Difference in degree distributions of generated and empirical networks. X-axis shows the Jensen Shannon Divergence, JSD , of the degree distribution between the generated and the target network. Colors show different algorithms for generating networks. Y-axis shows the different empirical networks considered. Dots show the result from individual simulations, while boxplots show the median and quartiles of the distribution of values.

A.2 Supplementary information for co-evolution's effect on opinion dynamics

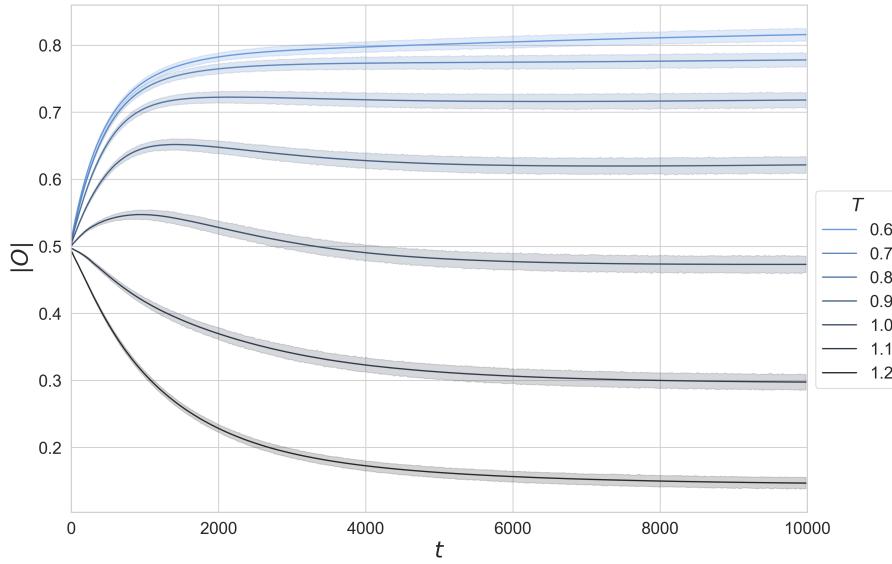


Figure A.11: T 's effect on polarization over time. The x-axis shows the time-steps and the y-axis shows the absolute value of the opinions of agents. Colors indicate different values of T . Lines indicate the mean value of each time-step, aggregated over all parameters excluding T . Shaded regions are the 95% confidence intervals.

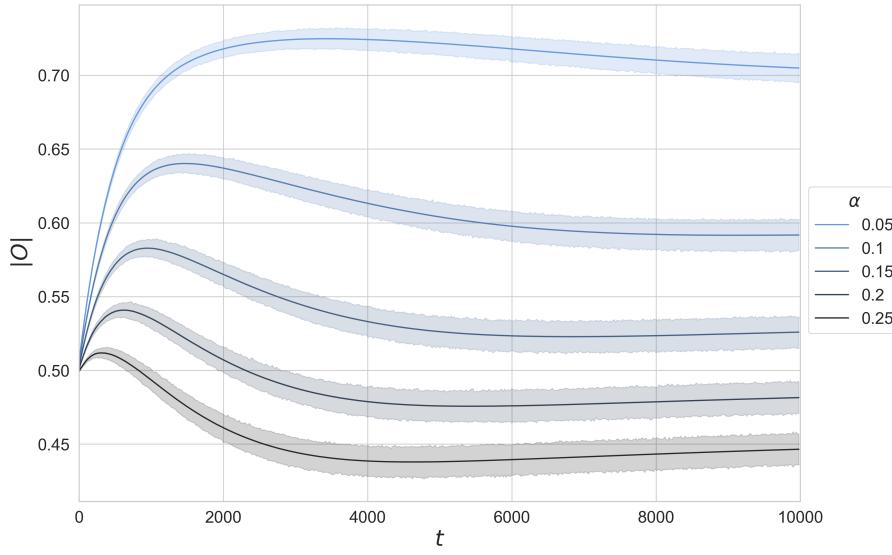


Figure A.12: α 's effect on polarization over time. The x-axis shows the time-steps and the y-axis shows the absolute value of the opinions of agents. Colors indicate different values of α . Lines indicate the mean value of each time-step, aggregated over all parameters excluding α . Shaded regions are the 95% confidence intervals.

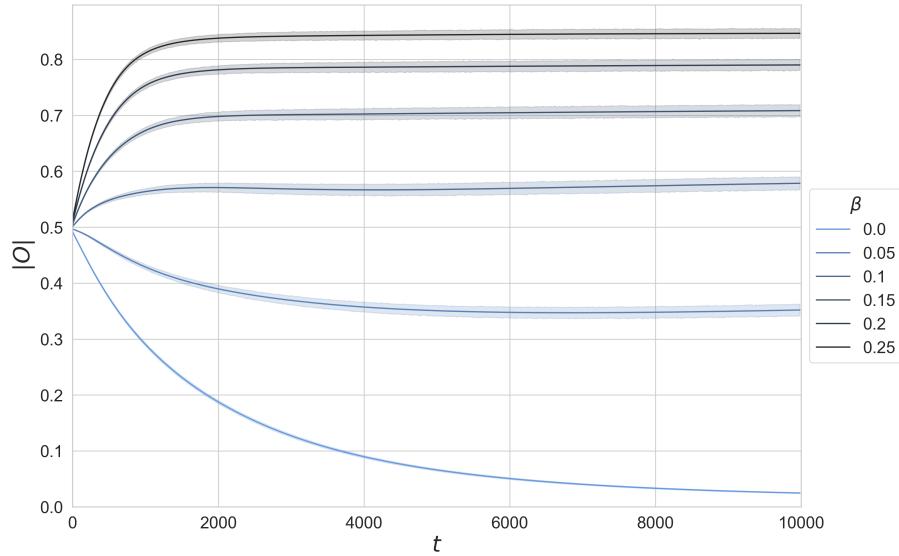


Figure A.13: β 's effect on polarization over time. The x-axis shows the time-step and the y-axis shows the absolute value of the opinions of agents. Colors indicate different values of β . Lines indicate the mean value of each time-step, aggregated over all parameters excluding β . Shaded regions are the 95% confidence intervals.

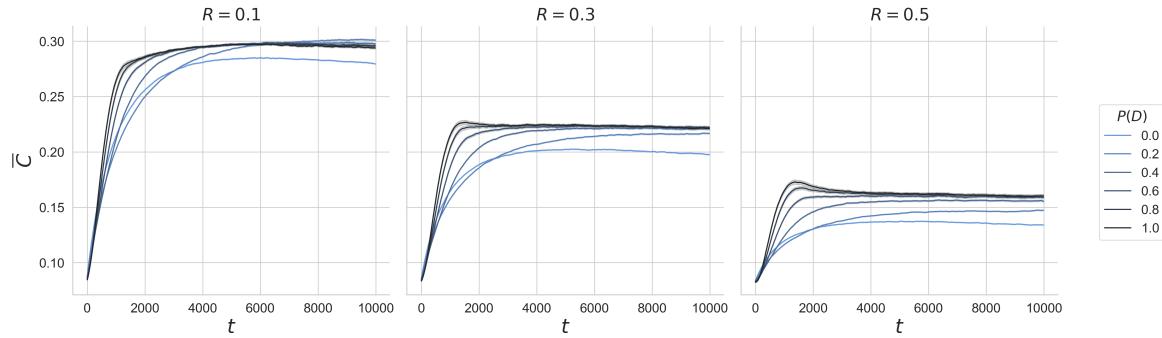


Figure A.14: Average clustering coefficient over time. X-axis shows time-steps, and the y-axis shows the average clustering coefficient of the network. Colors indicate different probabilities for negative tie deletion. Lines show the average value over all other parameters than $P(D)$ and R . Shaded areas indicate 95% confidence intervals. Each panel shows different values of R .

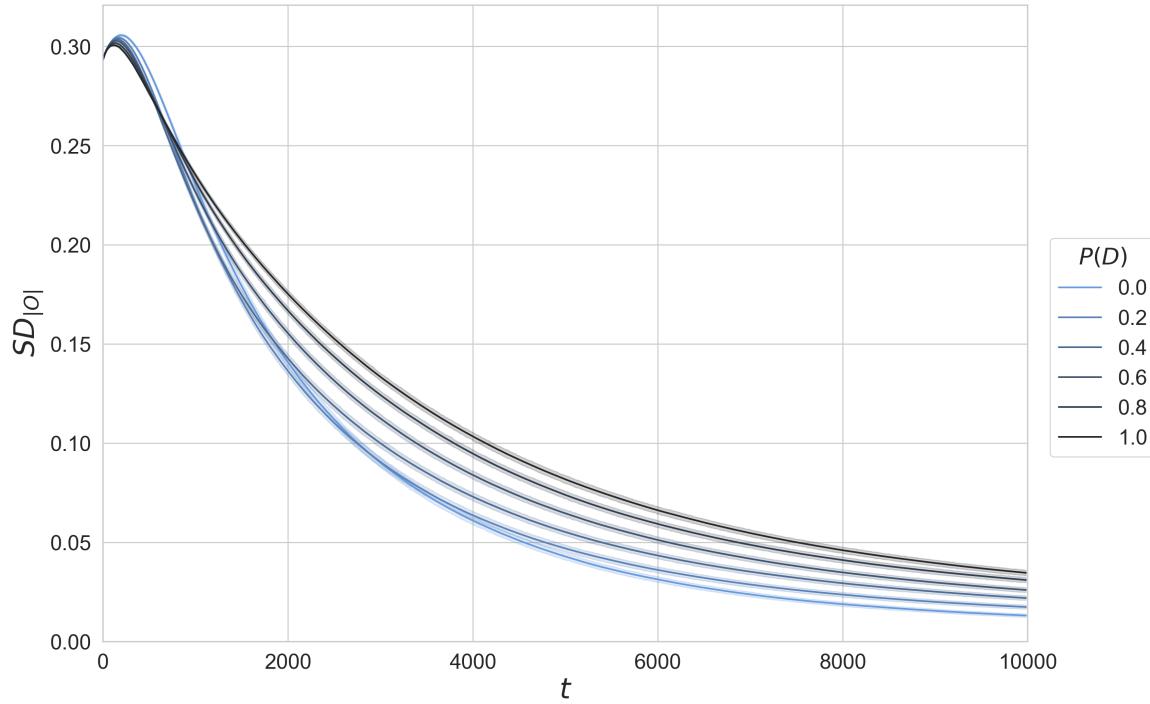


Figure A.15: $P(D)$'s effect on variance of opinions over time. The x-axis shows the time-steps and the y-axis shows the standard deviation of the absolute value of opinions. Colors indicate different probabilities of tie-deletion of dissimilar agents. Lines indicate the mean value of each time-step, aggregated over all parameters excluding the probability of tie-deletion. Shaded regions are the 95% confidence intervals.

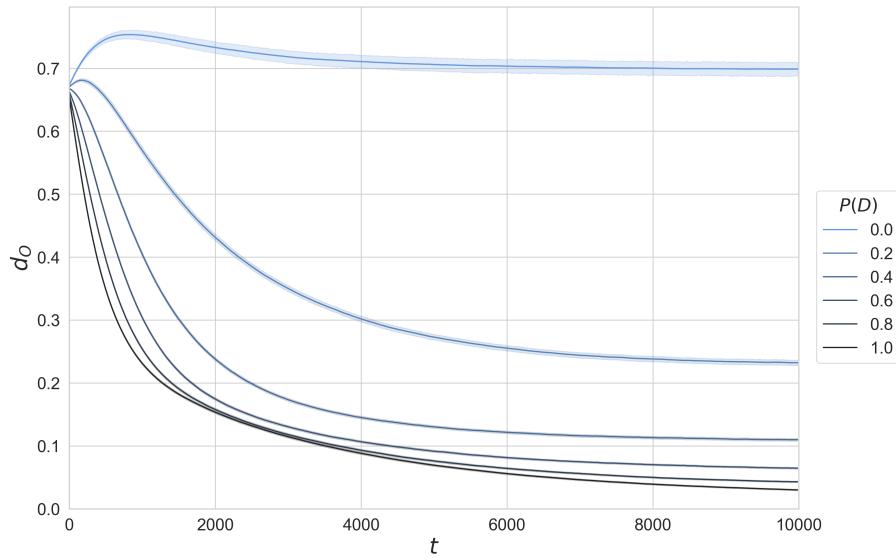


Figure A.16: Average distance to neighbors' opinions. Different colors show different values of the probability of tie-deletion ($P(D)$). The x-axis shows the time-steps and the y-axis shows the mean distance to neighbors' opinions (d_O).

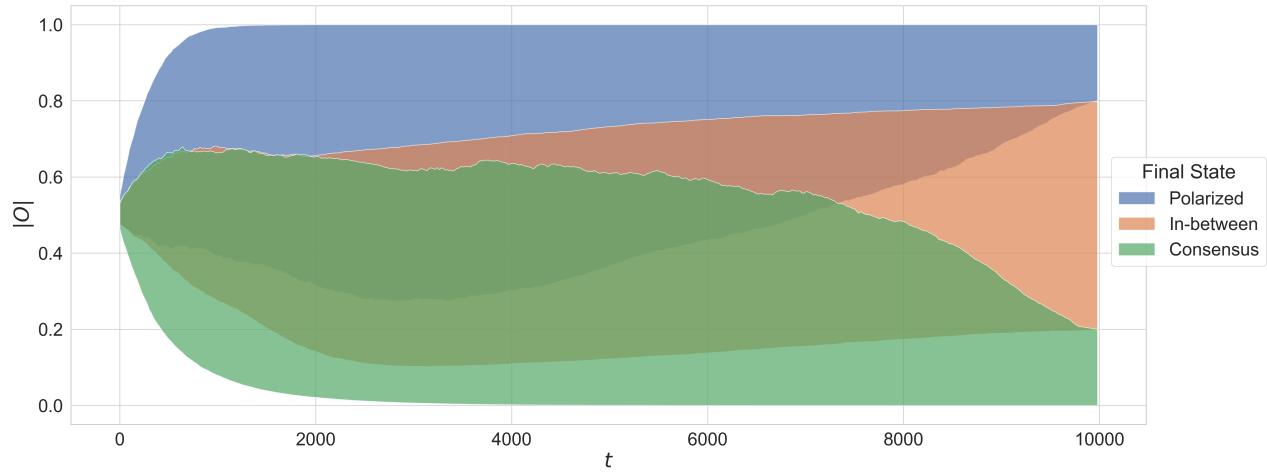


Figure A.17: Polarized, consensus and in-between conditions. X-axis shows time-steps, and the y-axis shows the average of the absolute value of opinions, $|O|$. Colors indicate whether simulations ended in consensus, polarization or something in-between. Different simulations were categorized into these different categories depending $|O|$ at the 10.000th time-step. Conditions were labelled consensus if $|O| \leq 0.2$, in-between if $0.2 < |O| < 0.8$ and polarized if $0.8 \leq |O|$. Areas on the plot were made by coloring the area between minimum and the maximum value for every time-step.

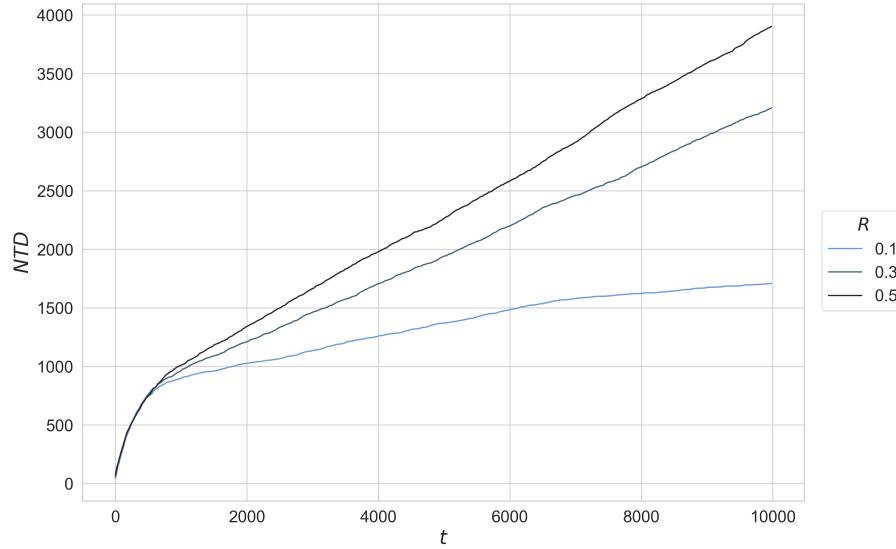


Figure A.18: Negative ties deleted in representative simulations over time. X-axis shows time-steps, and the y-axis shows the cumulative frequency of negative ties deleted. Colors indicate different levels of randomness. Lines show single simulations over time which were simulated using the same random seed. These simulations had parameter values of $T = 0.8$, $\alpha = 0.15$, $\beta = 0.1$ and $P(D) = 1$.

A.2.1 Negative tie-deletion

The more random new ties are, the more ties will be deleted on average. Interestingly, higher probabilities of tie-deletion almost always result in a lower frequency of deleted ties over time (see Appendix A.19).

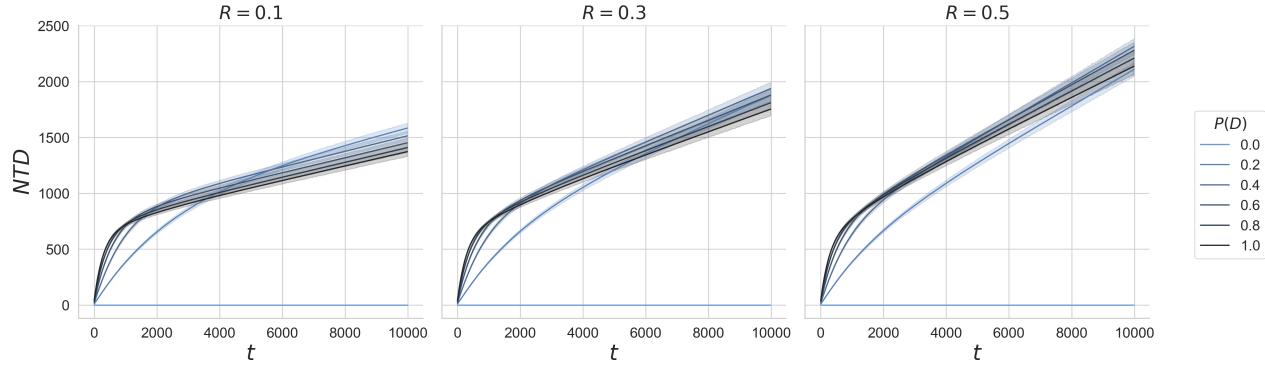


Figure A.19: Cumulative frequency of negative ties deleted. The x-axis shows the time-steps and the y-axis shows the average cumulative frequency of deleted ties. Colors indicate different probabilities of tie-deletion of dissimilar agents. Lines indicate the mean value of each time-step, aggregated over all parameters excluding the probability of tie-deletion and randomness. Shaded regions are the 95% confidence intervals. The three different plots show the different values of the probability of new connections being generated randomly. Notice that when $P(D) = 0$, no ties are deleted, which results in a bright blue line on the X-axis on all of the plots.

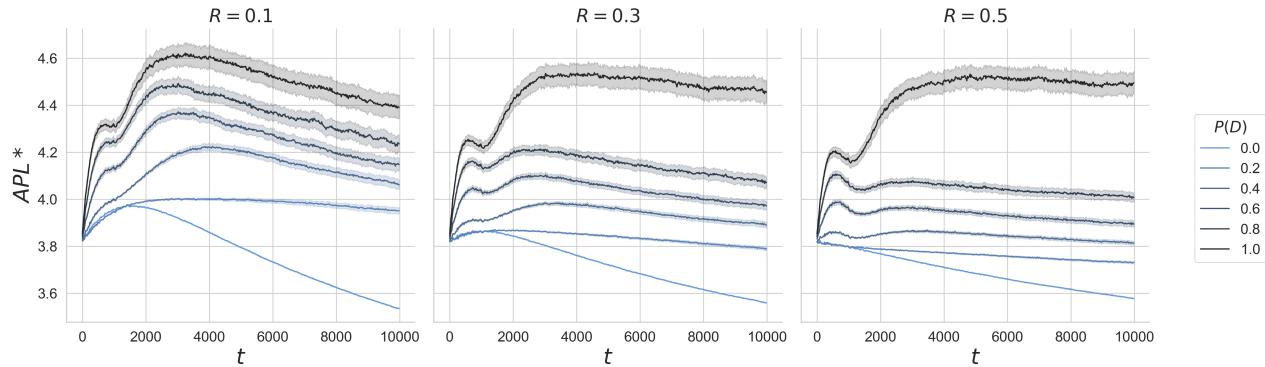


Figure A.20: $P(D)$'s effect on average path length over time. The x-axis shows the time-steps and the y-axis shows the estimated average path length of the simulated network. Colors indicate different probabilities of tie-deletion of dissimilar agents. Lines indicate the mean value of each time-step, aggregated over all parameters excluding the probability of tie-deletion and randomness. Shaded regions are the 95% confidence intervals. The three different plots show the different values of probability of new connections to be generated randomly.

A.3 The effect of initial opinions

The correlation between the initial opinion and the final opinion of agents, ρ_{O_I, O_F} , is dictated by the values of β and T . When β is low and T is high, there is little correlation between initial and final opinions. When β is high and T is low, final and initial values are highly correlated (see Appendix A.21).

The conditions where there is a strong positive correlation between initial and final opinions are conditions with the combination of strong negative social influence ($\beta \geq 0.15$) and lower threshold values ($T \leq 1$) (see Appendix A.21). These conditions are likely to lead to high degrees of polarization. The results indicate that in polarized conditions, agents will tend to move less than in conditions where they agents reach consensus. In conditions where agents do reach consensus, your starting position does not matter much. The final opinion of agents will by definition be very similar to all other agents. However, in polarizing conditions, polarization is primarily caused by negative social influence. These simulations are simulations with parameter values of $\beta \gg 0$. Under such conditions, the initial opinion and the final opinion of agents correlate positively. This is mainly due to the fact that agents will tend to move across the aisle, but be pushed to the nearest extreme (see Figure 14). Therefore, the initial opinion of agents becomes predictive of which extreme they will end up with.

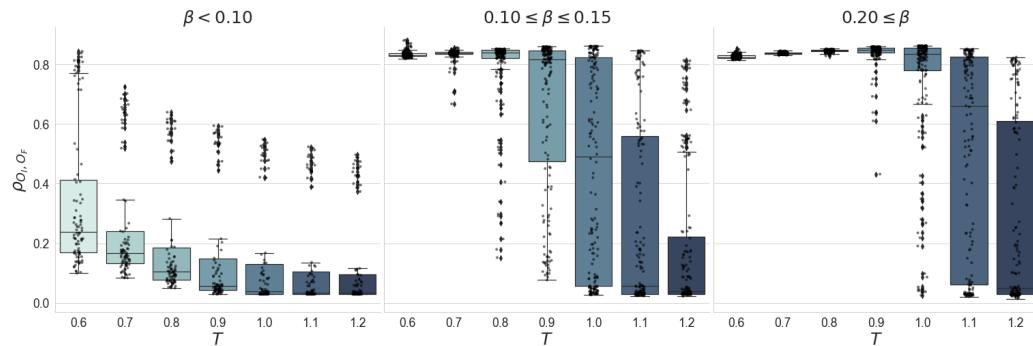


Figure A.21: Correlations between initial opinions and final opinions. X-axis shows different values of threshold and the y-axis shows the Pearson Correlation Coefficient between the initial and final opinion of agents. Boxes indicate the quartiles and whiskers indicate the 1.5 of the inter-quartile range outside of the quartiles. Points outside this range are depicted as diamonds. Smaller dark dots show the correlation coefficient for individual simulations. Columns show low, medium and high values of β .